

Lab3: Process Management

Yihua Xu

Overview

- Process scheduling algorithm:
 - Round-robin scheduling v.s. Priority scheduling
- Hints for lab 3
- Xv6 book link: <https://pdos.csail.mit.edu/6.828/2018/xv6/book-rev10.pdf>


Quick review

- Scheduler
 - OS module that manipulates the job queues, moving jobs to and from;
- Scheduling algorithm:
 - Determines which jobs are chosen to run next and what queues they wait on;
- Scheduler runs
 - When a job switches from running to waiting;
 - When an interrupt occurs;
 - When a job is created or terminated

scheduling algorithm

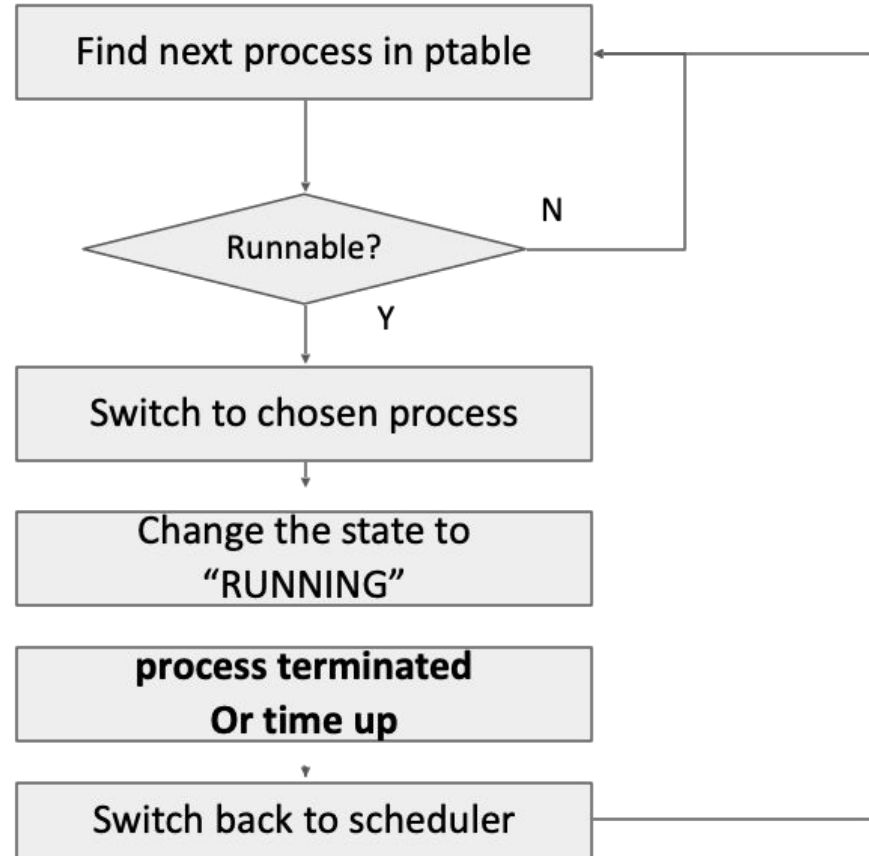

- Round Robin (RR)
 - Each task gets resource for a fixed period of time (time quantum). Unfinished task goes back in line

○



RR scheduler (current implementation)

when process calls
exit(), wait(), sleep()
etc.

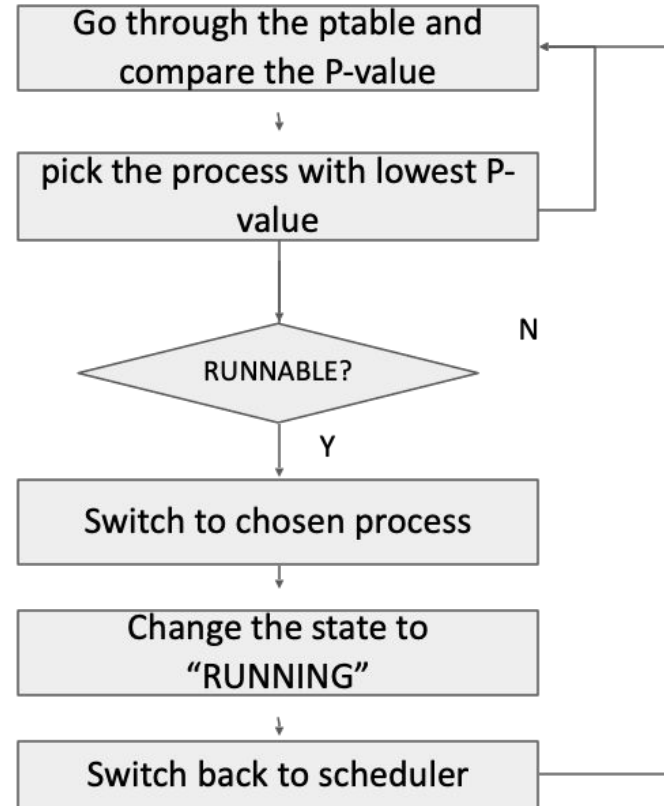


Priority scheduler (your task)

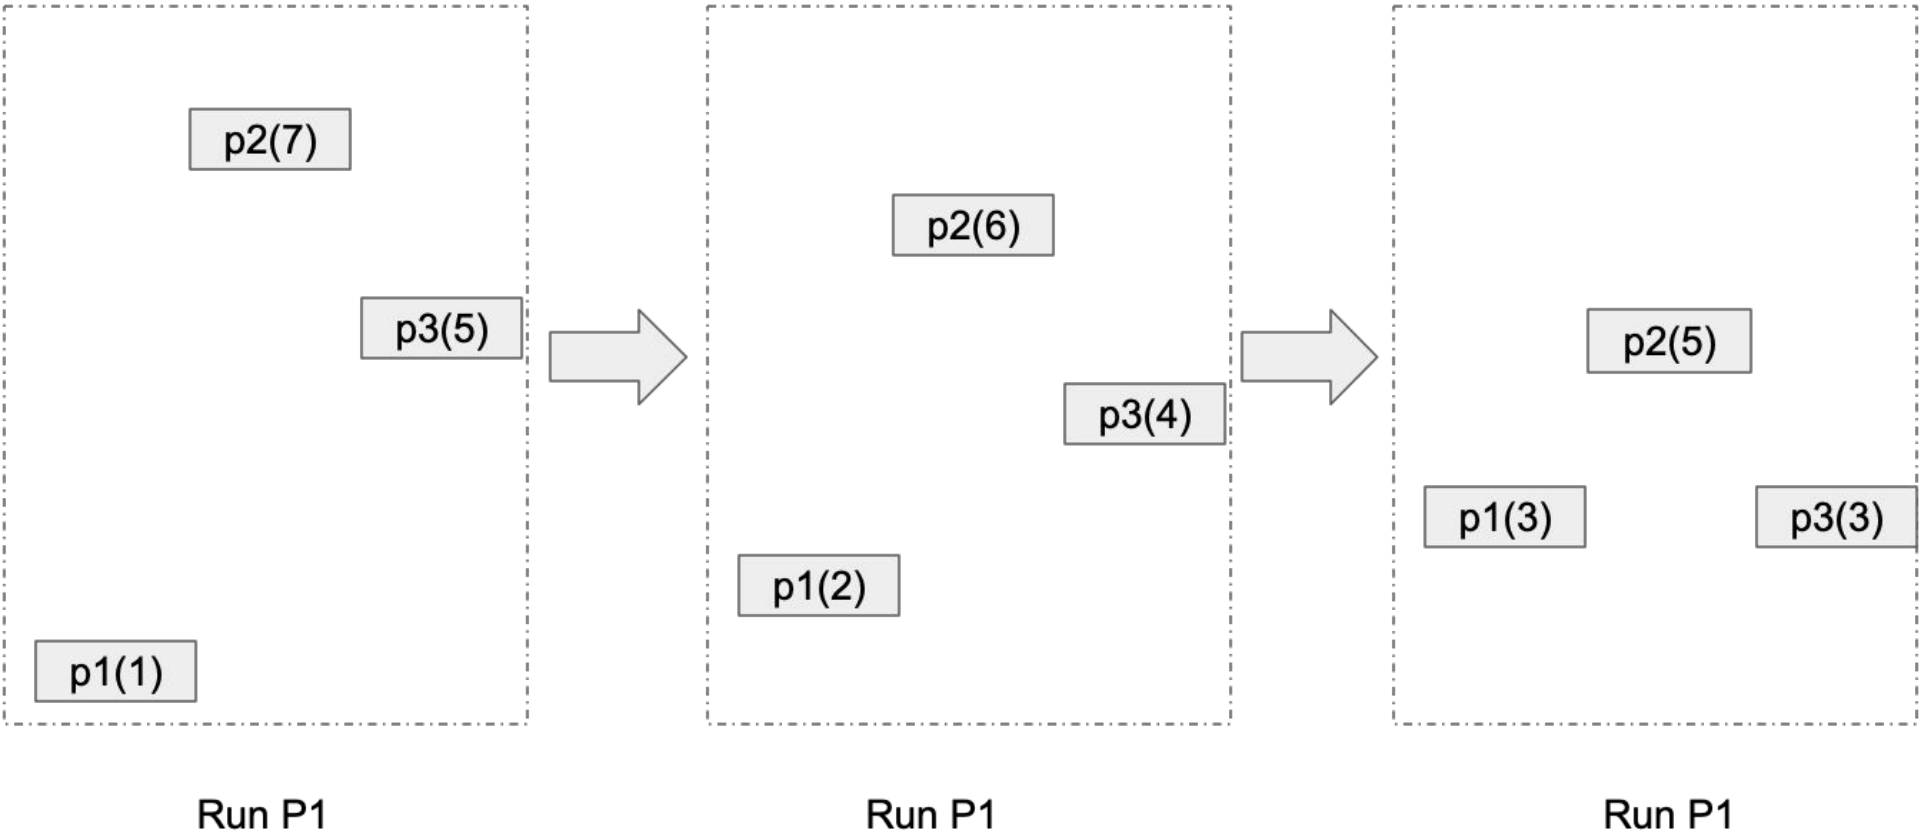
P-value range: [0, 31]

- 0: the highest priority;
- 31: the lowest priority;

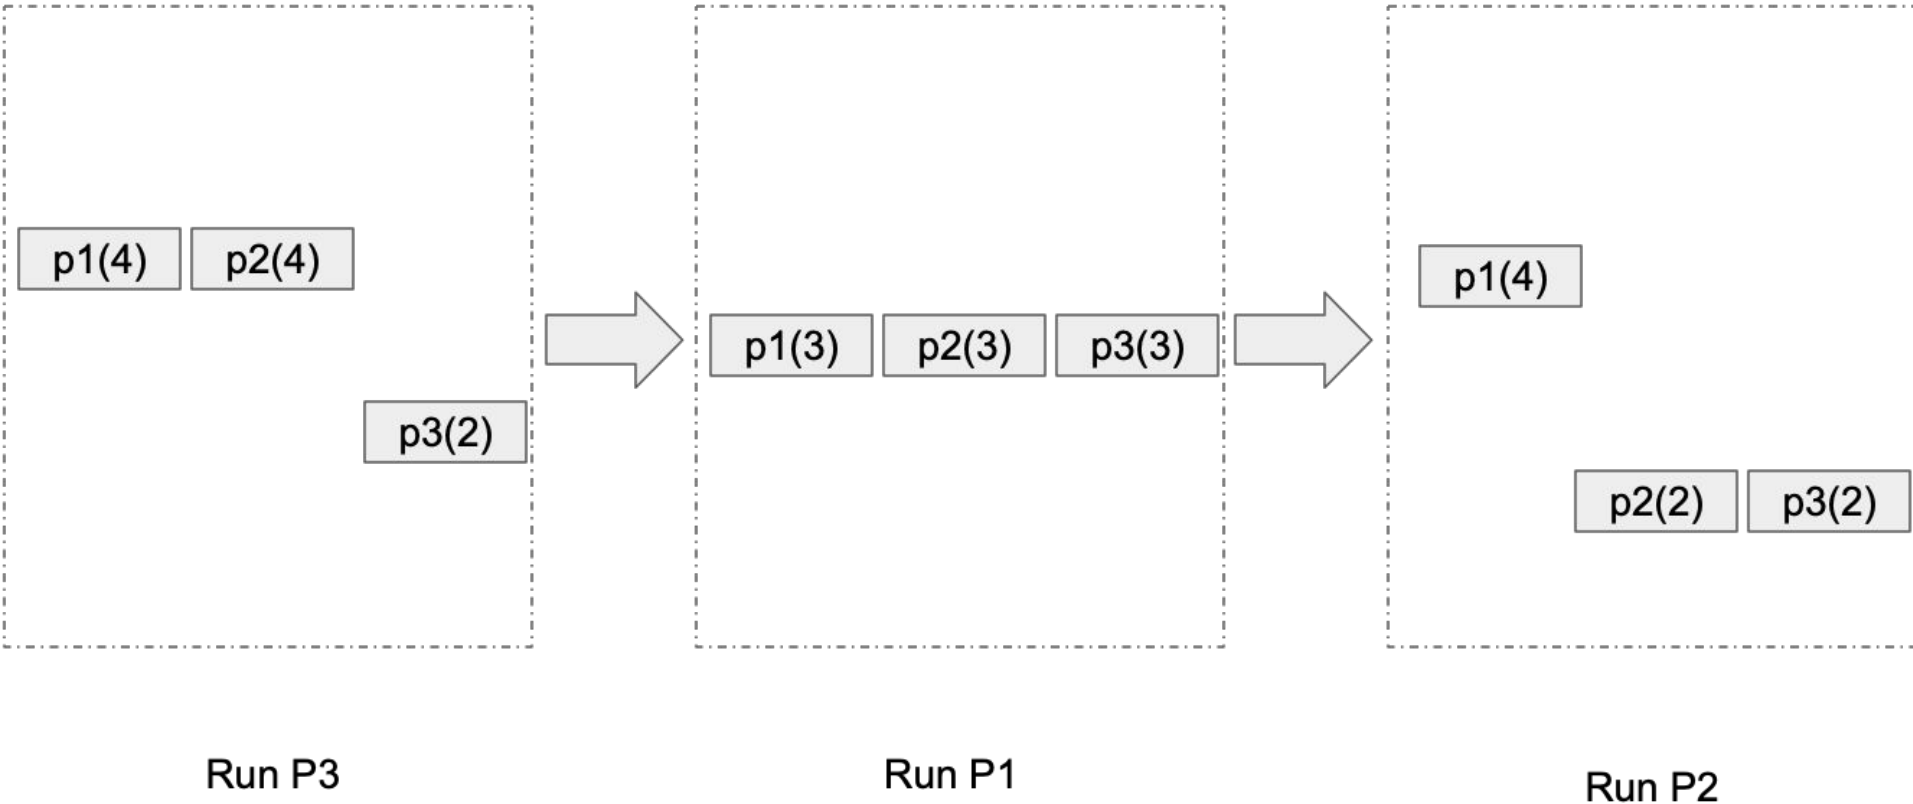
Implement aging of priority to avoid starvation.



How to age processes



How to age processes



Hints for lab3

1. Maintain a new field in proc structure to hold priority value;
2. add a new system call to modify the value anytime;
3. Modify round robin scheduler to priority scheduler;
4. Implement aging of priority to avoid starvation;
5. Compute the turnaround time and wait time for each process;
6. Develop a test program to illustrate priority scheduling.

Step 1: add new field to proc structure

- Add another field in proc struct:
 - E.g. int prior_val;
 - Value range: [0, 31]; 0: highest; 31: lowest
- Initialization of prior_val:
 - allocproc()
 - fork()
 - child should inherit parent's priority value
- Updating prior_val:
 - New system call;
 - Aging of priority;
 - Hint: to give up CPU resources when current proc's priority is lower than other procs:

```
C proc.h > proc
37 // Per-process state
38 struct proc {
39     uint sz; // Size of process memory (bytes)
40     pde_t* pgdir; // Page table
41     char *kstack; // Bottom of kernel stack for this process
42     enum procstate state; // Process state
43     int pid; // Process ID
44     struct proc *parent; // Parent process
45     struct trapframe *tf; // Trap frame for current syscall
46     struct context *context; // switch() here to run process
47     void *chan; // If non-zero, sleeping on chan
48     int killed; // If non-zero, have been killed
49     struct file *ofile[NOFILE]; // Open files
50     struct inode *cwd; // Current directory
51     char name[16]; // Process name (debugging)
52 };
```


Step 2: add system call to update priority value

- Follow the same steps in Lab2;
- `void update_prior(int prior_lvl)` OR `int update_prior(int prior_lvl)`
 - change the priority value of the current proc;
 - Then transfer control to scheduler immediately (call `sched()`) because the priority list has been changed;
 - Pay attention to :
 - the lock, or directly call `yield()`;
 - `sched()` is not `scheduler()`

Step 3: modify scheduler()

Round robin

- Loop over process table
 - Only choose the “RUNNABLE” process, which is ready to execute;
- Next RUNNABLE process p acquires CPU resource and start RUNNING;



Priority scheduling:

- Loop over process table
 - find RUNNABLE proc with the highest priority;
- The chosen RUNNABLE process acquires CPU resources and start RUNNING;

```
C proc.c > scheduler(void)
322 void
323 scheduler(void)
324 {
325     struct proc *p;
326     struct cpu *c = mycpu();
327     c->proc = 0;
328
329     for(;;){
330         // Enable interrupts on this processor.
331         sti();
332
333         // Loop over process table looking for process to run.
334         acquire(&ptable.lock):
335         for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
336             if(p->state != RUNNABLE)
337                 continue;
338
339             // Switch to chosen process. It is the process's job
340             // to release ptable.lock and then reacquire it
341             // before jumping back to us.
342             c->proc = p;
343             switchuvm(p);
344             p->state = RUNNING;
345
346             switch(&(c->scheduler), p->context);
347             switchkvm();
348
349             // Process is done running for now.
350             // It should have changed its p->state before coming back.
351             c->proc = 0;
352         }
353         release(&ptable.lock);
354     }
355 }
356 }
```

Step 4: aging of priority

- If a process waits increase its priority (decrease its value). When it runs, decrease it (increase its value).
- Modify scheduler() function to update procs' priority value;
- Pay attention to the legal range of priority value: [0, 31]

Step 5: turnaround & waiting time

- Turnaround time = $T_{\text{finish}} - T_{\text{start}}$
 - Add fields inside proc struct for tracking scheduling performance of each proc
 - Acquire time ticket from `exec()` => proc start time
 - Acquire time ticket from `exit()` => proc finish time
- Waiting time = Turnaround time - burst time
 - burst time: how many times it's been scheduled, each time is a time ticket;

```
49  switch(tf->trapno){
50  case T_IRQ0 + IRQ_TIMER:
51      if(cpuid() == 0){
52          acquire(&tickslock);
53          ticks++;
54          wakeup(&ticks);
55          release(&tickslock);
56      }
```

Step 6: test your scheduler

- Test programs should be able to run for a while;
- To run multiple programs simultaneously:
 - \$ proc1&;proc2&;proc3
 - Try different order (\$ proc3&;proc1&;proc2)
- Compare the turnaround time and waiting time of processes with different priority;
- Illustrate priority scheduling with the result;

```
5 int main(int argc, char *argv[]) {
6     if (argc >= 1){
7         //int new = atoi(argv[1]);
8         update_prior_val(1);
9         int limit = 4300;
10        int i, j;
11        for (i=0;i<limit;i++) {
12            asm("nop");
13            for (j=0;j<limit;j++) {
14                asm("nop");
15            }
16        }
17    }
18    exit();
19 }
```