

## LAB 2: FUN WITH SYSTEM CALLS

a) Change the exit system call signature to void exit (int status):

Changes made in the following files:

```
Open  user.h
~/xv6-public

1 struct stat;
2 struct rtcdate;
3
4 // system calls
5 int fork(void);
6 int exit(int) __attribute__((noreturn));
```

```
Open  defs.h
~/xv6-public

102 int pipewrite(struct pipe, char, int);
103
104 //PAGEBREAK: 16
105 // proc.c
106 int cputid(void);
107 void exit(int);
```

```
Open  *sysproc.c
~/xv6-public

15
16 int
17 sys_exit(void)
18 {
19     int exit_status;
20     argint(0, &exit_status);
21     exit(exit_status);
22     return 0; // not reached
23 }
```

```
Open  proc.h
~/xv6-public  Save

40 pde_t* pgdir; // Page table
41 char *kstack; // Bottom of kernel stack for this process
42 enum procstate state; // Process state
43 int pid; // Process ID
44 struct proc *parent; // Parent process
45 struct trapframe *tf; // Trap frame for current syscall
46 struct context *context; // switch() here to run process
47 void *chan; // If non-zero, sleeping on chan
48 int killed; // If non-zero, have been killed
49 struct file *ofile[NOFILE]; // Open files
50 struct inode *cwd; // Current directory
51 char name[16]; // Process name (debugging)
52 int exit_status; // exit status for proc
53 ;;
```

```

pen  ▼  [icon]  *proc.c
                  ~/xv6-public
void
exit(int status)
{
    struct proc *curproc = myproc();
    struct proc *p;
    int fd;

    if(curproc == initproc)
    {
        curproc->exit_status = 1; //error give 1 exit status
        panic("init exiting");
    }
    // Close all open files.
    for(fd = 0; fd < NOFILE; fd++){
        if(curproc->ofile[fd]){
            fileclose(curproc->ofile[fd]);
            curproc->ofile[fd] = 0;
        }
    }

    begin_op();
    iput(curproc->cwd);
    end_op();
    curproc->cwd = 0;

    acquire(&ptable.lock);

    // Parent might be sleeping in wait().
    wakeup1(curproc->parent);

    // Pass abandoned children to init.
    for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
        if(p->parent == curproc){
            p->parent = initproc;
            if(p->state == ZOMBIE)
                wakeup1(initproc);
        }
    }

    // Jump into the scheduler, never to return.
    curproc->state = ZOMBIE;
    sched();
    panic("zombie exit");
}

```

The following changes were made for the exit () function:

- In cat.c lines 15, 20, and 37 were changed to exit (-1); Lines 31 and 42 were changed to exit (0)
- In echo.c line 12 was changed to exit (0)
- In forktest.c lines 33, 39 and 45 were changed to exit (-1); Lines 28 and 55 were changed to exit (0)
- In grep.c lines 46, 52, and 58 were changed to exit (-1); Line 63 was changed to exit (0)
- In init.c lines 27 and 32 were changed to exit (-1)
- In kill.c lines 12 and 16 were changed to exit (0)
- In ln.c lines 10 was changed to exit (-1); Line 14 was changed to exit (0)
- In ls.c lines 80 was changed to exit (27); Line 84 was changed to exit (1)
- In mkdir.c lines 22 was changed to exit (0); Line 12 was changed to exit (-1)
- In rm.c lines 22 was changed to exit (0); Line 12 was changed to exit (-1)
- In sh.c lines 77, 87, 171, and 178 were changed to exit (0); Line 68 and 130 was changed to exit (1)
- In stressfs.c line 48 was changed to exit (0)

- In trap.c lines 41, 45, 101, and 111 were changed to exit (0)
- In usertests.c lines 24, 28, 32, 36, 52, 57, 61, 65, 92, 97, 110, 127, 150, 155, 159, 169, 176, 182, 197, 204, 213, 222, 227, 239, 273, 278, 283, 288, 299, 313, 324, 346, 352, 450, 511, 533, 540, 547, 566, 574, 596, 607, 614, 634, 637, 665, 673, 677, 686, 690, 694, 714, 718, 724, 730, 736, 740, 790, 810, 814, 824, 832, 868, 904, 915, 927, 944, 950, 962, 968, 976, 981, 987, 992, 1001, 1005, 1009, 1013, 1019, 1023, 1095, 1099, 1107, 1111, 1130, 1158, 1164, 1172, 1179, 1185, 1189, 1196, 1213, 1217, 1222, 1228, 1250, 1254, 1266, 1278, 1293, 1324, 1354, 1358, 1402, 1408, 1431, 1439, 1445, 1458, 1468, 1473, 1481, 1486, 1493, 1502, 1507, 1516, 1540, 1585, 1603, 1633, 1639, 1718, 1721, and 1733 were changed to exit(1); Lines 67, 103, 105, 133, 232, 327, 421, 454, 487, 550, 618, 746, 752, 757, 795, 847, 886, 957, 996, 1029, 1034, 1038, 1042, 1046, 1050, 1054, 1058, 1062, 1066, 1070, 1074, 1078, 1082, 1086, 1090, 1103, 1137, 1234, 1238, 1258, 1262, 1270, 1274, 1298, 1303, 1308, 1312, 1316, 1320, 1330, 1335, 1391, 1396, 1448, 1575, 1630, 1755, and 1802 were changed to exit(0)
- In wc.c lines 30 and 48 were changed to exit (1); Lines 42 and 53 were changed to exit (0)
- In zombie.c line 13 was changed to exit (0);

The xv6 compiles once make qemu or make qemu-nox is entered.

```

apalayil@ubuntu: ~/xv6-public
apalayil@ubuntu:~/xv6-public$ make qemu
qemu-system-i386 -serial mon:stdio -drive file=fs.img,index=1,media=disk,format=
xv6...
cpu1: starting 1
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap sta8
init: starting sh
$

```

b) Update the wait system call signature to it wait (int \*status):

```

Open  user.h
~/xv6-public
1 struct stat;
2 struct rtcdate;
3
4 // system calls
5 int fork(void);
6 int exit(int) __attribute__((noreturn));
7 int wait(int*);

```

```

Open  defs.h
~/xv6-public
115 void scheduler(void) __attribute__((noreturn));
116 void sched(void);
117 void setproc(struct proc*);
118 void sleep(void*, struct spinlock*);
119 void userinit(void);
120 int wait(int*);

```

```

Open  ▾  [+]
```

```

24  exit(exit_status);
25  return 0; // not reached
26 }
27
28 int
29 sys_wait(void)
30 {
31     char *pointer; //arg ptr requires char ptr
32     argptr(0,&pointer,32);
33     return wait((int *)pointer);
34 }

```

\*sysproc.c  
~/xv6-public

```

Open  ▾  [+]
```

```

5 //TODO consider status update
5 int
7 wait(int *status)
8 {
9     struct proc *p;
9     int havekids, pid;
1    struct proc *curproc = myproc();
2
3    acquire(&ptable.lock);
4    for(;;){
5        // Scan through table looking for exited children.
5        havekids = 0;
7        for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
8            if(p->parent != curproc)
9                continue;
9            havekids = 1;
1           if(p->state == ZOMBIE){
2               // Found one.
3               //for exit status
3               pid = p->pid;
4               kfree(p->kstack);
5               p->kstack = 0;
6               freevm(p->pgdir);
7               p->pid = 0;
8               p->parent = 0;
9               p->name[0] = 0;
1              p->killed = 0;
2              p->state = UNUSED;
3              if (status != 0) {
4                  *status = p->exit_status; //Exit Status
5              }
5              p->exit_status = 0;
7              release(&ptable.lock);
8              return pid;
9          }
10     }
11
12     // No point waiting if we don't have any children.
13     if(!havekids || curproc->killed){
14         release(&ptable.lock);
15         return -1;
16     }
17
18     // Wait for children to exit. (See wakeup1 call in proc_exit.)
19     sleep(curproc, &ptable.lock); //DOC: wait-sleep
20 }

```

\*proc.c  
~/xv6-public

The following changes were made for the wait () function:

- In sh.c lines 96, 120, 121, and 169 were changed to wait (&stat)
- In stressfs.c line 46 was changed to wait (&stat)
- In forktest.c lines 37 and 43 was changed to wait (&stat)

```

apalayil@ubuntu: ~/xv6-public
```

```

apalayil@ubuntu:~/xv6-public$ make qemu
qemu-system-i386 -serial mon:stdio -drive file=fs.img,index=1,media=disk,format=
xv6...
cpu1: starting 1
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap sta8
init: starting sh
$

```

c) Add a waitpid system call: int waitpid (int pid, int \*status, int options):

```
Open user.h ~/xv6-public
23 char *strk(int);
24 int sleep(int);
25 int uptime(void);
26 int waitpid(int, int*, int);
```

```
Open defs.h ~/xv6-public
0 int wait(int*);
1 void wakeup(void*);
2 void yield(void);
3 int waitpid(int, int*, int);
```

```
Open *sysproc.c ~/xv6-public-apalayil Save
*mytests.c *sysproc.c
0
1 int
2 sys_waitpid(void)
3 {
4     int pid, options; // default value
5     int* status;
6     if(argint(0, &pid) < 0){
7         return -1;
8     }
9     if(argptr(1, (void*)&status, sizeof(status)) < 0){
10        return -1;
11    }
12    argint(2, &options);
13    return waitpid(pid, (int*)status, options);
14 }
```

```
Open *proc.c ~/xv6-public Save
580 // timer to and sleep call
581 p->parent = curproc;
582 sleep(curproc, &table.lock); //does not go away - goes to
583 }
584 */
585 int
586 waitpid(int pid, int* status, int options)
587 {
588     struct proc *p, *curproc = myproc();
589     int found_pid;
590     acquire(&table.lock);
591     //Loops continuously till the process with given pid is terminated
592     for(;;) {
593         found_pid = 0;
594         for(p = table.proc; p < &table.proc[NPROC]; p++) {
595             //processes will check if ZOMBIE state, else it will inherit it.
596             if(p->pid != pid) continue;
597             found_pid = 1;
598             if(p->state == ZOMBIE) {
599                 //Found the process with the given pid that has exited.
600                 kfree(p->kstack);
601                 p->kstack = 0;
602                 freevm(p->pgdir);
603                 p->pid = 0;
604                 p->parent = 0;
605                 p->name[0] = 0;
606                 p->killed = 0;
607                 p->state = UNUSED;
608                 if(status) *status = p->exitstatus;
609                 p->exitstatus = 0;
610                 release(&table.lock);
611                 return pid;
612             } else if(options == 1) {
613                 release(&table.lock);
614                 return 0;
615             }
616         }
617         // or the current process is killed.
618         if(!found_pid || curproc->killed) {
619             release(&table.lock);
620             return -1;
621         }
622         // Wait for the process with the given pid to exit.
623         sleep(curproc, &table.lock);
624     }
625 }
```



```
apalayil@ubuntu: ~/xv6-public
SeaBIOS (version 1.13.0-1ubuntu1.1)

iPXE (http://ipxe.org) 00:03.0 CA00 PCI2.10 PnP PMM+1FF8CA10+1FECCA10 CA00

Booting from Hard Disk..xv6...
cpu1: starting 1
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap sta8
init: starting sh
$
```

d) Write an example program to illustrate that your waitpid works:

```
Open mytests.c Save
~/xv6-public/apalayil

1
2#include "types.h"
3#include "stat.h"
4#include "user.h"
5#define NULL 0
6
7void exitTest(int status) {
8    int pid;
9    int return_pid, fork_count = 3;
10   int exit_status;
11   int i;
12   // use this part to test exit(int status) and wait(int* status)
13   for(i = 0; i < fork_count; i++)
14   {
15       pid = fork();
16       if(pid == 0){ // only the child executed this code
17           if(i == 0){
18               printf(1, "\nThis is child with PID# %d ,exit with status %d\n", getpid(), 0);
19               exit(0);
20           }
21           else{
22               printf(1, "\nThis is child with PID# %d ,exit with status %d\n", getpid(),
23                   -1);
24               exit(-1);
25           }
26       } else if (pid > 0){ // only the parent executes this code
27           return_pid = wait(&exit_status);
28           printf(1, "\nThis is the parent: child with PID# %d has exited with status
29               %d\n", return_pid, exit_status);
30       } else{ // something went wrong with fork system call
31           printf(2, "\n Error using fork\n");
32           exit(-1);
33       }
34   }
35 }

36void waitpidTest(void) {
37   int exit_status;
38   int i;
39   int ret_fork;
40   // use this part to test wait(int pid, int* status, int options)
41   int cpid;
42   // fork 6 children
43   for(i=0; i<6;i = i +1) {
44       ret_fork = fork();
45       if(ret_fork == 0) { // for some reason ref pid[_] in any way breaks the loop
46           printf(1, "Child Process: %d\n",i);
47           int exit_stat = 100 + i;
48           exit(exit_stat); //exit right away
49       }
50   }
51   cpid = waitpid(6,&exit_status,0);//note the pids will range for 4 to 10
52   printf(1,"Child with pid: %d, exited with status: %d\n",cpid, exit_status);
53   cpid = waitpid(8,&exit_status,0);//note the pids will range for 4 to 10
54   printf(1,"Child with pid: %d, exited with status: %d\n",cpid, exit_status);
55   //clean up the rest of the processes
56   int has_children = 1;
57   while(has_children != -1) {
58       has_children = wait(NULL);
59   }
60 }
61
62int main() {
63   exitTest(NULL);
64   waitpidTest();
65   exit(0);
66 }
```

Compiling xv6-public and running the test files:

```
apalayil@ubuntu: ~/xv6-public-apalayil
SeaBIOS (version 1.13.0-1ubuntu1.1)

iPXE (http://ipxe.org) 00:03.0 CA00 PCI2.10 PnP PMM+1FF8CA10+1FECCA10 CA00

Booting from Hard Disk..xv6...
cpu1: starting 1
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap sta8
init: starting sh
$ mytest
```

```
apalayil@ubuntu: ~/xv6-public-apalayil
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap sta8
init: starting sh
$ mytests

This is child with PID# 4 ,exit with status 0
This is the parent: child with PID# 4 has exited with status 0
This is child with PID# 5 ,exit with status -1
This is the parent: child with PID# 5 has exited with status -1
This is child with PID# 6 ,exit with status -1
This is the parent: child with PID# 6 has exited with status -1
Child Process: 0
Child Process: 1
Child Process: 2
Child Process: 3
Child Process: 4
CChhllldd Pwriotche spisd: : 5-1
, exited with status: -1
Child with pid: 8, exited with status: 101
$
```

Thus the results show when the parent process waits for the child process to finish printing after which it exits.