

ILLINOIS TECH

College of Computing


CS 450 Operating Systems Concurrency Bugs

Yue Duan

Recap

- Semaphores are equivalent to locks + condition variables
 - Can be used for both mutual exclusion and ordering
- Semaphores contain **state**
 - How they are initialized depends on how they will be used
 - Init to 0: Join (1 thread must arrive first, then other)
 - Init to N: Number of available resources
- `sem_wait()`: Decrement and waits **if** value < 0
- `sem_post()`: Increment value, then wake a single waiter (**atomic**)
- Use semaphores in
 - producer/consumer
 - reader/writer

Concurrency Bugs



- Lu et al. [ASPLOS 2008]:
 - For four major projects, search for concurrency bugs among >500K bug reports.
 - Analyze small sample to identify common types of concurrency bugs.

Atomicity: MySQL

Thread 1:

```
if (thd->proc_info) {  
    ...  
    fputs(thd->proc_info, ...);  
    ...  
}
```

Thread 2:

```
thd->proc_info = NULL;
```

- Test (`thd->proc_info != NULL`) and set (writing to `thd->proc_info`) should be **atomic**

Atomicity: MySQL

Thread 1:

```
pthread_mutex_lock(&lock);  
if (thd->proc_info) {  
    ...  
    fputs(thd->proc_info, ...);  
    ...  
}  
pthread_mutex_unlock(&lock);
```

Thread 2:

```
pthread_mutex_lock(&lock);  
thd->proc_info = NULL;  
pthread_mutex_unlock(&lock);
```

- Fix it with locks

Ordering: Mozilla

Thread 1:

```
void init() {  
    ...  
    mThread =  
        PR_CreateThread(mMain, ...);  
    ...  
}
```

Thread 2:

```
void mMain(...) {  
    ...  
    mState = mThread->State;  
    ...  
}
```

- What's wrong?
 - Thread 1 sets value of mThread needed by Thread2
 - How to ensure that reading mThread happens after mThread initialization?

Ordering: Mozilla

Thread 1:

```
void init() {  
    ...  
  
    mThread =  
    PR_CreateThread(mMain, ...);  
  
    pthread_mutex_lock(&mtLock);  
    mtInit = 1;  
    pthread_cond_signal(&mtCond);  
    pthread_mutex_unlock(&mtLock);  
  
    ...  
}
```


Thread 2:

```
void mMain(...) {  
    ...  
  
    mutex_lock(&mtLock);  
    while (mtInit == 0)  
        Cond_wait(&mtCond, &mtLock);  
    Mutex_unlock(&mtLock);  
  
    mState = mThread->State;  
    ...  
}
```

- Fix ordering bugs with condition variables

Dining philosophers problem

- Five philosopher setting around a table
- Five forks
 - each between two philosophers
- Each philosopher will do:
 - think: no fork needed
 - get forks: two forks needed
 - eat
 - put down forks



```
while (1) {  
    think();  
    get_forks(p);  
    eat();  
    put_forks(p);  
}
```


Dining philosophers problem

- Write the routines **get_forks()** and **put_forks()** such that:
 - no deadlock
 - no starvation
 - high concurrency
- Use five semaphores
 - one for each fork `sem_t forks[5]`
- Some helper functions

```
int left(int p) { return p; }  
int right(int p) { return (p + 1) % 5; }
```

Dining philosophers problem

- First attempt
 - initialize each semaphore to a value of 1
 - assume that each philosopher knows its own number (p)

```
void get_forks(int p) {  
    sem_wait(&forks[left(p)]);  
    sem_wait(&forks[right(p)]);  
}  
  
void put_forks(int p) {  
    sem_post(&forks[left(p)]);  
    sem_post(&forks[right(p)]);  
}
```

- Does it work?
 - no, but what's wrong?

Dining philosophers problem

- A solution: breaking the dependency
 - change how forks are acquired by **at least one** of the philosophers

```
void get_forks(int p) {  
    if (p == 4) {  
        sem_wait(&forks[right(p)]);  
        sem_wait(&forks[left(p)]);  
    } else {  
        sem_wait(&forks[left(p)]);  
        sem_wait(&forks[right(p)]);  
    }  
}
```

Deadlock

- No progress can be made because **two or more threads are waiting for the other** to take some action and thus neither ever does

Thread 1:


```
lock(&A);  
lock(&B);
```

Thread 2:

```
lock(&B);  
lock(&A);
```

Deadlock

- Circular Dependency



Fix Deadlock Code

Thread 1:

```
lock(&A);  
lock(&B);
```

Thread 2:

```
lock(&B);  
lock(&A);
```


Thread 1

```
lock(&A);  
lock(&B);
```

Thread 2

```
lock(&A);  
lock(&B);
```

Non-circular Dependency (fine)



Deadlock


```
set_t *set_intersection (set_t *s1, set_t *s2) {  
    set_t *rv = malloc(sizeof(*rv));  
    mutex_lock(&s1->lock);  
    mutex_lock(&s2->lock);  
    for(int i=0; i<s1->len; i++) {  
        if(set_contains(s2, s1->items[i])  
            set_add(rv, s1->items[i]);  
    mutex_unlock(&s2->lock);  
    mutex_unlock(&s1->lock);  
}
```

Thread 1: `rv = set_intersection(setA, setB);`

Thread 2: `rv = set_intersection(setB, setA);`

Deadlock Theory

- Deadlocks can only occur when **all four** conditions are true:
 - 1) Mutual exclusion
 - 2) Hold-and-wait
 - 3) Circular wait
 - 4) No preemption
- Eliminate deadlock by eliminating **any one** condition



1). Mutual Exclusion

- Definition: “Threads claim **exclusive control of resources** that they require (e.g., thread grabs a lock)”
- Strategy: **eliminate locks!**
 - try to replace locks with atomic instructions

```
int CompAndSwap(int *addr, int expected, int new);  
Returns 0: fail, 1: success
```

Code with locks

```
void add (int *val, int amt)  
{  
    mutex_lock(&m);  
    *val += amt;  
    mutex_unlock(&m);  
}
```

Code with Compare-and-Swap (CAS)

```
void add (int *val, int amt)  
{  
    do {  
        int old = *value;  
    } while(!CAS(val, old, old+amt));  
}
```

Example: Lock-Free Linked List Insert

Code with locks

```
void insert (int val)
{
    node_t *n =
        malloc(sizeof(*n));
    n->val = val;
    mutex_lock(&m);
    n->next = head;
    head = n;
    mutex_unlock(&m);
}
```

Code with Compare-and-Swap (CAS)

```
void insert (int val)
{
    node_t *n = malloc(sizeof(*n));
    n->val = val;
    do {
        n->next = head;
    } while (!CAS(&head, n->next, n));
}
```

2). Hold-and-Wait

- Definition: “Threads **hold** resources allocated to them (e.g., locks they have already acquired) **while waiting for additional resources** (e.g., locks they wish to acquire).”
- Strategy: release **currently held resources** when waiting for new ones

2). Hold-and-Wait

- How? Use a meta lock

```
lock(&meta);  
lock(&L1);  
lock(&L2);  
...  
unlock(&meta);  
  
// Critical section code  
  
unlock(...);
```

2). Hold-and-Wait

- Disadvantages?
 - Must know ahead of time which locks will be needed
 - Must be conservative (acquire any lock possibly needed)
 - Degenerates to just having one big lock

3). No preemption

- Problem: Resources (e.g., locks) **cannot be forcibly removed** from threads that are holding them
- Strategy: if thread can't get what it wants, release what it holds

```
top:
    lock(A);
    if (trylock(B) == -1) {
        unlock(A);
        goto top;
    }
```

3). No preemption

- Potential issue:
 - **livelock**: no process makes forward progress, but the state of involved processes constantly changes
 - Can happen if all processes release resources and then try to re-acquire, fail, and keep doing this
 - Classic solution: back-off techniques
 - **Random back-off**: wait for a random amount of time before retrying
 - **Exponential back-off**: wait for exponentially increasing amount of time before retrying

4). Circular Wait

- Definition: “There exists a **circular chain of threads** such that each thread holds a resource (e.g., lock) being requested by next thread in the chain.”
- Usually the easiest deadlock requirement to attack
- Strategy: impose a well-documented order of acquiring locks
 - decide which locks should be acquired before others
 - if A before B, never acquire A if B is already held!
 - document this, and write code accordingly

4). Circular Wait

Thread 1

```
lock (&A) ;  
lock (&B) ;
```

Thread 2

```
lock (&B) ;  
lock (&A) ;
```

How would you fix this code?

Thread 1

```
lock (&A) ;  
lock (&B) ;
```

Thread 2

```
lock (&A) ;  
lock (&B) ;
```

Lock Ordering in Linux

- Works well if system has distinct layers

```
In linux-3.2.51/include/linux/fs.h  
/* inode->i_mutex nesting subclasses for the lock  
 * validator:  
 * 0: the object of the current VFS operation  
 * 1: parent  
 * 2: child/target  
 * 3: quota file  
 * The locking order between these classes is  
 * parent -> child -> normal -> xattr -> quota  
 */
```

THANK YOU!