

ILLINOIS TECH

College of Computing

CS 450 Operating Systems **Condition Variables**

Yue Duan

Concurrency Goals

- Mutual Exclusion
 - Keep two threads from executing in a critical section concurrently
 - We solved this with **locks**
- Dependent Events
 - We want a thread to wait until some particular event has occurred
 - Or some condition has been met
 - Solved with **condition variables and semaphores**

Example: join()

```
pthread_t p1, p2;  
pthread_create(&p1, NULL, mythread, "A");  
pthread_create(&p2, NULL, mythread, "B");  
// join waits for the threads to finish  
pthread_join(p1, NULL);  
pthread_join(p2, NULL);  
printf("Main: done\n [balance: %d]\n", balance);  
return 0;
```

join(): parent must wait for child thread to finish

How to implement?

Waiting for an Event

- Parent thread has to wait until child terminates
- **Option 1:** spin until that happens
 - Waste of CPU time
- **Option 2:** wait (sleep) in a queue until that happens
 - Better use of CPU time
 - Similar to the idea in queue-based lock in the previous lecture
 - Child thread will signal the parent to wake up before its termination

Condition Variables

- CV:
 - queue of waiting threads
- **B** waits for a signal on CV before running
 - `wait(CV, ...);`
- **A** sends `signal()` on CV when time for **B** to run
 - `signal(CV, ...);`

API

- **cond_wait**(cond_t * cv, mutex_t * lock)
 - assumes lock is held when wait() is called
 - puts caller to sleep + releases the lock (atomically)
 - when awoken, reacquires lock before returning
- **cond_signal**(cond_t * cv)
 - wake a single waiting thread (if ≥ 1 thread is waiting)
 - if there is no waiting thread, NOP

Thread Join: Attempt 1

Parent

```
void thr_join() {  
    cond_wait(&c);  
}
```

Child

```
void thr_exit() {  
    cond_signal(&c);  
}
```

- Does this work? If not, what's the problem?
- Child may run and call **cond_signal()** before parent called **cond_wait()**
 - Parent will sleep indefinitely.

Thread Join: Attempt 2

Parent

```
void thr_join() {  
    if (done == 0) {  
        cond_wait(&c);  
    }  
}
```

Child

```
void thr_exit() {  
    done = 1;  
    cond_signal(&c);  
}
```

- Let's keep some **state** then
- Is there a problem here?

Parent:

a

b

Child:

x

y

Using Locks to Achieve Atomicity

Waiting Thread

```
mutex_lock(&m);  
if (!check_cond())  
    cond_wait(&c, &m);  
...  
mutex_unlock(&m);
```

Waking Thread

```
mutex_lock(&m);  
set_cond();  
cond_signal(&c);  
...  
mutex_unlock(&m);
```

- Need a lock (mutex) to ensure two things
 - Checking condition (waiting thread) & modifying it (waking thread) remain mutually exclusive
 - Checking condition & putting thread to sleep (waiting thread) remain atomic

Using Locks to Achieve Atomicity

Waiting Thread

```
mutex_lock(&m);  
if (!check_cond())  
    cond_wait(&c, &m);  
...  
mutex_unlock(&m);
```

Waking Thread

```
mutex_lock(&m);  
set_cond();  
cond_signal(&c);  
...  
mutex_unlock(&m);
```

- **cond_wait()** should unlock mutex atomically w/ going to sleep
 - If mutex not released, waking thread cannot make progress
 - If release is not atomic, we get a race condition. Can you identify it?

Recap: CV Rules of Thumb (Take 1)

- Shared state determines if condition is true or not
- Check the state before waiting on cv
- Use a mutex to protect
 - 1) the shared state on which condition is based, as well as
 - 2) operations on the cv
- Remember to acquire the mutex before calling `cond_signal()`

Bounded Buffer (Producer/Consumer)

- Multiple producers and multiple consumers communicate using a shared, finite-size buffer
- Producers add items to buffer
 - If buffer is full, producer has to wait until there is free space
- Consumers remove items from buffer
 - If buffer is empty, consumer has to wait until one or more items are added
- Common examples:
 - Unix pipe: bounded buffer in kernel (multiple producers & consumers)
 - Work queue in a web server (one producer, multiple consumers)

Bounded Buffer (Producer/Consumer)



Bounded Buffer (Producer/Consumer)

Bounded Buffer (producer-consumer queue)




t1: put()



Bounded Buffer (Producer/Consumer)




Bounded Buffer (Producer/Consumer)



t2: take()

Bounded Buffer (Producer/Consumer)




t2: take()

Bounded Buffer (Producer/Consumer)



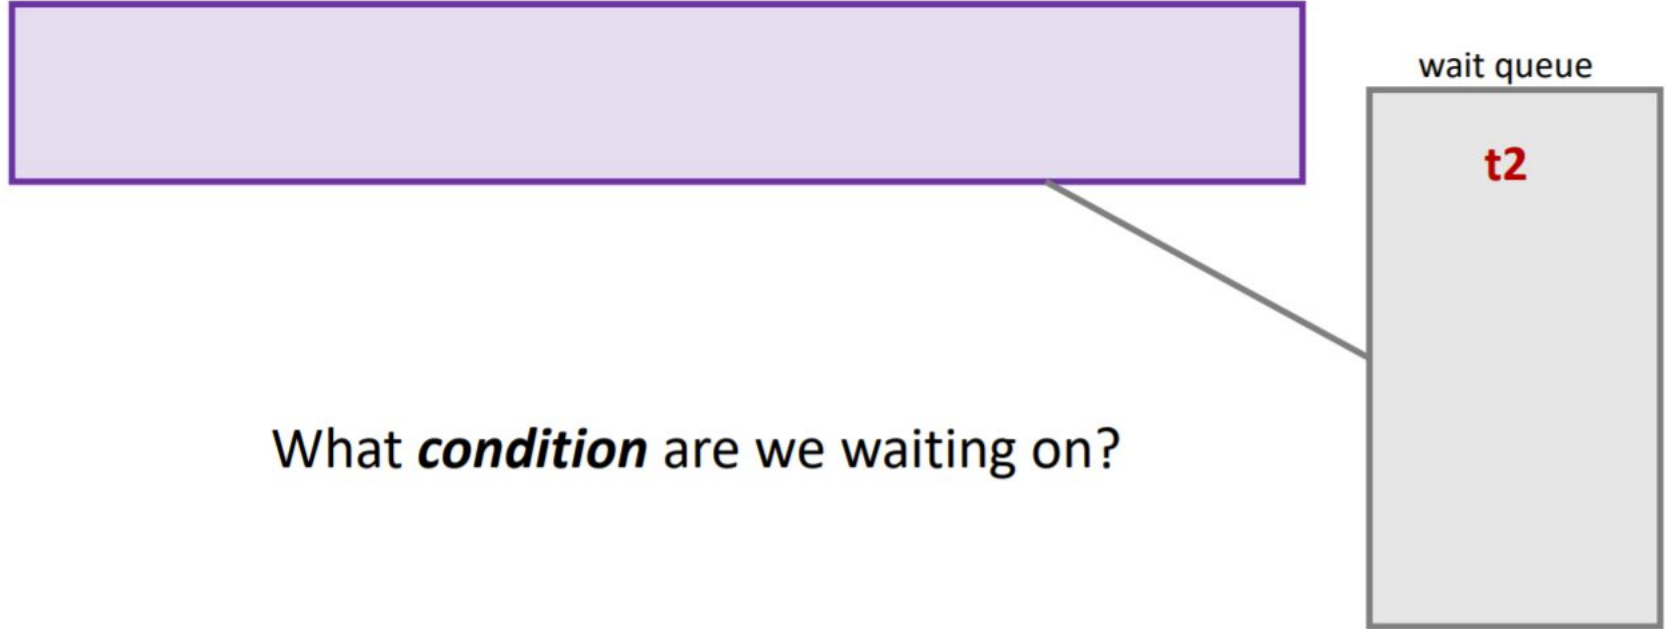
Bounded Buffer (Producer/Consumer)



t2: take()

???

Bounded Buffer (Producer/Consumer)



What *condition* are we waiting on?

Bounded Buffer (Producer/Consumer)



Bounded Buffer (Producer/Consumer)




wait queue

t2


What should happen?

t1: put()


Bounded Buffer (Producer/Consumer)



Bounded Buffer (Producer/Consumer)



Bounded Buffer (Producer/Consumer)




Bounded Buffer (Producer/Consumer)

- The queue is a **circular** queue
- Sometimes also called a **ring buffer**

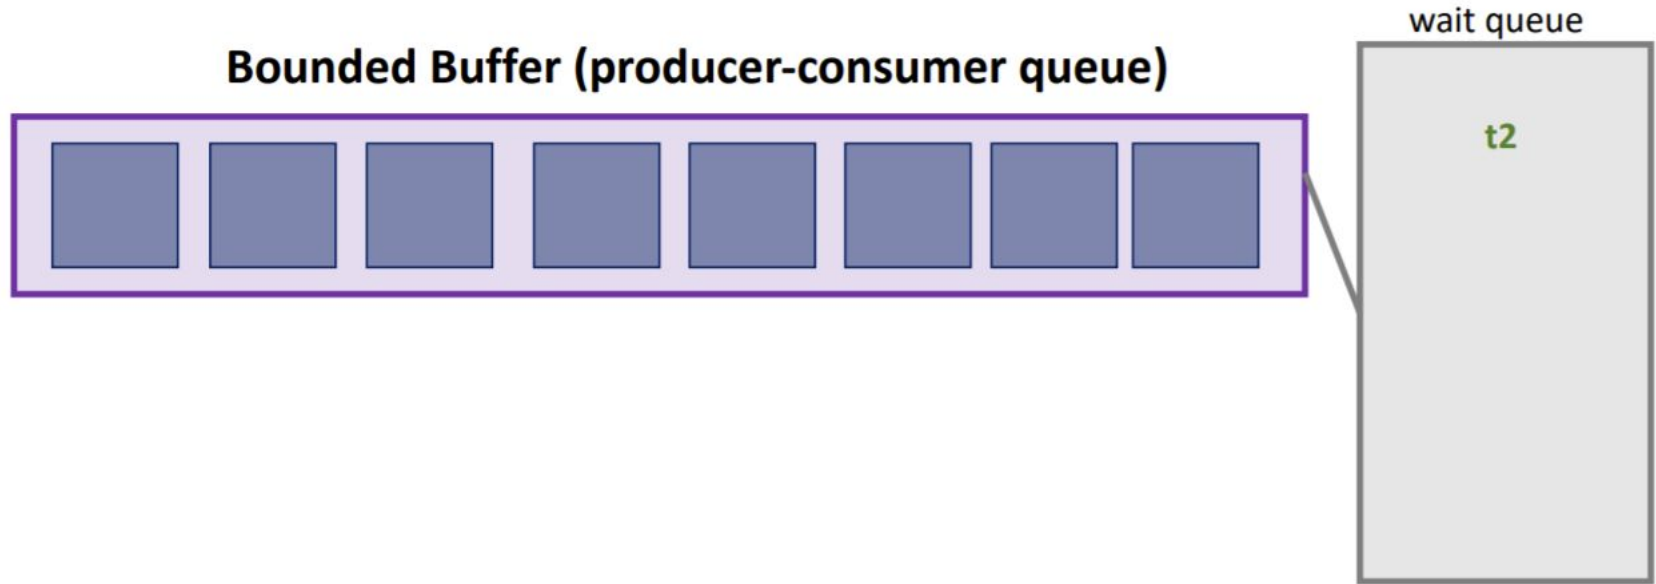


We're full

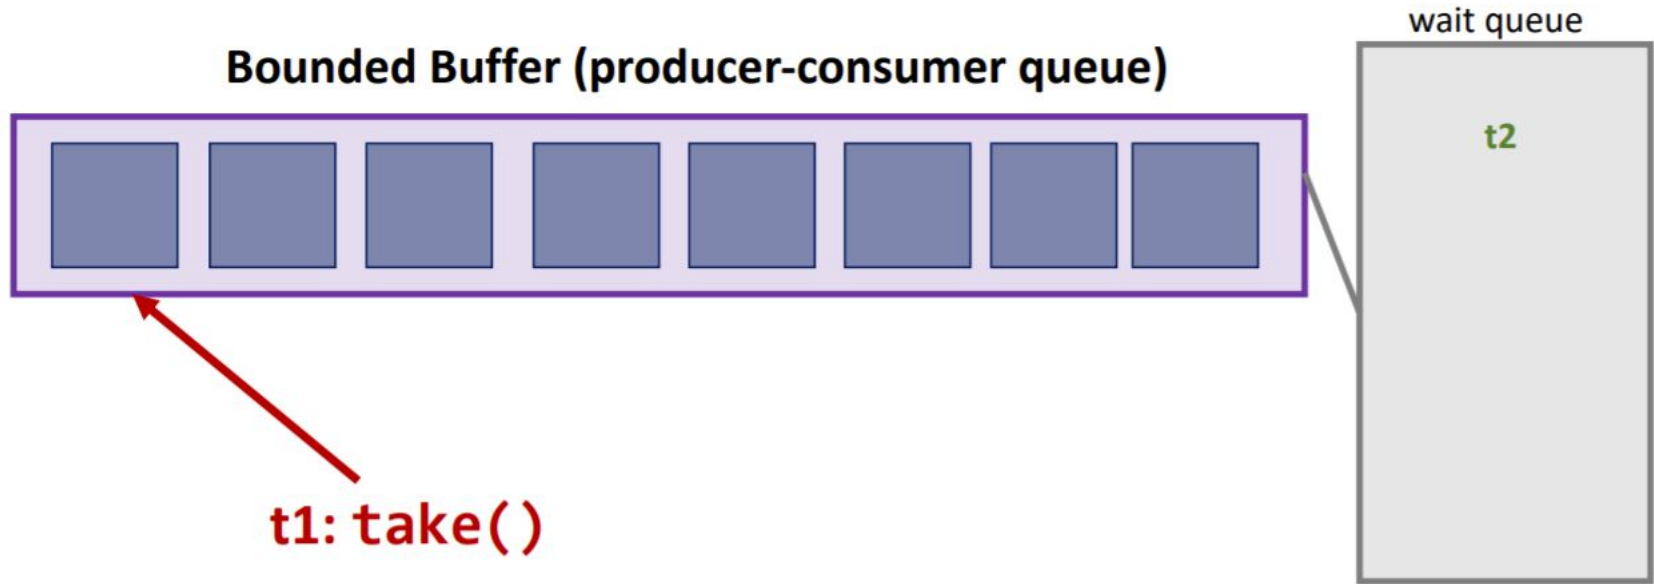
Bounded Buffer (Producer/Consumer)



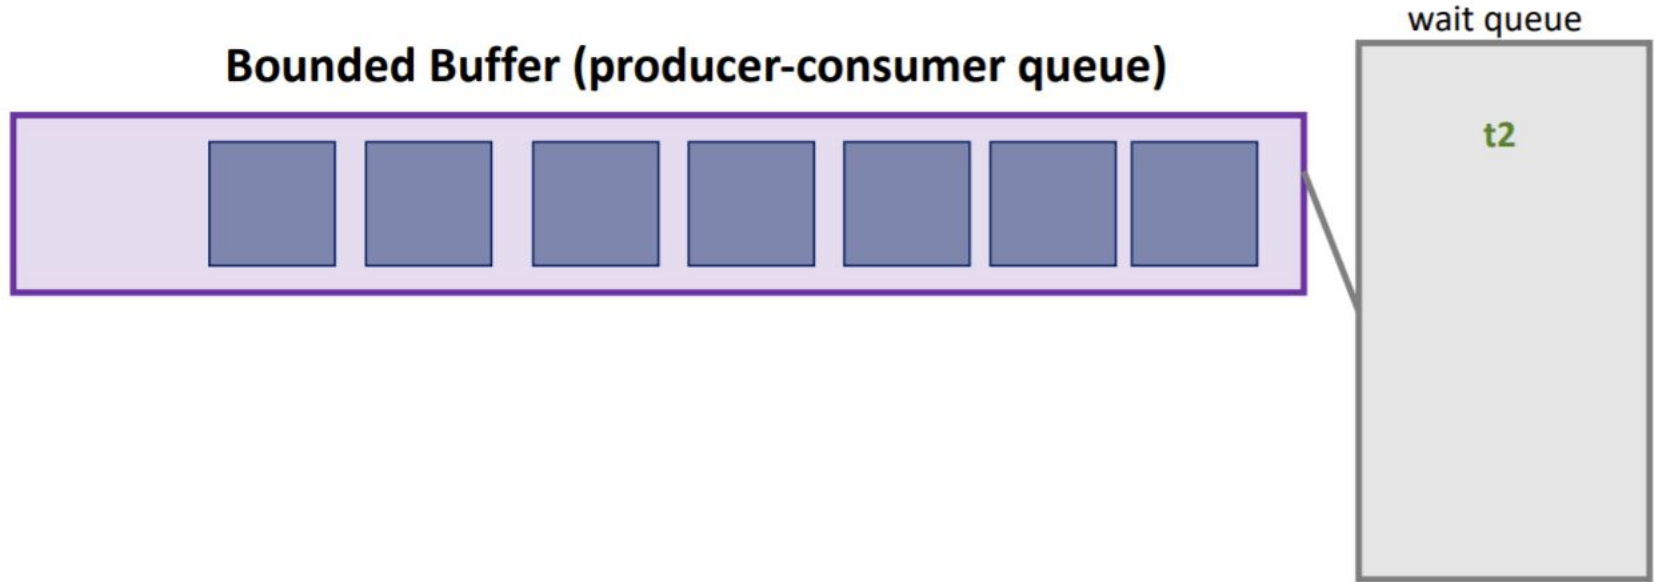
Bounded Buffer (Producer/Consumer)




Bounded Buffer (Producer/Consumer)



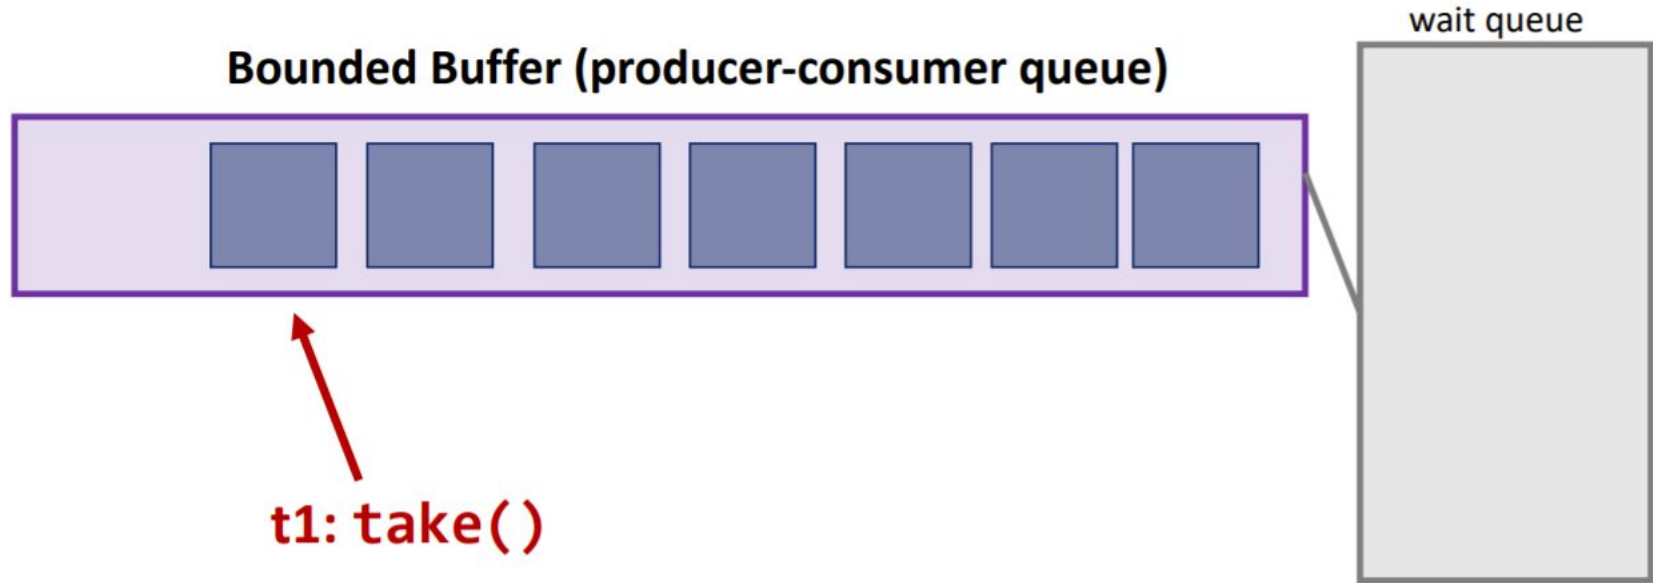
Bounded Buffer (Producer/Consumer)



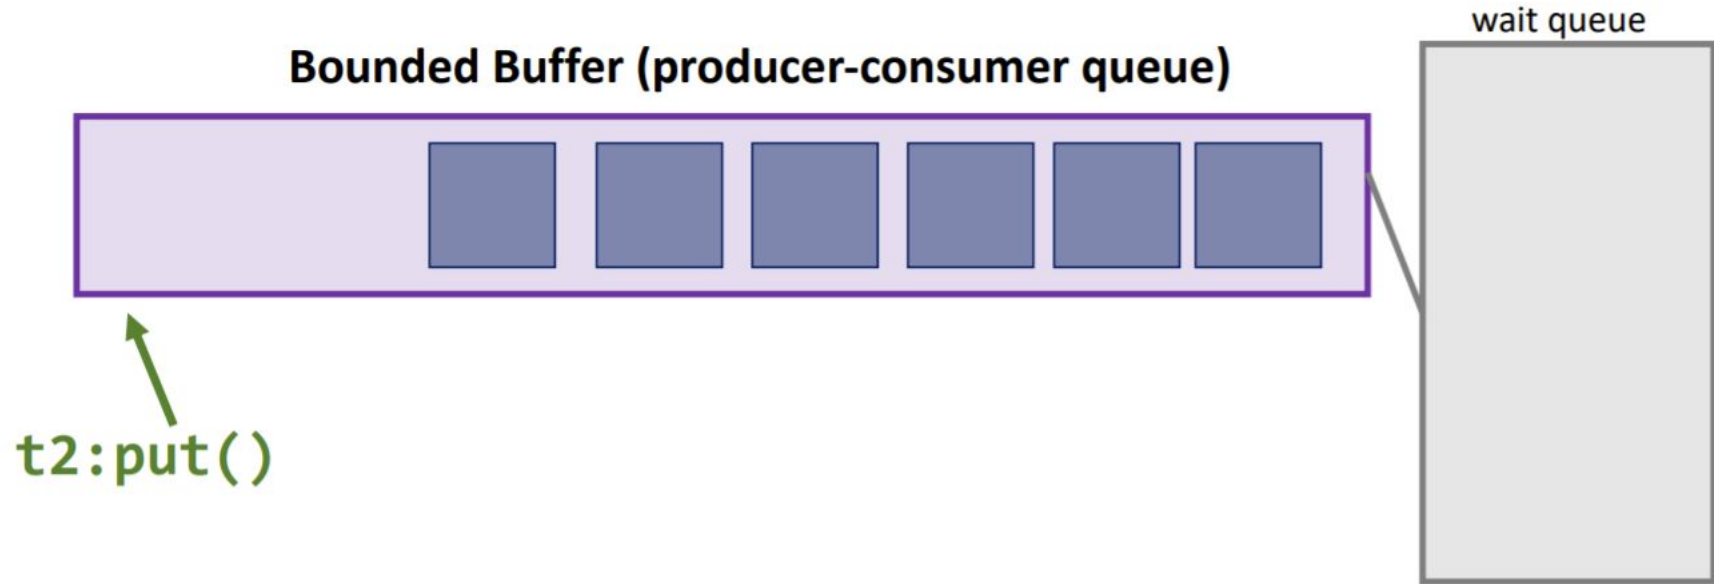
Bounded Buffer (Producer/Consumer)



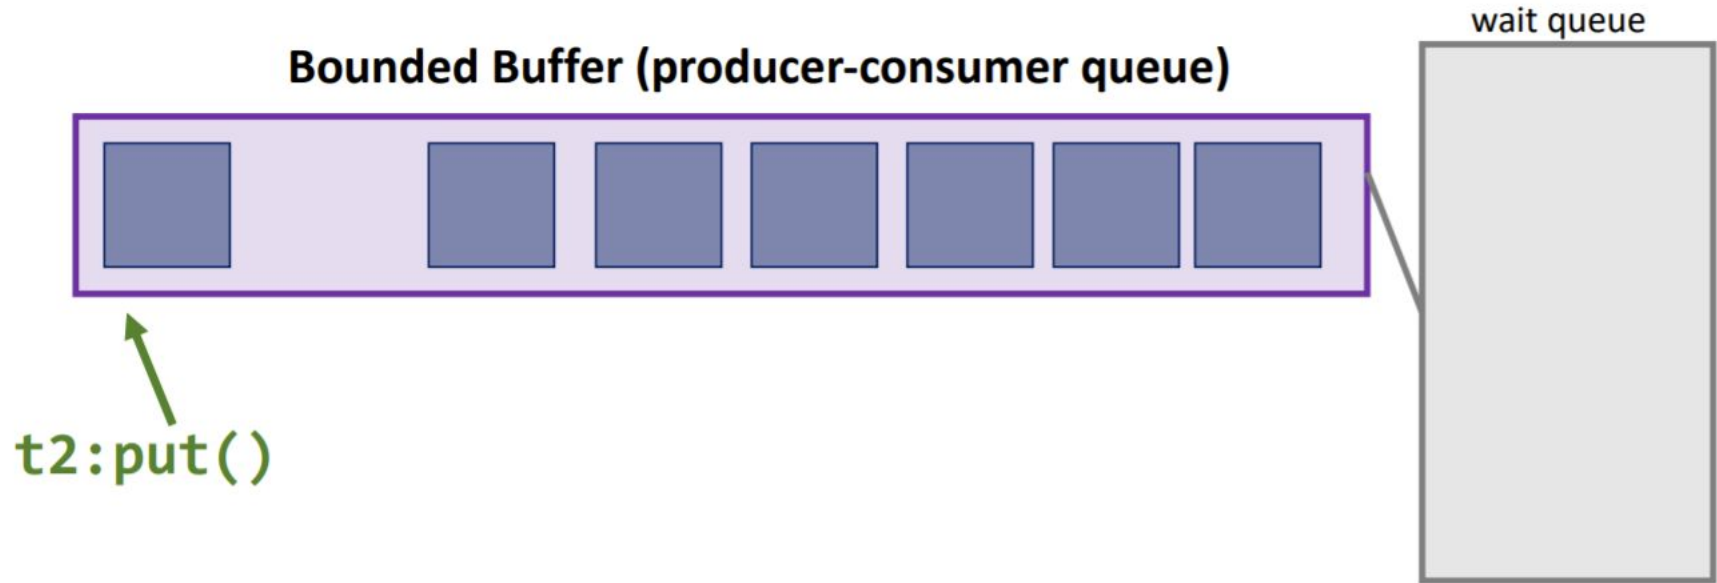
Bounded Buffer (Producer/Consumer)



Bounded Buffer (Producer/Consumer)



Bounded Buffer (Producer/Consumer)



Bounded Buffer (Producer/Consumer)

- Reads and writes to buffer require locking
- when buffers are full
 - writers wait
- when buffers are empty
 - readers wait

Bounded Buffer (Producer/Consumer)

- **Producers** generate data
- **Consumers** take data and process it
- Very frequent situation encountered in systems programming
- General strategy:
 - use CVs to notify:
 - waiting readers when data is available
 - waiting writers when slots are free

Bounded Buffer: Attempt 1

Producer

```
for (int i=0; i<loops; i++) {  
    mutex_lock(&m);  
    while (numfull == MAX)  
        cond_wait(&cond, &m);  
    do_fill(i);  
    cond_signal(&cond);  
    mutex_unlock(&m);  
}
```

Consumer

```
while(1) {  
    mutex_lock(&m);  
    while (numfull == 0)  
        cond_wait(&cond, &m);  
    int tmp = do_get();  
    cond_signal(&cond);  
    mutex_unlock(&m);  
    printf("%d\n", tmp);  
}
```

- Starting simple: assume one producer, one consumer
- **numfull**: number of elements in the buffer
- Does this code work for 1P and 1C?
 - Yes 😊

Bounded Buffer: Attempt 1

Producer

```
for (int i=0; i<loops; i++) {  
    mutex_lock(&m);  
    while (numfull == MAX)  
        cond_wait(&cond, &m); //a  
    do_fill(i);                //b  
    cond_signal(&cond);        //c  
    mutex_unlock(&m);  
}
```

Consumer

```
while(1) {  
    mutex_lock(&m);  
    while (numfull == 0)  
        cond_wait(&cond, &m); //x  
    int tmp = do_get();        //y  
    cond_signal(&cond);        //z  
    mutex_unlock(&m);  
    printf("%d\n", tmp);  
}
```

- How about 1P and 2C?
 - No 😞, **why?**

Bounded Buffer: Attempt 1

Producer

```
for (int i=0; i<loops; i++) {  
    mutex_lock(&m);  
    while (numfull == MAX)  
        cond_wait(&cond, &m); //a  
    do_fill(i);                //b  
    cond_signal(&cond);        //c  
    mutex_unlock(&m);  
}
```

Consumer

```
while(1) {  
    mutex_lock(&m);  
    while (numfull == 0)  
        cond_wait(&cond, &m); //x  
    int tmp = do_get();        //y  
    cond_signal(&cond);        //z  
    mutex_unlock(&m);  
    printf("%d\n", tmp);  
}
```

- Assume C1 and C2 are all waiting at x due to empty queue
- 1) P adds an item to buffer (**line b**), signals cond (**line c**), waking up **C1**, waits on cond until signaled (**line a**)
- 2) C1 is awoken, removes item from buffer (**line y**), signals cond (**line z**), waking up C2, finds buffer empty, goes to sleep (**line x**)
- 3) C2, being woken up by C1, finds buffer empty, goes to sleep waiting on cond (**line x**)

Bounded Buffer: Attempt 1

Producer

```
for (int i=0; i<loops; i++) {  
    mutex_lock(&m);  
    while (numfull == MAX)  
        cond_wait(&cond, &m); //a  
    do_fill(i);                //b  
    cond_signal(&cond);        //c  
    mutex_unlock(&m);  
}
```

Consumer

```
while(1) {  
    mutex_lock(&m);  
    while (numfull == 0)  
        cond_wait(&cond, &m); //x  
    int tmp = do_get();        //y  
    cond_signal(&cond);        //z  
    mutex_unlock(&m);  
    printf("%d\n", tmp);  
}
```

- Everyone is sleeping → P can't produce → no forward progress
- C1's signal was meant to awaken P but it awoke C2

Solution 1: Wake up Everyone

- When not sure if next waiting thread is the right one to wake up, just wake up all
- Not the most elegant solution (that's Solution 2)
 - Probably bad for performance: all awoken threads will compete for mutex again
 - But a good fallback mechanism to ensure correctness
- Need a new API: `cond_broadcast(cv)`
 - Semantic: wakes up all the queues waiting on cv

Solution 2: Use Multiple CVs

- Identify different conditions that need waiting for
- Use a separate CV for each condition using `cond_wait()` and `cond_signal()`
- More elegant, better-performing solution than using `cond_broadcast()`
- Different conditions in bounded buffer problem?
 - Two
 - Waiting for queue to become non-full
 - Waiting for queue to become non-empty

Bounded Buffer: Correct & Elegant Solution

Producer

```
for (int i=0; i<loops; i++) {  
    mutex_lock(&m);  
    while (numfull == MAX)  
        cond_wait(&non_full, &m);  
    do_fill(i);  
    cond_signal(&non_empty);  
    mutex_unlock(&m);  
}
```

Consumer

```
while(1) {  
    mutex_lock(&m);  
    while (numfull == 0)  
        cond_wait(&non_empty, &m);  
    int tmp = do_get();  
    cond_signal(&non_full);  
    mutex_unlock(&m);  
    printf("%d\n", tmp);  
}
```

- Would it be okay also to use two mutexes?
- Why not?
 - Because mutex protects associated with the shared state (buffer, in this case)

Summary: CV Rules of Thumb

- Keep state in addition to CVs
- Always `cond_wait()` or `cond_signal()` with lock held
- Use different CVs for different conditions
- Recheck state assumptions when waking up from waiting

THANK YOU!