

ILLINOIS TECH

College of Computing

CS 450 Operating Systems


Fast File System

Yue Duan


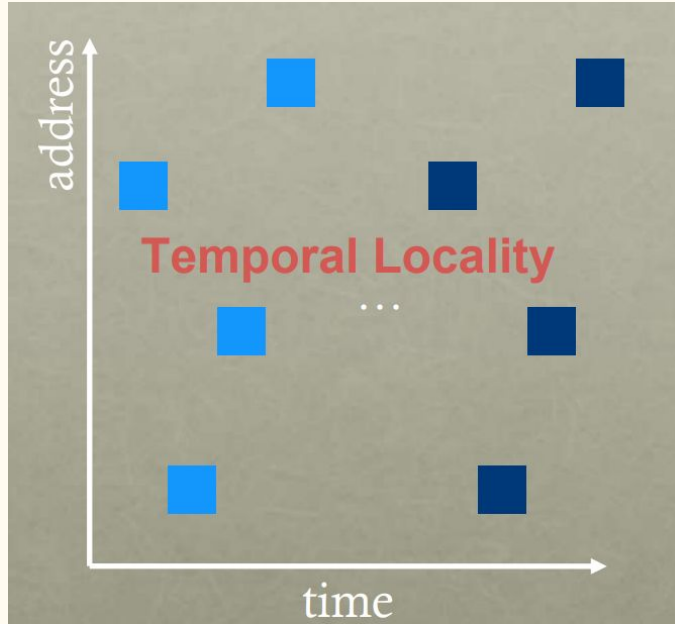
File System Case Study

- Local
 - **FFS: Fast File System**
 - LFS: Log-Structured File System
- Network
 - **NFS: Network File System**
 - AFL: Andrew File System

Recap: Basic Layout





Locality Types




- Which type of locality is most interesting with a disk?

Order Matters




- Implication for disk schedulers?

Policy: Choose Inode, Data Blocks




- Assuming all free, which should be chosen?


Bad File Layout



Better File Layout



Best File Layout



- Can't do this for all files

Fast File System

- System Building
 - 1. identify existing state of the art
 - 2. measure it, identify and understand problems
 - 3. get idea (solutions often flow from deeply understanding problem)
 - 4. build it
- Measure then build!


Measure Old FS

- State of the art: original UNIX file system
- Measure throughput for whole sequential file reads/writes
- Compare to theoretical max, which is the **disk bandwidth**
- Old UNIX file system: achieved only **2%** of potential. Why?



Measurement 1: Aging

- What is performance before/after aging?
 - New FS: 17.5% of disk bandwidth
 - Few weeks old: 3% of disk bandwidth
- Problem: FS becomes fragmented over time
 - free list makes contiguous chunks hard to find



Measurement 2: Block Size

- How does block size affect performance?
 - try doubling it!
- Result: Performance **more** than doubled
- Why double the performance?
 - logically adjacent blocks not physically adjacent
 - only half as many seeks+rotations now required
- Why **more** than double the performance?
 - smaller blocks require more indirect blocks


Old FS Summary

- Free list becomes scrambled \implies random allocations
- Small blocks (512 bytes)
- Blocks laid out poorly
 - long distance between inodes/data
 - related inodes not close to one another
 - inodes in same directory (`ls -l`)
- Result: 2% of potential performance! (and worse over time)
- **Problem:** old FS treats disk like RAM!

Solution: Disk-Awareness


- Primary File System Design Questions:
 - **Where** to place meta-data and data on disk?
 - **How** to use big blocks without wasting space?

Placement Technique 1: Bitmaps




- Use bitmaps instead of free list
- Provides better speed, with more global view
- Faster to find contiguous free blocks

Placement Technique 2: Groups




- How to keep inode close to data?

Placement Technique 2: Groups




- How to keep inode close to data?

Placement Technique 2: Groups




- How to keep inode close to data?

Placement Technique 2: Groups



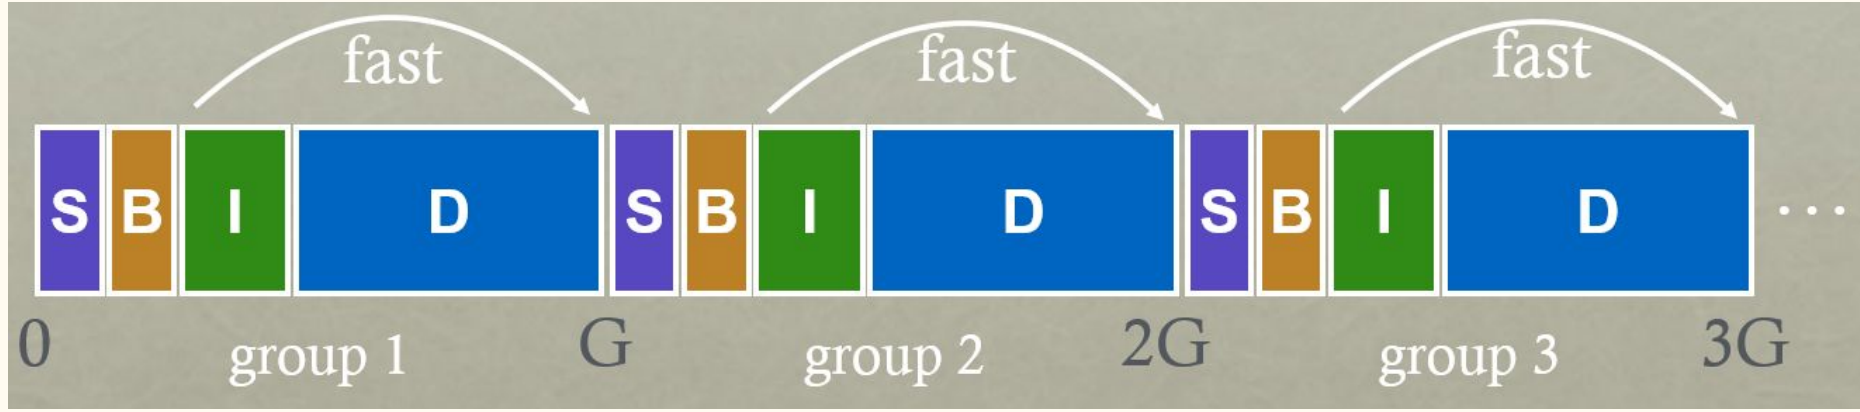
- How to keep inode close to data?

Placement Technique 2: Groups



- How to keep inode close to data?
 - key idea: keep inode close to data
 - use groups across disks
 - try to place inode and data in same group


Placement Technique 2: Groups




- Strategy: allocate inodes and data blocks in same group.

Groups

- In FFS, groups were ranges of cylinders
 - called **cylinder group**
- In ext2-4, groups are ranges of blocks
 - called **block group**




Placement Technique 3: Super Rotation



- Is it useful to have multiple replicated super blocks?
 - Yes, if some (but not all) fail.


Problem




- Old FS: All super-block copies are on the **top platter**
- Correlated failures! What if **top platter** dies?
- Solution:
 - for each group, store super-block at different offset
 - super rotation:
 - use multiple super blocks for each group using different offset

Technique 4: Larger Blocks

- Observation: Doubling block size for old FS over doubled performance
- Why not make blocks huge?
- Most file are very small, even today!
- Large block:
 - lots of waste due to internal fragment in most blocks



Solution: Fragments



- Solution:
 - hybrid – combine best of large blocks and best of small blocks
- Use large block when file is large enough
- Introduce “fragment” for files that use parts of blocks
- Only tail of file uses fragments

file, size 5KB


file, size 2KB

AAAA

B

A

B



append A to first file


file, size 6KB

file, size 2KB

AAAA


B A B A






append A to first file
Not allowed to use fragments across multiple blocks!

What to do instead?



append A to first file,
copy to fragments to new block

Smart Allocation Policy



- Where should new inodes and data blocks go?
 - strategy: put related pieces of data near each other.
- Rules:
 - put directory entries near directory inodes
 - put inodes near directory entries
 - put data blocks near inodes.

Smart Allocation Policy

- Strategy: put related pieces of data near each other.
- Rules:
 - put directory entries near directory inodes
 - put inodes near directory entries
 - put data blocks near inodes.
- Sound good?
- Problem:
 - file system is one big tree
 - all directories and files have a common root
 - all data in same FS is related in some way
- Trying to put everything near everything else

Revised Strategy

- Put more-related pieces of data near each other
- Put less-related pieces of data **far** from each other

Preferences

- File inodes:
 - allocate in **same** group with dir
- Dir inodes:
 - allocate in **new** group with fewer used inodes than average group
- First data block:
 - allocate **near** inode
- Other data blocks:
 - allocate **near** previous block

Problem: Large Files

- Single large file can fill nearly all of a group
- Displaces data for many small files
- Better to do one seek for large file
 - than one seek for each of many small files
- Define “large” as requiring an **indirect block**

```
group inodes      data
  0 /a----- /aaaaaaaaa aaaaaaaaaa aaaaaaaaaa a-----
  1 -----
  2 -----
  ...
```

Revised Strategy

group	inodes	data			
0	/a-----	/aaaaa-----	-----	-----	-----
1	-----	aaaaa-----	-----	-----	-----
2	-----	aaaaa-----	-----	-----	-----
3	-----	aaaaa-----	-----	-----	-----
4	-----	aaaaa-----	-----	-----	-----
5	-----	aaaaa-----	-----	-----	-----
6	-----	-----	-----	-----	-----

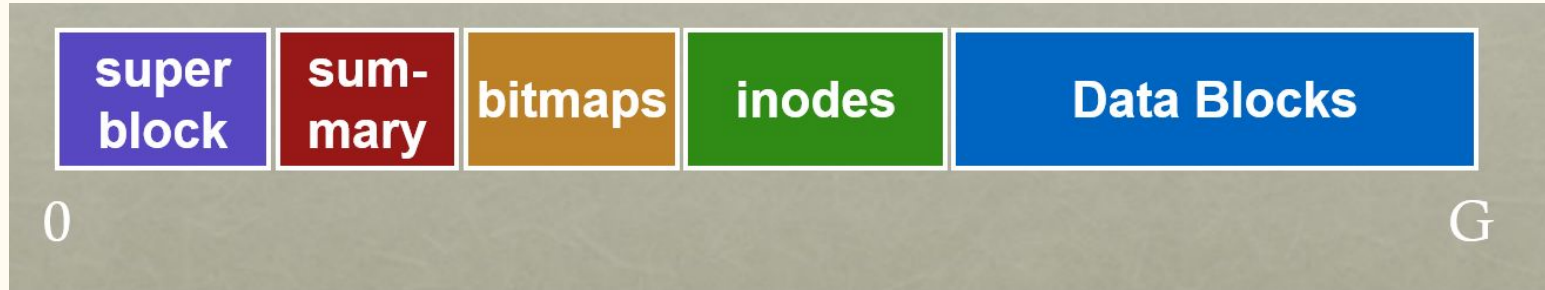
- Large files: where to cut the tree and start growing into another group?
- Starting at indirect (e.g., after 48 KB) put blocks in a new block group.
- Each chunk corresponds to one indirect block
- Block size 4KB, 4 byte per address
 - \Rightarrow 1024 address per indirect block
 - $1024 * 4KB = 4MB$ contiguous “chunk”

Preferences

- File inodes:
 - allocate in **same** group with dir
- Dir inodes:
 - allocate in **new** group with fewer used inodes than average group
- First data block:
 - allocate **near** inode
- Other data blocks:
 - allocate **near** previous block
- Large file data blocks:
 - after 48KB, go to **new** group
 - move to **another** group (w/ fewer than avg blocks) every subsequent 1MB.

Group Descriptor (Summary Block)

- How does file system know which new group to pick?
 - summary block: tracks number of free inodes and data blocks



Conclusion

- FFS inspired modern files systems, including ext2 and ext3
- First disk-aware file system
 - bitmaps
 - locality groups
 - rotated superblocks
 - large blocks
 - fragments
 - smart allocation policy
 - ability to rebuild free lists after crash (fsck)

THANK YOU!