

ILLINOIS TECH

College of Computing

CS 450 Operating Systems More Locks

Yue Duan

Recap

- We went through some software implementation of locks
 - Good concurrency practice
- Ended up with Dekker and Peterson's algorithms
 - Work under assumptions of atomic and in order memory system
 - So, they do not work in practice
 - Compiler reorders
 - memory system is not ordered
- Need hardware support for synchronization
 - Two flavors:
 - Atomic instructions that read and update a variable
 - E.g., test-and-set, xchange, ...
 - Disable interrupts

Using Interrupts

- Turn off interrupts for critical sections
 - Prevent dispatcher from running another thread
 - Code between interrupts executes atomically

```
void acquire(lock_t *l) {  
    disableInterrupts();  
}  
  
void release(lock_t *l) {  
    enableInterrupts();  
}
```

Using Interrupts

- Disadvantages
 - Only works on uniprocessors
 - Process can keep control of CPU for arbitrary length
 - Cannot perform other necessary work

```
void acquire(lock_t *l) {  
    disableInterrupts();  
}  
  
void release(lock_t *l) {  
    enableInterrupts();  
}
```

xchg: atomic exchange, or test-and-set

```
// xchg(int *addr, int newval)  
// return what was pointed to by addr  
// at the same time, store newval into addr
```

```
int xchg(int *addr, int newval) {  
    int old = *addr;  
    *addr = newval;  
    return old;  
}
```

```
static inline unsigned  
xchg(volatile unsigned int *addr, unsigned int newval) {  
    unsigned result;  
    asm volatile("lock; xchgl %0, %1" :  
        "+m" (*addr), "=a" (result) :  
        "1" (newval) : "cc");  
    return result;  
}
```

xchg: atomic exchange, or test-and-set

```
typedef struct __lock_t {  
    int flag;  
} lock_t;  
  
void init(lock_t *lock) {  
    lock->flag = ??;  
}  
  
void acquire(lock_t *lock) {  
    ???  
    // spin-wait (do nothing)  
}  
  
void release(lock_t *lock) {  
    lock->flag = ??;  
}
```

```
int xchg(int *addr, int newval)
```

xchg: atomic exchange, or test-and-set

```
typedef struct __lock_t {  
    int flag;  
} lock_t;  
  
void init(lock_t *lock) {  
    lock->flag = 0;  
}  
  
void acquire(lock_t *lock) {  
    while (xchg(&lock->flag, 1) == 1);  
    // spin-wait (do nothing)  
}  
  
void release(lock_t *lock) {  
    lock->flag = 0;  
}
```

Other atomic HW instructions

```
int CompareAndSwap(int *ptr, int expected, int new) {  
    int actual = *addr;  
    if (actual == expected)  
        *addr = new;  
    return actual;  
}
```

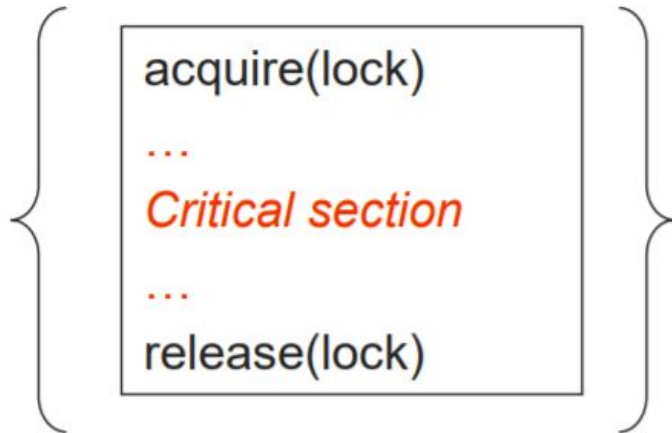
```
void acquire(lock_t *lock) {  
    while(CompareAndSwap(&lock->flag, 0, 1) == 1) ;  
    // spin-wait (do nothing)  
}
```


Summarize Where We Are

- Goal: Use mutual exclusion to protect critical sections of code that access shared resources
- Method: Use locks (spinlocks or disable interrupts)
- Problem: Critical sections can be long

Spinlocks:

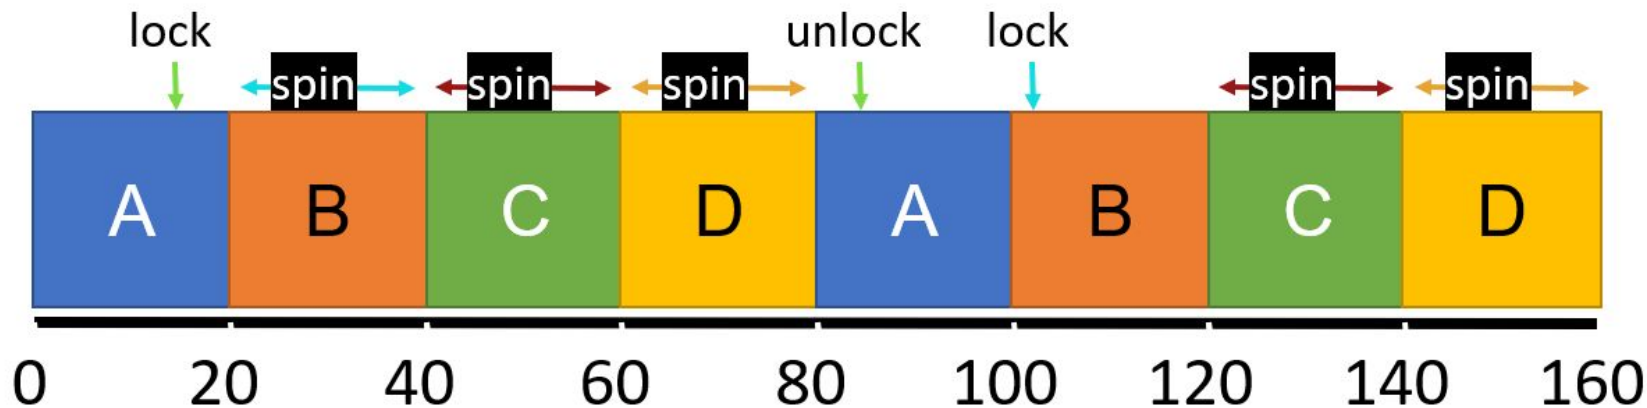
- ▢ Threads waiting to acquire lock spin in test-and-set loop
- ▢ Wastes CPU cycles
- ▢ Longer the CS, the longer the spin
- ▢ Greater the chance for lock holder to be interrupted
- ▢ Memory consistency model causes problems (out of scope of this class)



Disabling Interrupts:


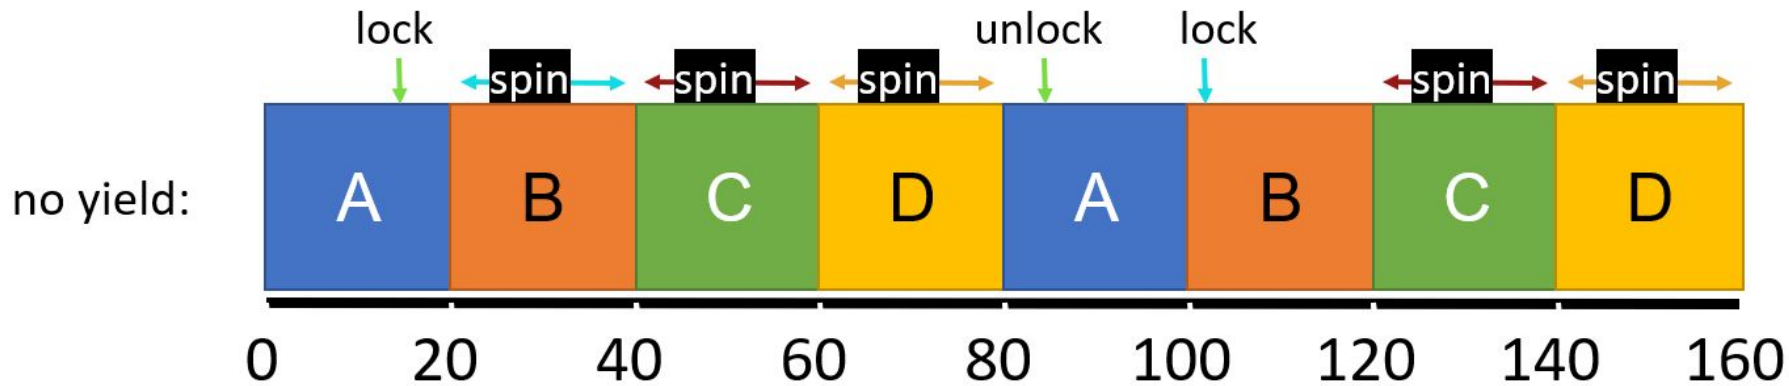
- ▢ Should not disable interrupts for long periods of time
- ▢ Can miss or delay important events (e.g., timer, I/O)

CPU Scheduler is Ignorant



- CPU scheduler may run B instead of A even though **B is waiting for A**

yield() Instead of spin



Spinlock Performance

- Waste...
 - Without yield: $O(\text{threads} * \text{time_slice})$
 - With yield: $O(\text{threads} * \text{context_switch})$
 - So even with yield, spinning is slow with high thread contention
- Next improvement:
 - Block and put thread on waiting queue instead of spinning

Lock implementation: block when waiting

- Lock implementation removes waiting threads from the ready queue
 - e.g. with `park()`: removes current thread from run queue
 - and `unpark(threadID)`: returns thread tid to run queue
- Scheduler only runs ready threads
- Separates concerns
 - programmer doesn't need to deal with yield vs not/spinning issue
- **Quiz**: where should locks be implemented?

Example

Ready: {A, B, C, D}

Waiting: {}

Running: {}




Example

Ready: {B, C, D}

Waiting: {}

Running: {A}




Example

Ready: {A, C, D}

Waiting: {}

Running: {B}




Example

Ready: {A, C, D}

Waiting: {B}

Running: {}




Example

Ready: {A, D}

Waiting: {B}

Running: {C}




Example

Ready: {A, C}

Waiting: {B}

Running: {D}




Example

Ready: {A, C}

Waiting: {B, D}

Running: {}




Example

Ready: {C}

Waiting: {B, D}

Running: {A}




Example

Ready: {A}

Waiting: {B, D}

Running: {C}




Example

Ready: {C}

Waiting: {B, D}

Running: {A}

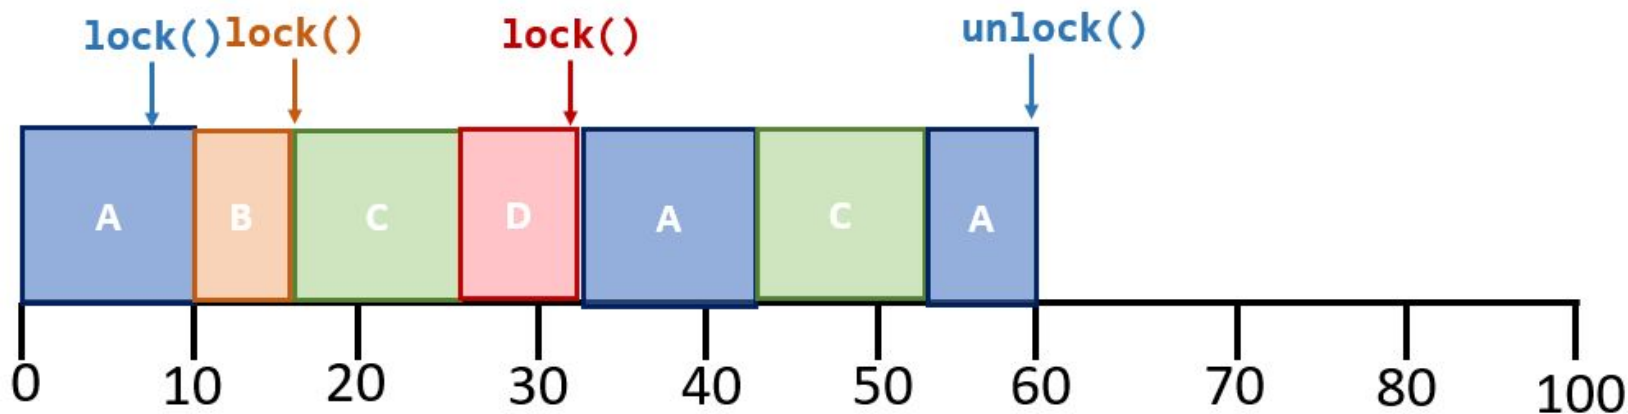


Example

Ready: {C, B, D}

Waiting: {}

Running: {A}




Example

Ready: {C, B, D}

Waiting: {}

Running: {A}




Example

Ready: {B, D, A}

Waiting: {}

Running: {C}




Example

Ready: {D, A, C}

Waiting: {}

Running: {B}



Benefits

- Threads aren't contending on the lock when they shouldn't be (mostly)
- **Fewer** context switches

Lock: block when waiting

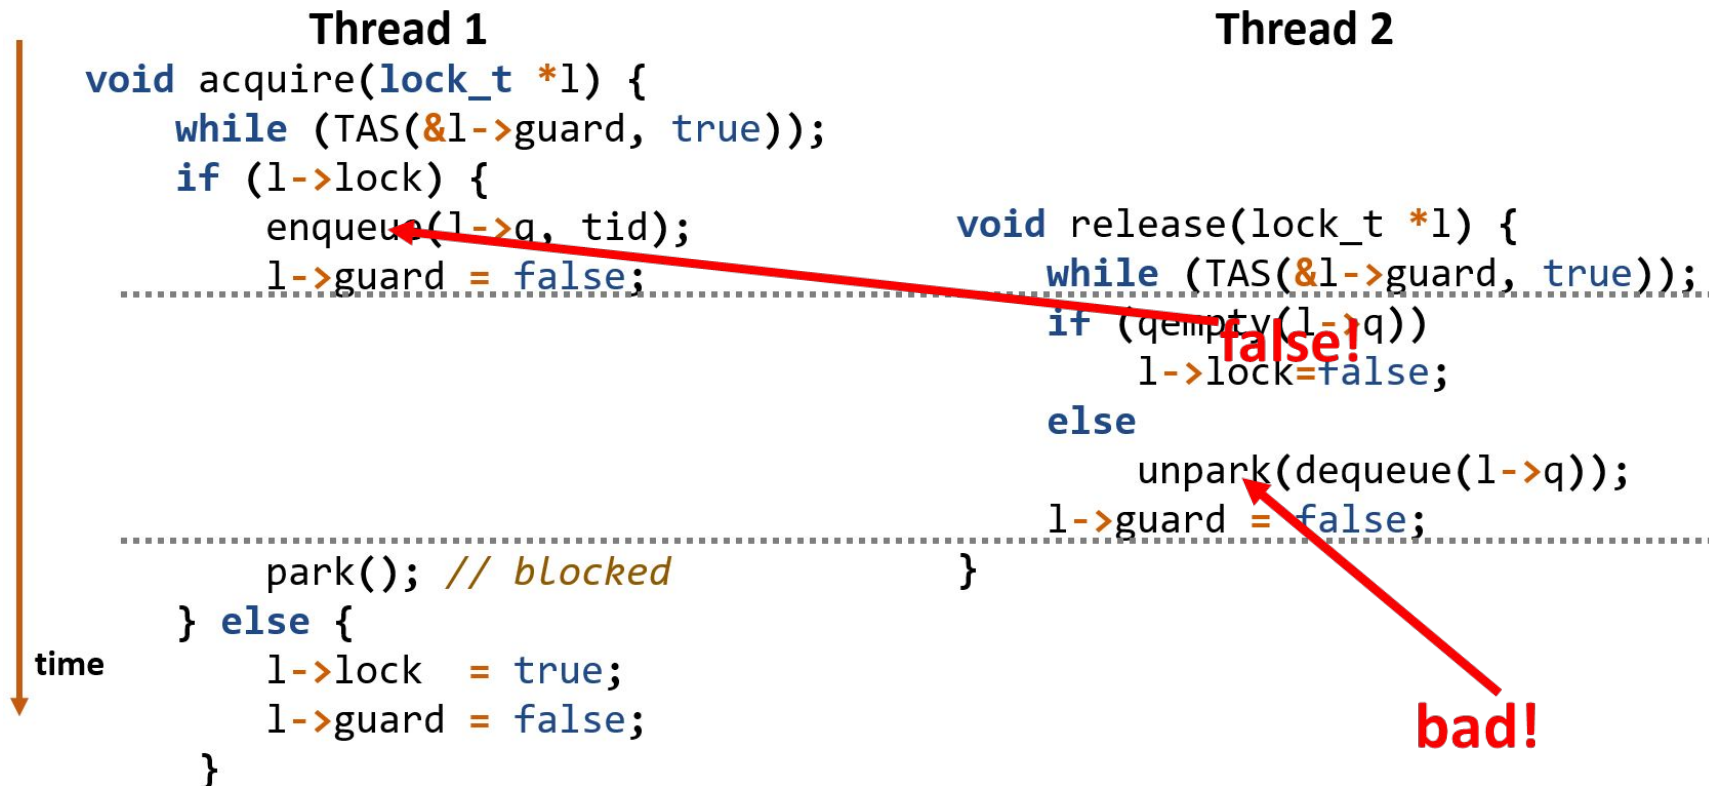
```
typedef struct {  
    bool lock = false;  
    bool guard = false;  
    queue_t q;  
} lock_t;
```

- Why do we need guard?
 - to keep queue ops from race
- Why OK to spin on guard?
 - waiting time is small
- In release(), why not set lock=false when unpark()?
- Is there a race condition?

```
void acquire(lock_t *l) {  
    while (TAS(&l->guard, true));  
    if (l->lock) {  
        enqueue(l->q, tid);  
        l->guard = false;  
        park(); // blocked  
    } else {  
        l->lock = true;  
        l->guard = false;  
    }  
}
```

```
void release(lock_t *l) {  
    while (TAS(&l->guard, true));  
    if (qempty(l->q))  
        l->lock=false;  
    else  
        unpark(dequeue(l->q));  
    l->guard = false;  
}
```

Race Condition



Lock: the fix

```
typedef struct {  
    bool lock = false;  
    bool guard = false;  
    queue_t q;  
} lock_t;
```

- setpark(): informs the OS of my plan to park() myself
- If there is an unpark() between my setpark() and park(), park() will return immediately (no blocking)
- **Why does this fix the race?**

```
void acquire(lock_t *l) {  
    while (TAS(&l->guard, true));  
    if (l->lock) {  
        enqueue(l->q, tid);  
        setpark(); // notify of plan  
        l->guard = false;  
        park(); // unless unpark()  
    } else {  
        l->lock = true;  
        l->guard = false;  
    }  
}  
void release(lock_t *l) {  
    while (TAS(&l->guard, true));  
    if (qempty(l->q))  
        l->lock=false;  
    else  
        unpark(dequeue(l->q));  
    l->guard = false;  
}
```

Spinning vs. Blocking

- Each approach is better under different circumstances
- **Uniprocessor**
 - Waiting process is scheduled → Process holding lock isn't
 - Therefore, waiting process should always relinquish processor
 - Associate queue of waiters with each lock (as in previous implementation)
- **Multiprocessor**
 - Waiting process is scheduled → Process holding lock **might be**
 - Spin or block depends on **how long before lock is released**
 - Lock is going to be released quickly → Spin-wait
 - Lock released slowly → Block

Oracle

- Suppose we know how long **(t) a lock will be held**
- And suppose we make our decision to block or spin based on **C, the cost of a context switch**

$$action = \begin{cases} t \leq C \rightarrow spin \\ t > C \rightarrow block \end{cases}$$

But we need to know the future!

Two-Phase Locking

- A hybrid approach that combines best of spinning and blocking
- Phase 1:
 - spin for a short time, hoping the lock becomes available soon
- Phase 2:
 - if lock not released after a short while, then block
- Question: **how long to spin for?**
 - There's a nice theory which is in practice hard to implement, **so just spin for a few iterations**

Concurrency Goals

- Mutual Exclusion
 - Keep two threads from executing in a critical section concurrently
 - We solved this with **locks**
- Dependent Events
 - We want a thread to wait until some particular event has occurred
 - Or some condition has been met
 - Solved with **condition variables and semaphores**

Example: join()

```
pthread_t p1, p2;  
pthread_create(&p1, NULL, mythread, "A");  
pthread_create(&p2, NULL, mythread, "B");  
// join waits for the threads to finish  
pthread_join(p1, NULL);  
pthread_join(p2, NULL);  
printf("Main: done\n [balance: %d]\n", balance);  
return 0;
```

join(): parent must wait for child thread to finish

Condition Variables

- CV:
 - queue of waiting threads
- **B** waits for a signal on CV before running
 - `wait(CV, ...);`
- **A** sends `signal()` on CV when time for **B** to run
 - `signal(CV, ...);`

API

- **wait**(cond_t * cv, mutex_t * lock)
 - assumes lock is held when wait() is called (why?)
 - puts caller to sleep + releases the lock (atomically)
 - when awoken, reacquires lock before returning
- **signal**(cond_t * cv)
 - wake a single waiting thread (if ≥ 1 thread is waiting)
 - if there is no waiting thread, NOP

THANK YOU!