

# ILLINOIS TECH

College of Computing

## CS 450 Operating Systems

### Mutual $\bar{E}$ xclusion

Yue Duan

# Recap

- Want 3 instructions to execute as an uninterruptible group
- That is, we want them to be an **atomic** unit

```
mov 0x123, %eax  
add %0x1, %eax  
mov %eax, 0x123
```

*critical section*

- More general:
  - Need mutual exclusion for critical sections
  - if process A is in critical section C, process B can't be (okay if other processes do unrelated work)

# Another Example

- Consider multi-threaded programs that do more than increment a shared balance
- E.g., multi-threaded program with a shared linked-list
  - All concurrent operations:
    - Thread **A** inserts element **a**
    - Thread **B** inserts element **b**
    - Thread **C** looks up element **c**

# Shared Linked List

```
void list_insert(list_t *L, int key) {
    node_t *new = malloc(sizeof(node_t));
    assert(new);
    new->key = key;
    new->next = L->head;
    L->head = new;
}

int list_lookup(list_t *L, int key) {
    node_t *tmp = L->head;
    while (tmp) {
        if (tmp->key == key)
            return 1;
        tmp = tmp->next;
    }
    return 0;
}
```

```
typedef struct __node_t {
    int key;
    struct __node_t *next;
} node_t;

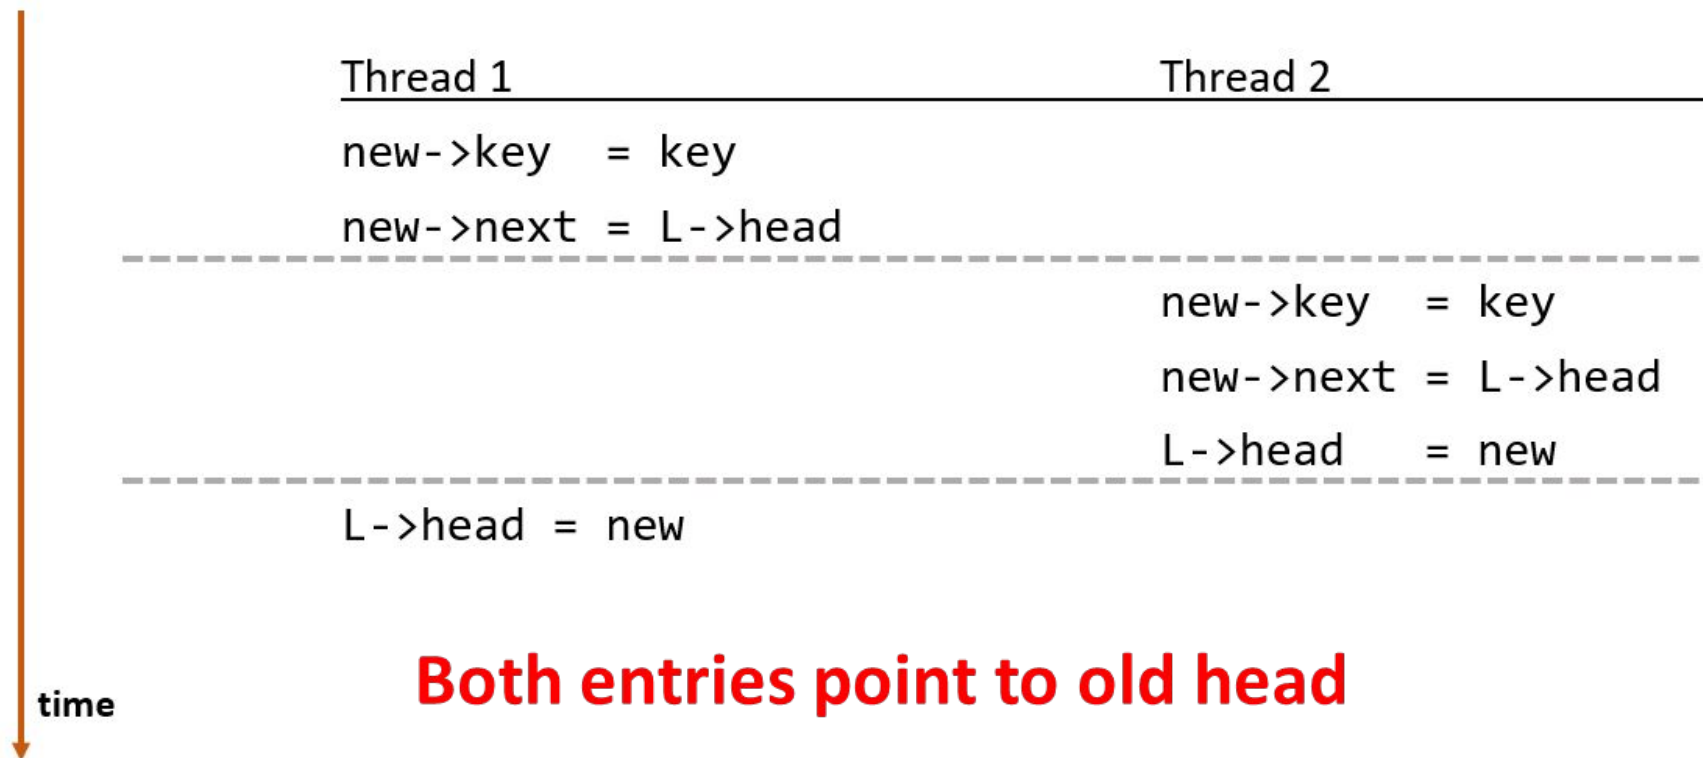
typedef struct __list_t {
    node_t *head;
} list_t;

void list_init(list_t *L) {
    L->head = NULL;
}
```

What can go wrong?

What schedule leads to a problem?


# Shared Linked List



**Both entries point to old head**

Only one entry (which one?) can be the new head.

# Shared Linked List



# Shared Linked List

```
void list_insert(list_t *L, int key) {
    node_t *new = malloc(sizeof(node_t));
    assert(new);
    new->key = key;
    new->next = L->head;
    L->head = new;
}

int list_lookup(list_t *L, int key) {
    node_t *tmp = L->head;
    while (tmp) {
        if (tmp->key == key)
            return 1;
        tmp = tmp->next;
    }
    return 0;
}
```

```
typedef struct __node_t {
    int key;
    struct __node_t *next;
} node_t;

typedef struct __list_t {
    node_t *head;
} list_t;

void list_init(list_t *L) {
    L->head = NULL;
}
```

How do we add locks to this?

# Concurrent Linked List

```
void list_insert(list_t *L, int key) {  
    node_t *new = malloc(sizeof(node_t));  
    assert(new);  
    new->key = key;  
    new->next = L->head;  
    L->head = new;  
}
```

```
int list_lookup(list_t *L, int key) {  
    node_t *tmp = L->head;  
    while (tmp) {  
        if (tmp->key == key)  
            return 1;  
        tmp = tmp->next;  
    }  
    return 0;  
}
```

```
typedef struct __node_t {  
    int key;  
    struct __node_t *next;  
} node_t;
```

```
typedef struct __list_t {  
    pthread_mutex_t lock;  
    node_t *head;  
} list_t;
```


```
void list_init(list_t *L) {  
    L->head = NULL;  
    pthread_mutex_init(&L->lock, NULL);  
}
```

**pthread\_mutex\_t lock;**  
**One lock per list**





# Locking Linked Lists : Approach #1

- Consider everything critical section

```
pthread_mutex_lock(&L->lock);  Void list_insert(list_t *L, int key) {  
    node_t *new =  
        malloc(sizeof(node_t));  
    assert(new);  
    new->key = key;  
    new->next = L->head;  
    L->head = new;  
    pthread_mutex_unlock(&L->lock);  
}
```

- Can critical section be smaller?

```
pthread_mutex_lock(&L->lock);  int list_lookup(list_t *L, int key) {  
    node_t *tmp = L->head;  
    while (tmp) {  
        if (tmp->key == key)  
            return 1;  
        tmp = tmp->next;  
    }  
    pthread_mutex_unlock(&L->lock);  return 0;  
}
```

# Locking Linked Lists : Approach #2

- Critical section as small as possible

```
Void list_insert(list_t *L, int key) {
    node_t *new =
        malloc(sizeof(node_t));
    assert(new);
    new->key = key;
    new->next = L->head;
    L->head = new;
}

int list_lookup(list_t *L, int key) {
    node_t *tmp = L->head;
    while (tmp) {
        if (tmp->key == key)
            return 1;
        tmp = tmp->next;
    }
    return 0;
}
```

`pthread_mutex_lock(&L->lock);` →

`pthread_mutex_unlock(&L->lock);` →

`pthread_mutex_lock(&L->lock);` →

`pthread_mutex_unlock(&L->lock);` →

# Locking Linked Lists : Approach #3

- What about `list_lookup()`?

`pthread_mutex_lock(&L->lock);`



```
Void list_insert(list_t *L, int key) {  
    node_t *new =  
        malloc(sizeof(node_t));  
    assert(new);  
    new->key = key;  
    new->next = L->head;  
    L->head = new;  
}
```

`pthread_mutex_unlock(&L->lock);`



`pthread_mutex_lock(&L->lock);`



```
int list_lookup(list_t *L, int key) {  
    node_t *tmp = L->head;  
    while (tmp) {  
        if (tmp->key == key)  
            return 1;  
        tmp = tmp->next;  
    }  
}
```

- If no `list_delete()`, locks not necessary

`pthread_mutex_unlock(&L->lock);`



```
}  
return 0;  
}
```

# Critical Section Requirements

- Mutual exclusion (mutex)
  - If one thread is in the critical section, then no other is
- Progress
  - A thread in the critical section will eventually leave the critical section
  - If some thread T is not in the critical section, then T cannot prevent other thread from entering the critical section
- Bounded waiting (no starvation)
  - If some thread T is waiting on the critical section, then T will eventually enter the critical section
- Performance
  - The overhead of entering and exiting the critical section is small with respect to the work being done within it

# Building Critical Sections

- To implement, need **atomic operations**
- Atomic operation:
  - guarantees no other instructions can be interleaved
- Examples of atomic operations
  - **Code between interrupts on uniprocessors**
    - Disable timer interrupts, don't do any I/O
  - **Loads and stores of words**
    - Load r1, B
    - Store r1, A
  - **Special hardware instructions**
    - **atomic** test & set
    - **atomic** compare & swap

# Building Critical Sections

- **Locks**
  - Primitive, minimal semantics, used to build others
- **Semaphores**
  - Basic, easy to get the hang of, but hard to program with
- **Monitors**
  - High-level, requires language support, operations implicit
- **Architecture help**
  - Atomic read/write » Can it be done?

# How to implement a lock? First try

```
pthread_trylock(mutex) {  
    if (mutex==0) {  
        mutex= 1;  
        return 1;  
    } else return 0;  
}
```

Thread 0, 1, ...

```
...//time to access critical region  
while(!pthread_trylock(mutex); // wait  
<critical region>  
pthread_unlock(mutex)
```

- Does this work?
  - Assume reads/writes are atomic
- The lock itself is a critical region!
  - Chicken and egg
- Computer scientist struggled with how to create software locks

# How to implement a lock? Second try

```
int turn = 1;
```

```
while (true) {  
    while (turn != 1) ;  
    critical section  
    turn = 2;  
    outside of critical section  
}
```

```
while (true) {  
    while (turn != 2) ;  
    critical section  
    turn = 1;  
    outside of critical section  
}
```

- This is called **alternation**
- It satisfies mutex:
  - If blue is in the critical section, then  $\text{turn} == 1$  and if yellow is in the critical section then  $\text{turn} == 2$
  - $(\text{turn} == 1) \equiv (\text{turn} != 2)$
- Is there anything wrong with this solution?



# How to implement a lock? Second try

```
int turn = 1;
```

P1

```
while (true) {  
    while (turn != 1) ;  
    critical section  
    turn = 2;  
    outside of critical section  
}
```

P2

```
while (true) {  
    while (turn != 2) ;  
    critical section  
    turn = 1;  
    outside of critical section  
}
```

- Is there anything wrong with this solution?
  - break the progress requirement
  - **how?**
  - strictly alternating order; may not map well to application needs
  - consider the situation where P2 wants to enter its CS earlier than P1

## Third try - Two Variables

Bool flag[2]

```
while (flag[1] != 0);  
flag[0] = 1;  
critical section  
flag[0]=0;  
outside of critical section
```

```
while (flag[0] != 0);  
flag[1] = 1;  
critical section  
flag[1]=0;  
outside of critical section
```

- We added two variables to try to break the race for the same variable
- **Is there anything wrong with this solution?**
  - break the mutual exclusion requirement
  - **how?**

## Third try - Two Variables

Bool flag[2]

```
while (flag[1] != 0);  
flag[0] = 1;  
critical section  
flag[0]=0;  
outside of critical section
```

```
while (flag[0] != 0);  
flag[1] = 1;  
critical section  
flag[1]=0;  
outside of critical section
```

- Is there anything wrong with this solution?
  - break the mutual exclusion requirement

## Fourth try – set before you check

Bool flag[2]

```
flag[0] = 1;  
while (flag[1] != 0);  
critical section  
flag[0]=0;  
outside of critical section
```

```
flag[1] = 1;  
while (flag[0] != 0);  
critical section  
flag[1]=0;  
outside of critical section
```

- Is there anything wrong with this solution?
  - **deadlock**
  - both threads could set their flag as true simultaneously and both will wait infinitely later on

# Fifth try – double check and back off

Bool flag[2]

```
flag[0] = 1;
while (flag[1] != 0) {
    flag[0] = 0;
    wait a short time;
    flag[0] = 1;
}
critical section
flag[0]=0;
outside of critical section
```

```
flag[1] = 1;
while (flag[0] != 0) {
    flag[1] = 0;
    wait a short time;
    flag[1] = 1;
}
critical section
flag[1]=0;
outside of critical section
```

- Is this a correct solution?
  - indefinite postponement

# Six try – Dekker's Algorithm

```
Bool flag[2];  
Int turn = 1;
```

```
flag[0] = 1;  
while (flag[1] != 0) {  
    if(turn == 2) {  
        flag[0] = 0;  
        while (turn == 2);  
        flag[0] = 1;  
    } //if  
}  
  
//while  
critical section  
flag[0]=0;  
turn=2;  
outside of critical section
```

```
flag[1] = 1;  
while (flag[0] != 0) {  
    if(turn == 1) {  
        flag[1] = 0;  
        while (turn == 1);  
        flag[1] = 1;  
    } //if  
}  
  
//while  
critical section  
flag[1]=0;  
turn=1;  
outside of critical section
```

# Another solution: Peterson's Algorithm

```
int turn = 1;  
bool try1 = false, try2 = false;
```

```
while (true) {  
    try1 = true;  
    turn = 2;  
    while (try2 && turn != 1) ;  
    critical section  
    try1 = false;  
    outside of critical section  
}
```

```
while (true) {  
    try2 = true;  
    turn = 1;  
    while (try1 && turn != 2) ;  
    critical section  
    try2 = false;  
    outside of critical section  
}
```

# Some observations

- This stuff (software locks) is hard
  - Hard to get right
  - Hard to prove right
- It also is inefficient
  - A spin lock – waiting by checking the condition repeatedly
- Even better, software locks don't really work
  - Compiler and hardware reorder memory references from different threads
  - So, we need to find a different way
- Hardware help



# Using Interrupts

- Turn off interrupts for critical sections
  - Prevent dispatcher from running another thread
  - Code between interrupts executes atomically

```
void acquire(lock_t *l) {  
    disableInterrupts();  
}  
  
void release(lock_t *l) {  
    enableInterrupts();  
}
```

# Using Interrupts

- Disadvantages
  - Only works on uniprocessors
  - Process can keep control of CPU for arbitrary length
  - Cannot perform other necessary work

```
void acquire(lock_t *l) {  
    disableInterrupts();  
}  
  
void release(lock_t *l) {  
    enableInterrupts();  
}
```

# xchg: atomic exchange, or test-and-set

```
// xchg(int *addr, int newval)  
// return what was pointed to by addr  
// at the same time, store newval into addr
```

```
int xchg(int *addr, int newval) {  
    int old = *addr;  
    *addr = newval;  
    return old;  
}
```

```
static inline unsigned  
xchg(volatile unsigned int *addr, unsigned int newval) {  
    unsigned result;  
    asm volatile("lock; xchgl %0, %1" :  
        "+m" (*addr), "=a" (result) :  
        "1" (newval) : "cc");  
    return result;  
}
```

# xchg: atomic exchange, or test-and-set

```
typedef struct __lock_t {  
    int flag;  
} lock_t;  
  
void init(lock_t *lock) {  
    lock->flag = ??;  
}  
  
void acquire(lock_t *lock) {  
    ???  
    // spin-wait (do nothing)  
}  
  
void release(lock_t *lock) {  
    lock->flag = ??;  
}
```

```
int xchg(int *addr, int newval)
```

# xchg: atomic exchange, or test-and-set

```
typedef struct __lock_t {  
    int flag;  
} lock_t;  
  
void init(lock_t *lock) {  
    lock->flag = 0;  
}  
  
void acquire(lock_t *lock) {  
    while (xchg(&lock->flag, 1) == 1);  
    // spin-wait (do nothing)  
}  
  
void release(lock_t *lock) {  
    lock->flag = 0;  
}
```

## Other atomic HW instructions

```
int CompareAndSwap(int *ptr, int expected, int new) {  
    int actual = *addr;  
    if (actual == expected)  
        *addr = new;  
    return actual;  
}
```

```
void acquire(lock_t *lock) {  
    while(CompareAndSwap(&lock->flag, 0, 1) == 1) ;  
    // spin-wait (do nothing)  
}
```

**THANK YOU!**