



ILLINOIS TECH

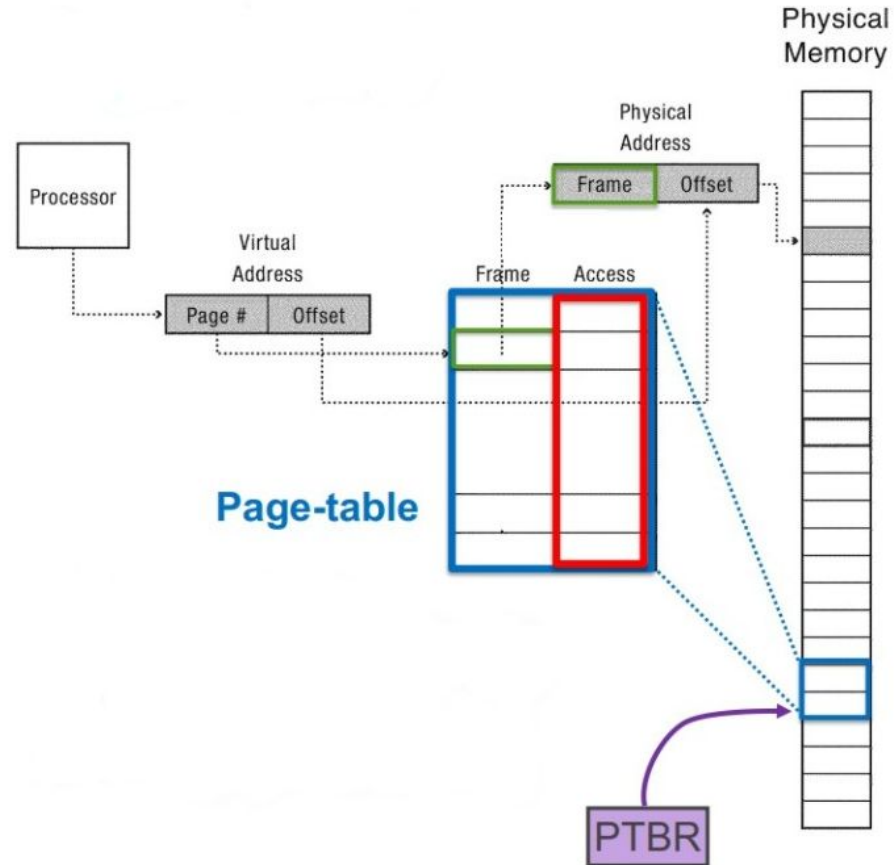
College of Computing

CS 450 Operating Systems

Paging and TLB

Yue Duan

Paging



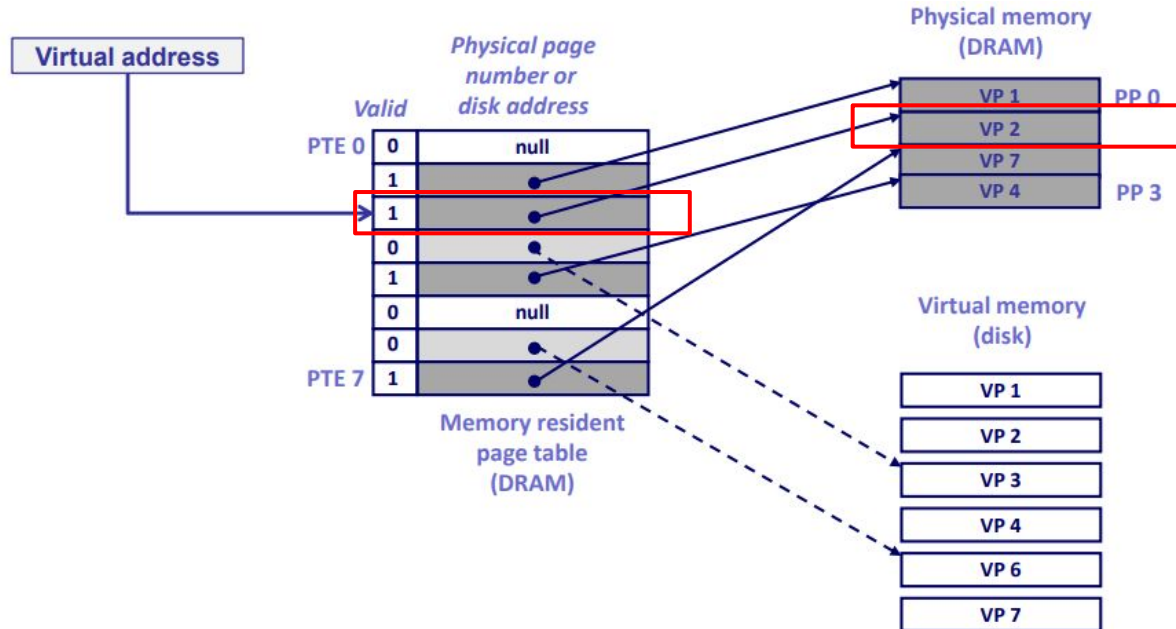
Paging Example



- Consider: 32-bit address space, pages are 4KB
 - Offset is 12 bits (because $4\text{KB} = 2^{12}$ bytes)
 - VPN is 20 bits
- Virtual address is 0x7468
 - Virtual page is 0x7, offset is 0x468
- Page table entry 0x7 contains 0x2
 - Page frame number is 0x2
 - Seventh virtual page is at address 0x2 (2nd physical page)
- Physical address = $0x2 \parallel 0x468 = 0x2468$

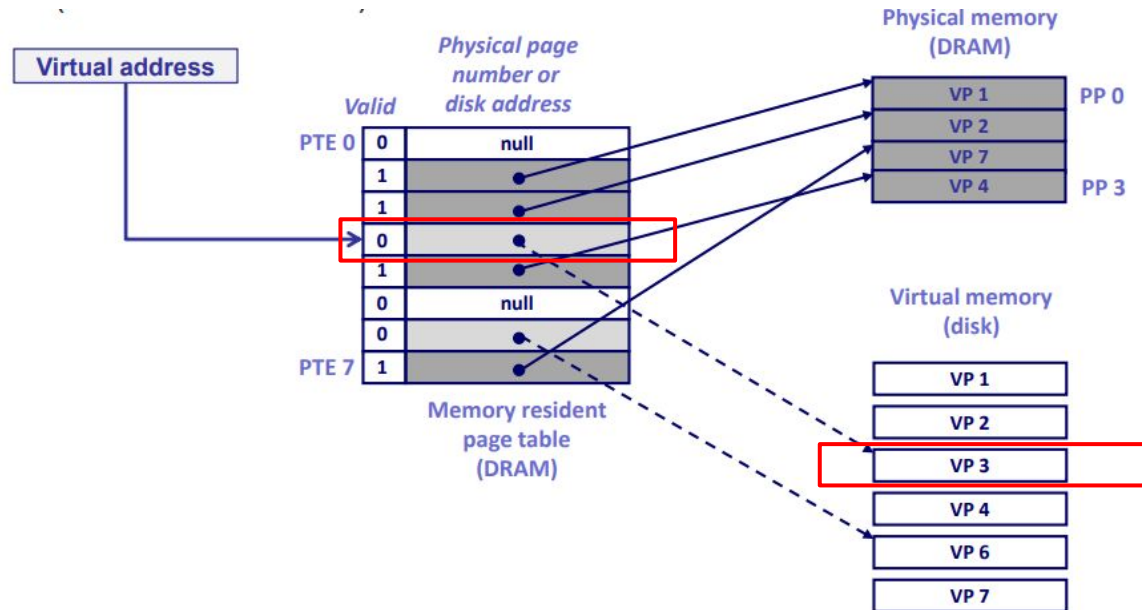
Paging Hit

- reference to VM word is in physical memory (DRAM cache hit)



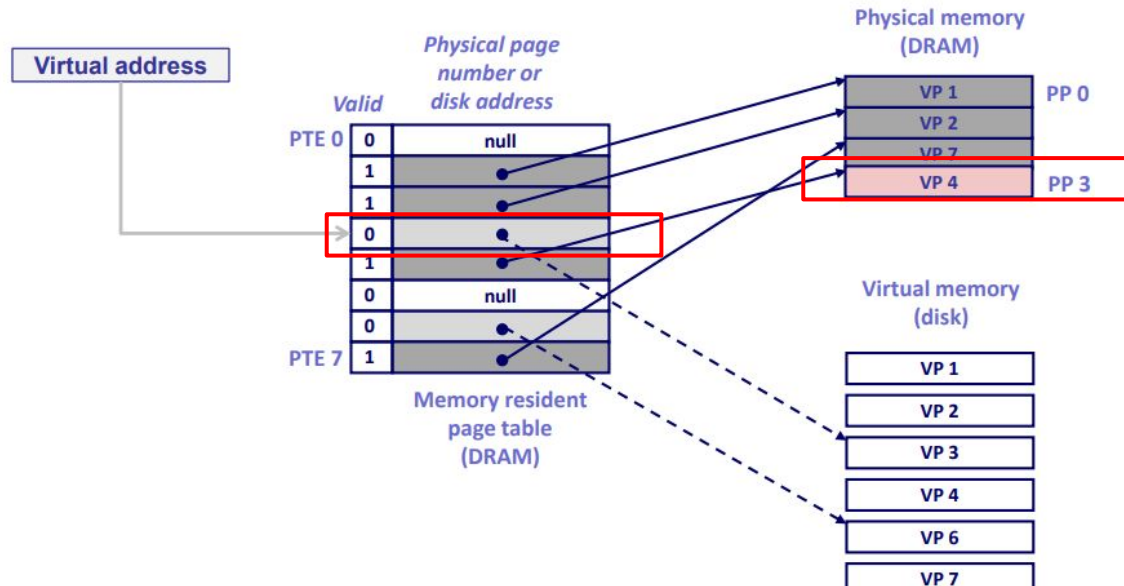
Paging Fault

- reference to VM word is NOT in physical memory (DRAM cache miss)



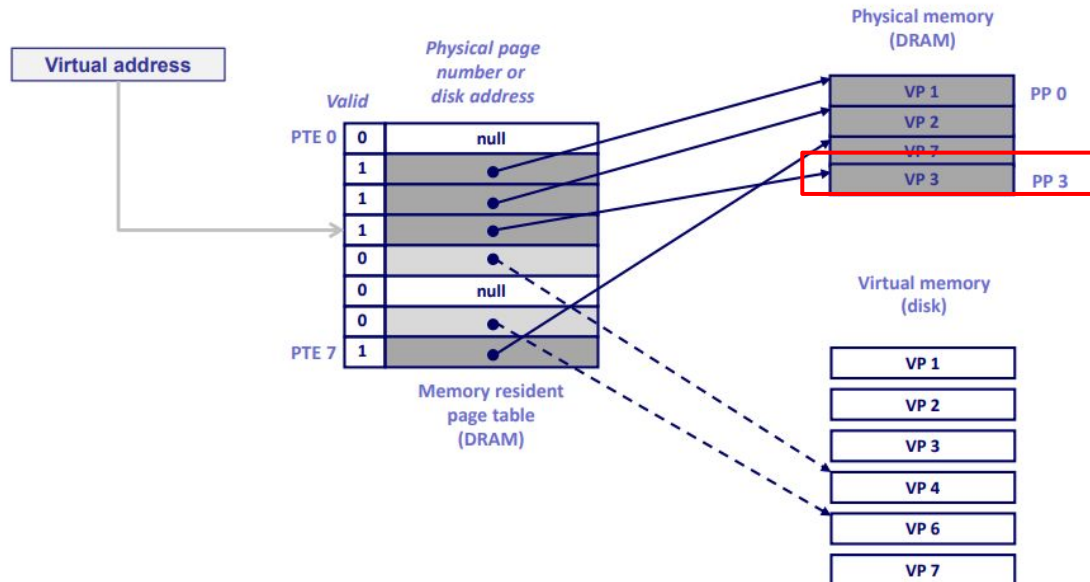
Handling Paging Fault

- Page miss causes page fault (an exception)
- Page fault handler selects a victim to be evicted (here VP 4)



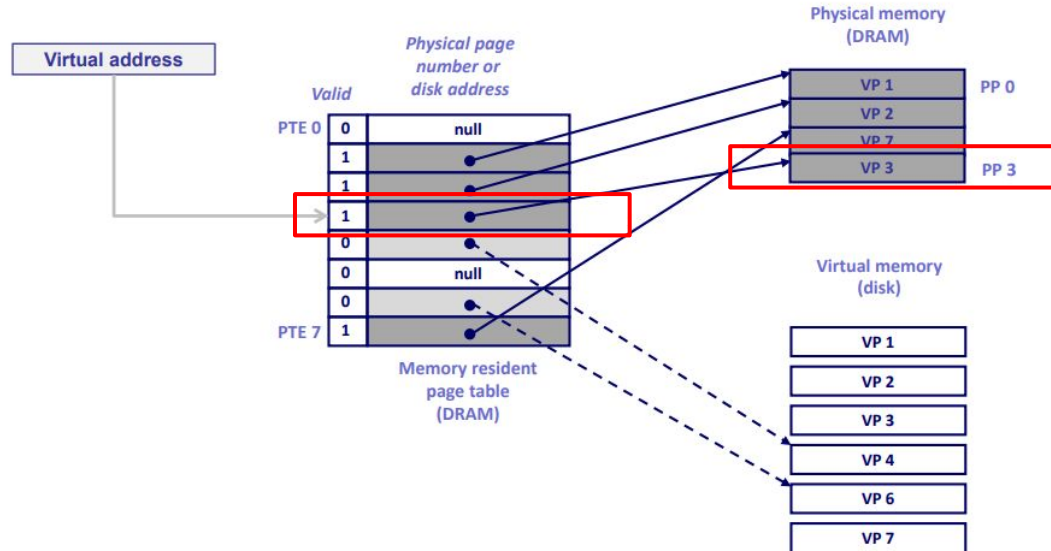
Handling Paging Fault

- Page miss causes page fault (an exception)
- Page fault handler selects a victim to be evicted (here VP 4)

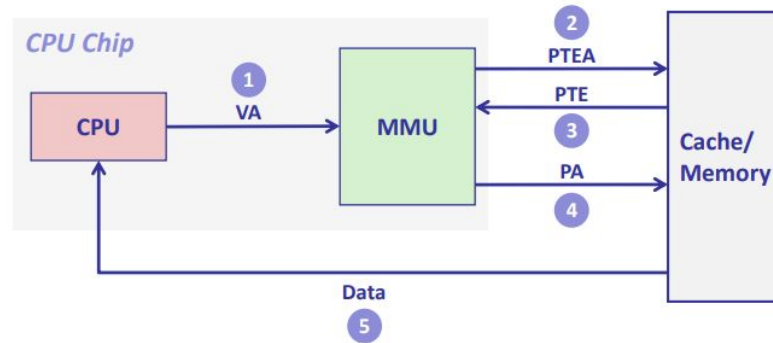


Handling Paging Fault

- Page miss causes page fault (an exception)
- Page fault handler selects a victim to be evicted (here VP 4)
- Offending instruction is restarted: page hit!

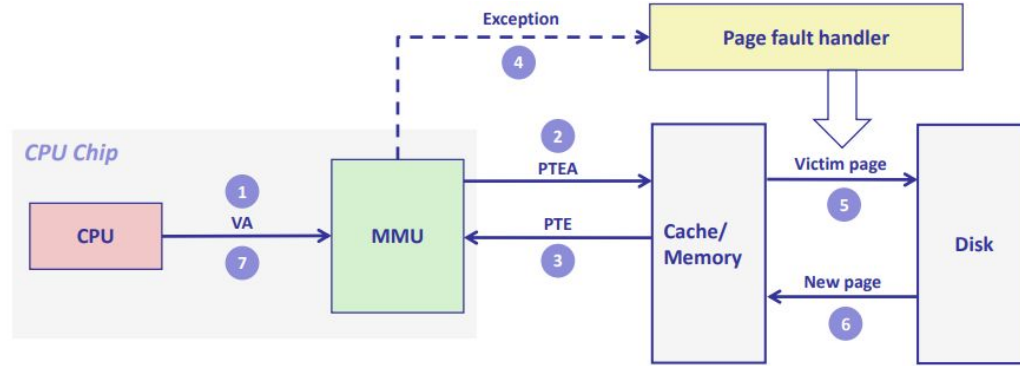


Address Translation: Page Hit



- 1) Processor sends virtual address to MMU
- 2-3) MMU fetches PTE from page table in memory
- 4) MMU sends physical address to cache/memory
- 5) Cache/memory sends data word to processor

Address Translation: Page Fault



- 1) Processor sends virtual address to MMU
- 2-3) MMU fetches PTE from page table in memory
- 4) Valid bit is zero, so MMU triggers page fault exception
- 5) Handler identifies victim (and, if dirty, pages it out to disk)
- 6) Handler pages in new page and updates PTE in memory
- 7) Handler returns to original process, restarting faulting instruction

Advantages of Paging



- Easily accommodates transparency, isolation, protection and sharing
- No external fragmentation
- Fast to allocate and free page frames
 - Alloc: No searching for suitable free space; pick the first free page frame
 - Free: Doesn't have to coalesce with adjacent free space; just add to the list of free page frames
 - Simple data structure (bitmap, linked list, etc.) to track free/allocated page frames

Disadvantages of Paging



- **Inefficiency:** memory references to page table
 - Page table must be stored in memory
 - MMU stores only base address of page table
- **Storage overhead:**
 - Internal fragmentation
 - Page size may not match size needed by process
 - Make pages smaller? But then...
 - page tables may be large
 - Simple page table: requires PTE for all pages in address space
 - Entry needed even if page not allocated ?

Page Translation Steps

%rip = 0x0010

```
0x0010: movl 0x1100, %edi
```

```
0x0013: addl $0x3, %edi
```

```
0x0019: movl %edi, 0x1100
```

Assume PT is at phys addr **0x5000**

Assume PTE's are **4 bytes**

Assume **4KB pages**

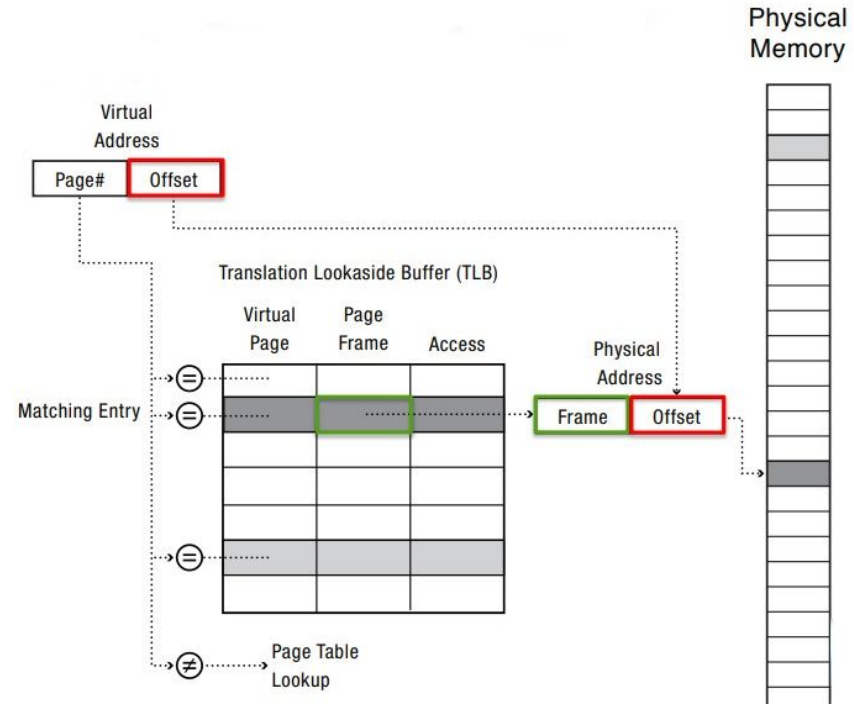
Simplified view
of page table

2
0
80
99

- Physical Memory Accesses with Paging
- Fetch instruction at virtual addr 0x0010
 - Access page table to get PFN for VPN 0
 - **Mem ref 1: 0x5000**
 - Learn VPN 0 is at PFN 2
 - Fetch instruction at 0x2010 (Mem ref 2)
- Exec, load from virtual addr 0x1100
 - Access page table to get PFN for VPN 1
 - **Mem ref 3: 0x5004**
 - Learn VPN 1 is at PFN 0
 - movl from 0x0100 into %edi (Mem ref 4)

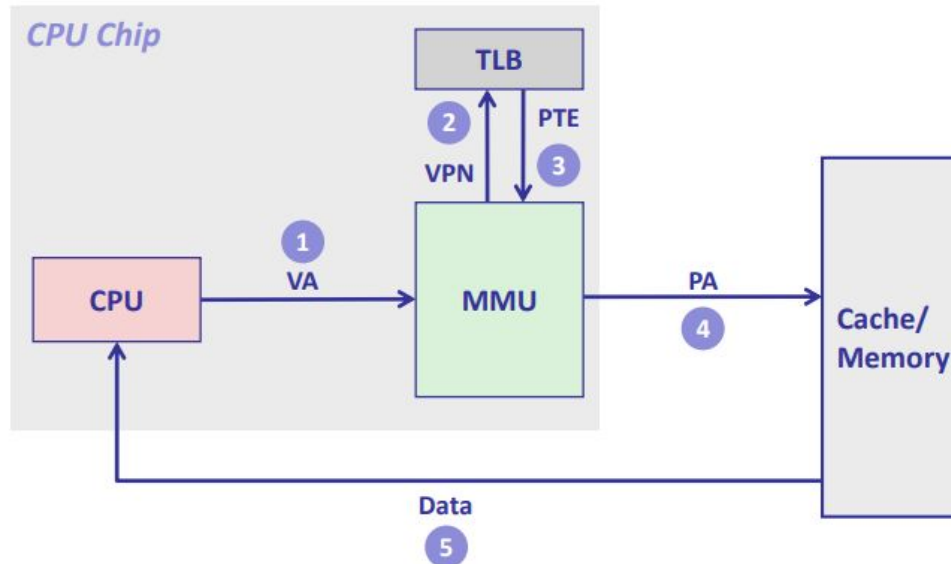
Speeding up Translation by Caching

- Translation Lookaside Buffer (TLB)
 - Small hardware cache in MMU
 - Maps virtual page numbers to physical page numbers
 - Contains complete page table entries for small number of pages



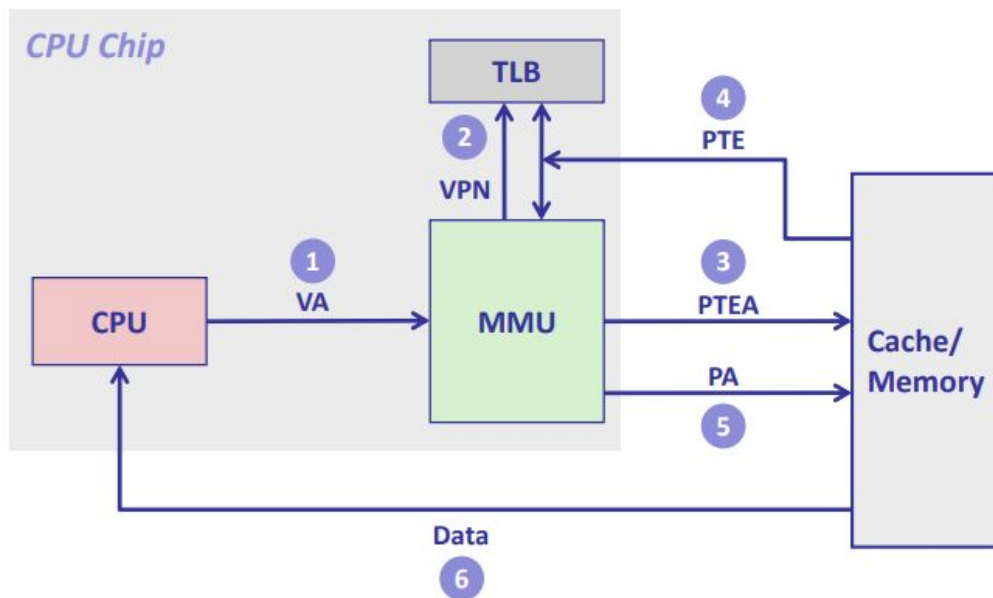
Address Translation with TLB

- Access TLB before you access memory
- A TLB hit eliminates a memory access



Address Translation with TLB

- A TLB miss incurs an additional memory access (the PTE)
- Fortunately, TLB misses are rare. Why?



TLB Entry



- TLB Entry
 - fully associative
 - any given translation can be anywhere in the TLB
 - hardware will search the entire TLB in parallel

Tag (virtual page number)	Physical page number (page table entry)
---------------------------	---

Array Iterator (w/ TLB)

```
int sum = 0;
for (i = 0; i < 2048; i++){
    sum += a[i];
}
```

Assume 'a' starts at 0x1000
Ignore instruction fetches
and access to 'i'

Assume following virtual address stream:

load 0x1000

load 0x1004

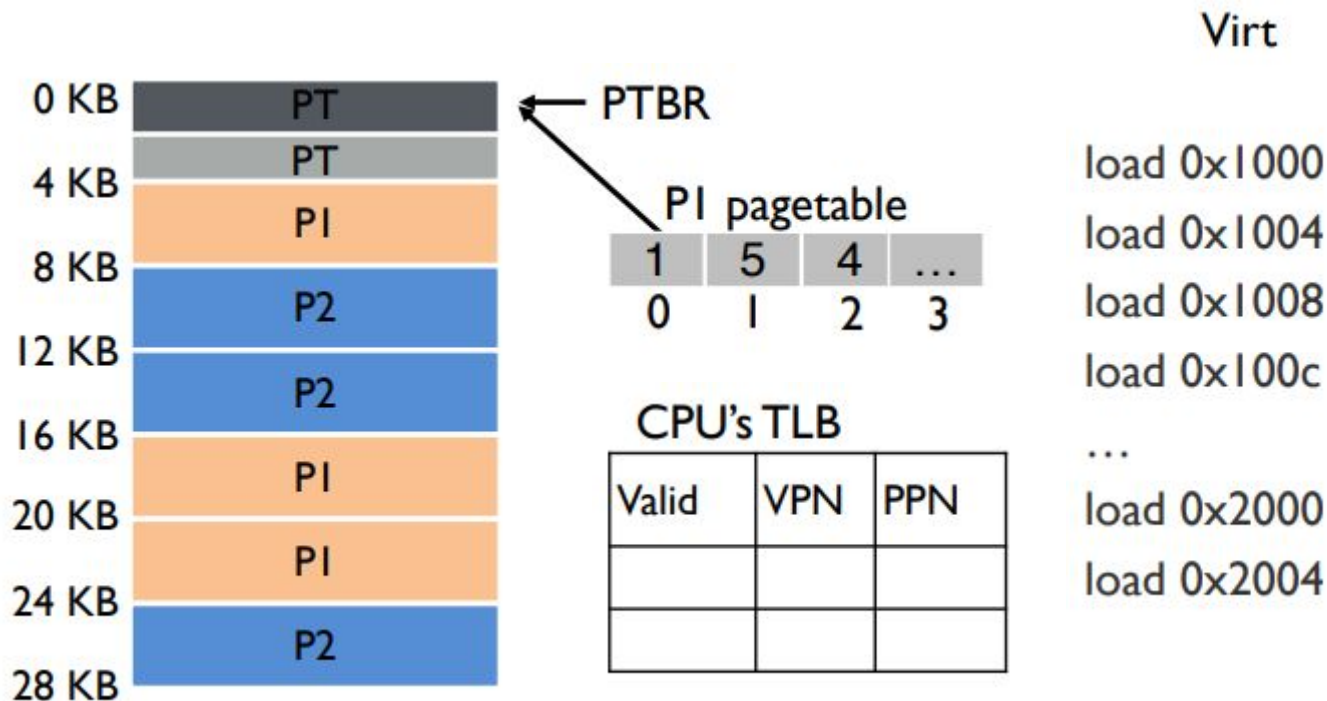
load 0x1008

load 0x100C

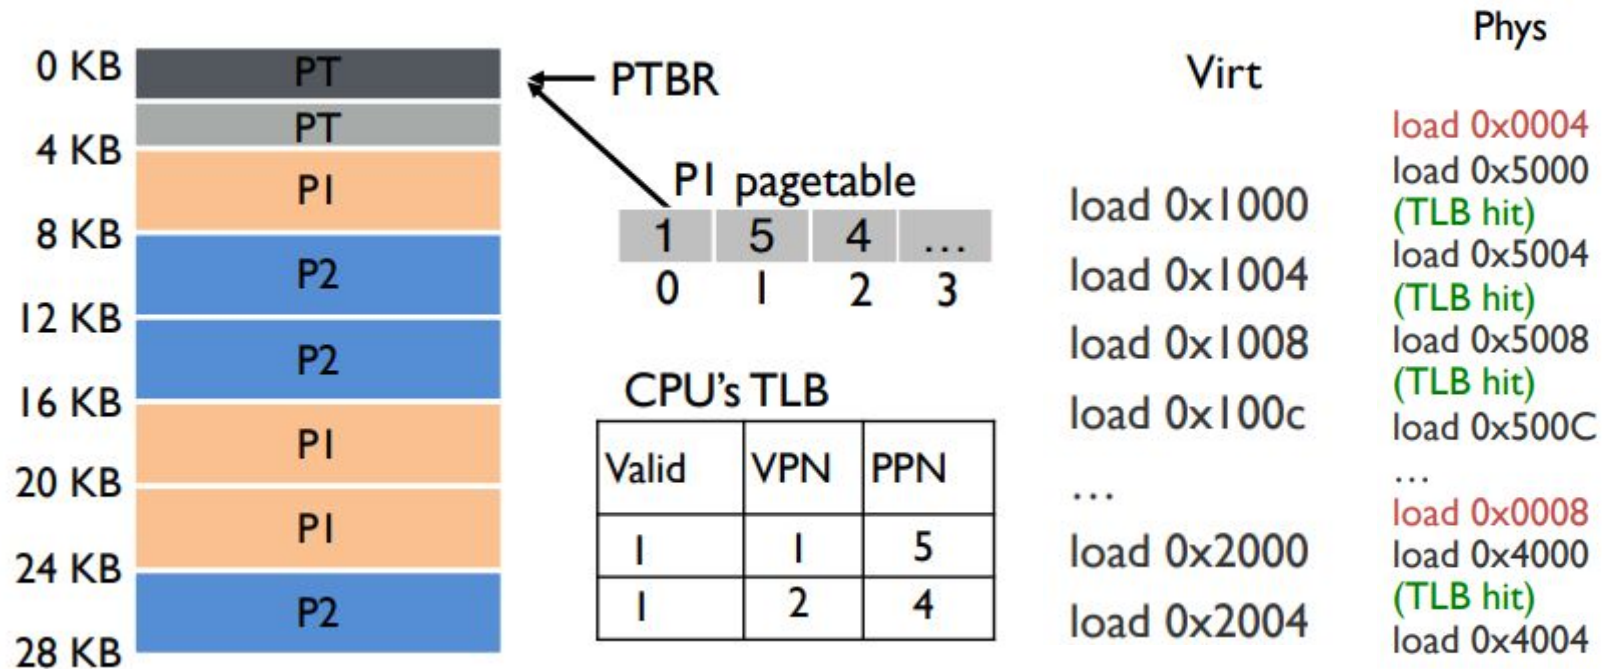
...

- What will TLB behavior look like?

Array Iterator (w/ TLB)



Array Iterator (w/ TLB)



Performance of TLB

```
int sum = 0;
for (i=0; i<2048; i++) {
    sum += a[i];
}
```

**Would hit rate get better or worse
with smaller pages?**

- Miss rate of TLB:
TLB misses / # TLB lookups
- # TLB lookups?
= number of accesses to a = 2048
- # TLB misses?
= number of unique pages accessed
= 2048 / (elements of 'a' per 4K page)
= 2K / (4K / sizeof(int)) = 2K / 1K
= 2
- Miss rate?
= 2/2048 = 0.1%
- Hit rate?
(1 - miss rate) = 99.9%

Workload Access Patterns





Workload A

```
int sum = 0;
for (i=0; i<2048; i++) {
    sum += a[i];
}
```

Workload B

```
int sum = 0;
srand(1234);
for (i=0; i<1000; i++) {
    sum += a[rand() % N];
}
srand(1234);
for (i=0; i<1000; i++) {
    sum += a[rand() % N];
}
```

Workload Access Patterns



Workload Locality



- Spatial Locality:
 - future access will be to nearby addresses
- Temporal Locality:
 - future access will be repeats to the same data
- What TLB characteristics are best for each type?

Workload Locality



- Spatial:
 - Access same page repeatedly
 - need same vpn => ppn translation
 - Same TLB entry re-used
- Temporal:
 - Access same address near in future
 - Same TLB entry re-used in near future
 - How near in future?
 - How many TLB entries are there?

TLB Replacement Policies

- LRU: evict Least-Recently Used TLB slot when needed



- Workload repeatedly accesses same offset (0x01) across 5 pages (strided access), but only 4 TLB entries
- What will TLB contents be over time?
- How will TLB perform?

TLB Replacement Policies



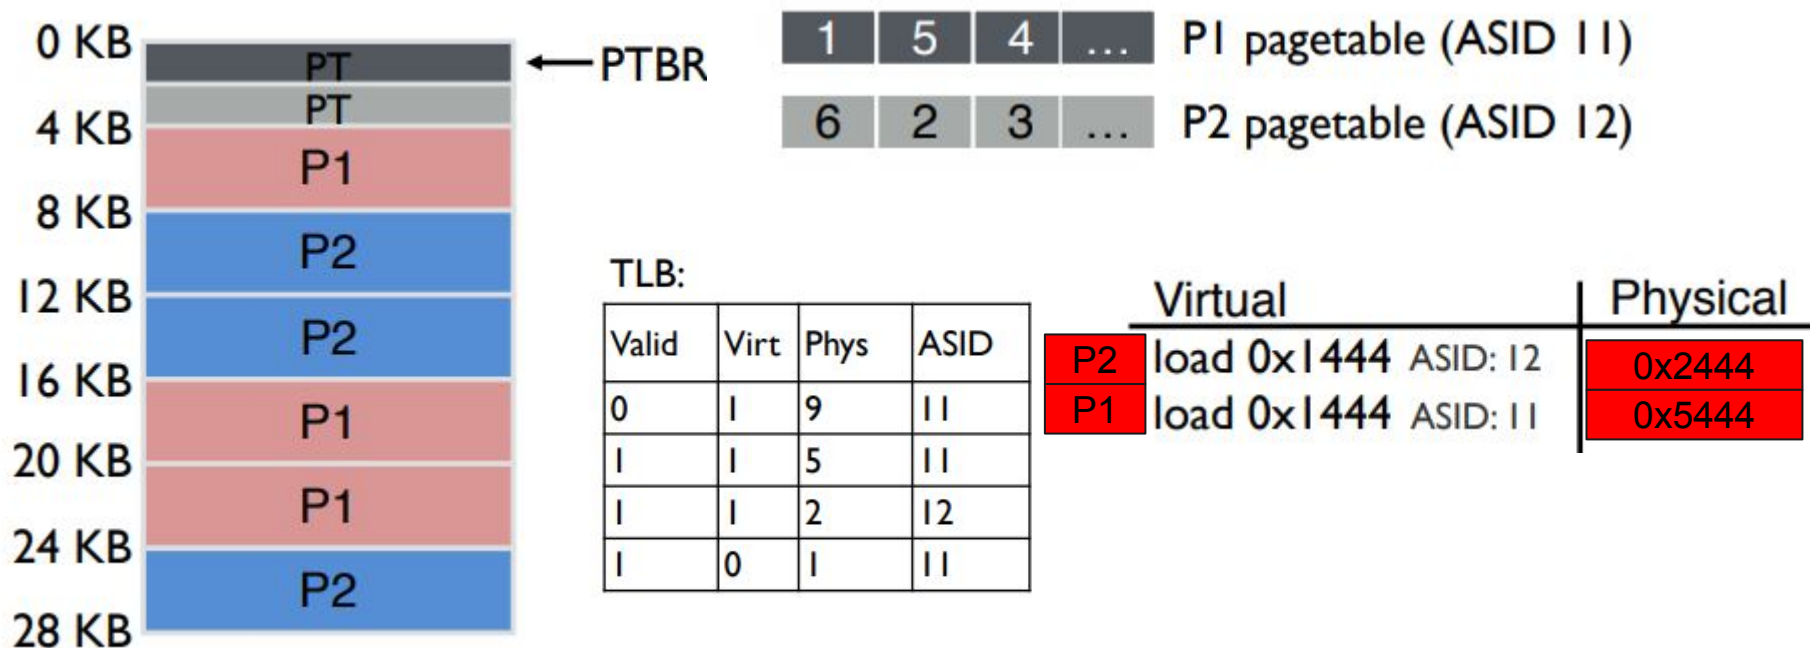
- Random: Evict randomly chosen entry
- Sometimes random is better than a “smart” policy!

Context Switches



- What happens if a process uses cached TLB entries from another process?
- Flush TLB on each switch
 - Costly \Rightarrow lose all recently cached translations
- Track which entries are for which process
 - Address Space Identifier
 - Tag each TLB entry with an 8-bit ASID

TLB Example with ASID



TLB Miss Handling



- If H/W handles TLB Miss
 - CPU must know where pagetables are
 - CR3 register on x86
 - Pagetable structure fixed and agreed upon between HW and OS
 - HW “walks” the pagetable and fills TLB
- If OS handles TLB Miss: “Software-managed TLB”
 - CPU traps into OS upon TLB miss.
 - OS interprets pagetables as it chooses
 - Modify TLB entries with privileged instruction

TLB Summary



- Pages are great, but accessing page tables for every memory access is slow
- Cache recent page translations \Rightarrow TLB
 - MMU performs TLB lookup on every memory access
- TLB performance depends strongly on workload
 - Sequential workloads perform well
 - Workloads with temporal locality can perform well
- In different systems, hardware or OS handles TLB misses
- TLBs increase cost of context switches
 - Flush TLB on every context switch
 - Add ASID to every TLB entry


Disadvantages of Paging



- **Inefficiency:** memory references to page table
 - Page table must be stored in memory
 - MMU stores only base address of page table
- **Storage overhead:**
 - Internal fragmentation
 - Page size may not match size needed by process
 - Make pages smaller? But then...
 - page tables may be large
 - Simple page table: requires PTE for all pages in address space
 - Entry needed even if page not allocated ?


Multi-Level Page Tables

- + Allocate only PTEs in use
- + Simple memory allocation
 - no large continuous memory space needed
- – more lookups per memory reference

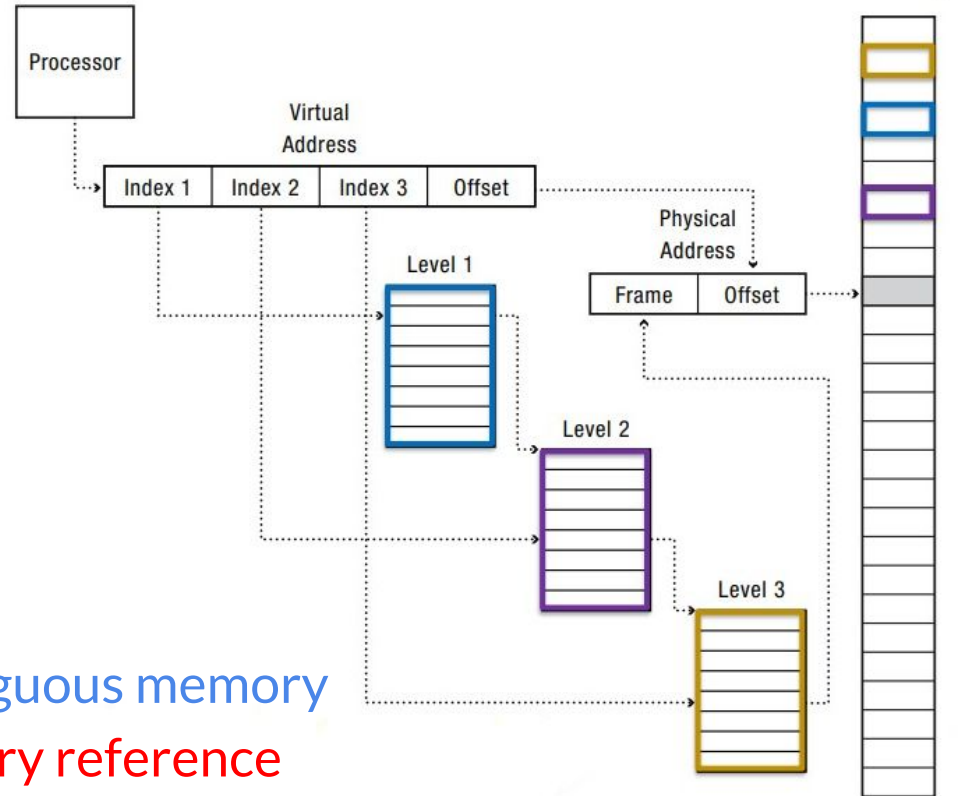


Two-Level Paging Example

- Assume 32-bit machine, 1KB page size
- Logical address is divided into:
 - a page offset of 10 bits ($1024 = 2^{10}$)
 - a page number of 22 bits ($32 - 10$)
- The page number is further divided into:
 - a 12-bit first index
 - a 10-bit second index



Three-Level Paging is also Possible



- + First Level requires less contiguous memory
- - even more lookups per memory reference



THANK YOU!