ILLINOIS TECH

College of Computing

CS 450 Operating Systems

Processes

Yue Duan

Process vs Program

- A program consists of code and data
  - specified in some programming language, such as C
- Typically stored in a file on disk
- “Running a program” = creating a process
- you can run a program multiple times!
  - one after another or even concurrently

Executable

- An executable is a file containing:
  - executable code
    - CPU instructions
  - data
    - information manipulated by these instructions
- Obtained by compiling a program
  - and linking with libraries

Executable

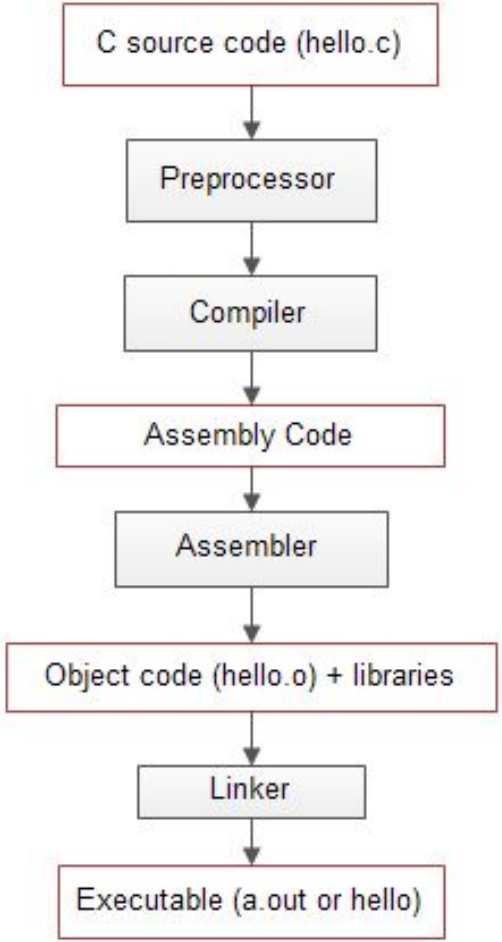
C source code (hello.c)
Preprocessor
Compiler
Assembly Code
Assembler
Object code (hello.o) + libraries
Linker
Executable (a.out or hello)

Executable

```c
#include<stdio.h>

int main ()
{
   printf("hello world!");
   return 0;
}
```

Compiler, assembler, linker

Process vs Program

- An executable running on an abstraction of a computer:
  - Address Space (memory) + Execution Context (registers incl. PC and SP)
    - manipulated through machine instructions
  - Environment (clock, files, network, ...)
    - manipulated through system calls
- Current state is called “image” in Thompson/Ritchie paper
- A good abstraction:
  - is portable and hides implementation details
  - has an intuitive and easy-to-use interface
  - can be instantiated many times

Process != Program

- A program is **passive**:
  - code + data
- A process is **alive**:
  - mutable data + registers + files + ...
- Same program can be run multiple time simultaneously (1 program, 2 processes)
  - > ./program &
  - > ./program &

A Day in the Life of a Program

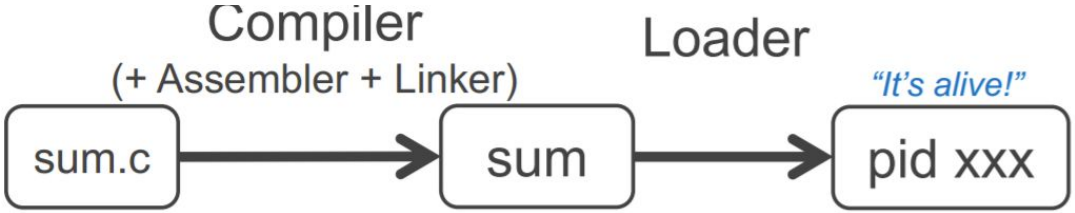


```c
#include <stdio.h>

int max = 10;

int main () {
   int i;
   int sum = 0;
   add(m, &sum);
   printf("%d",i);
   ...

}
```

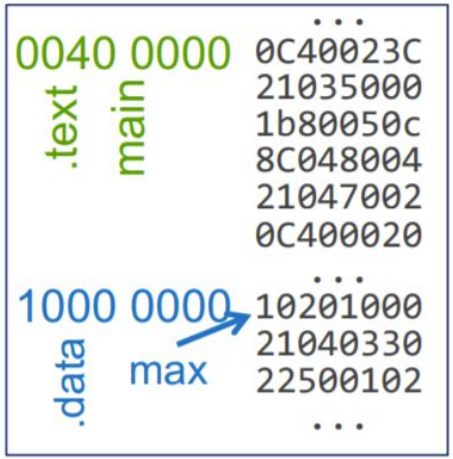

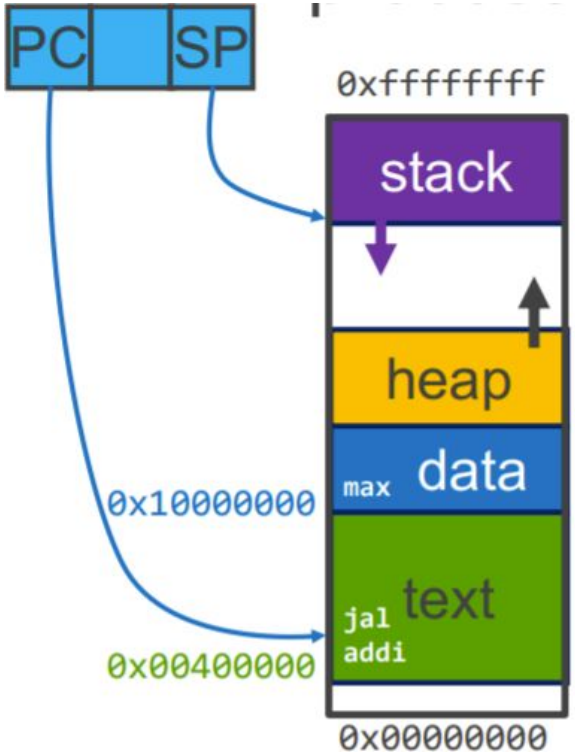

Process Creation

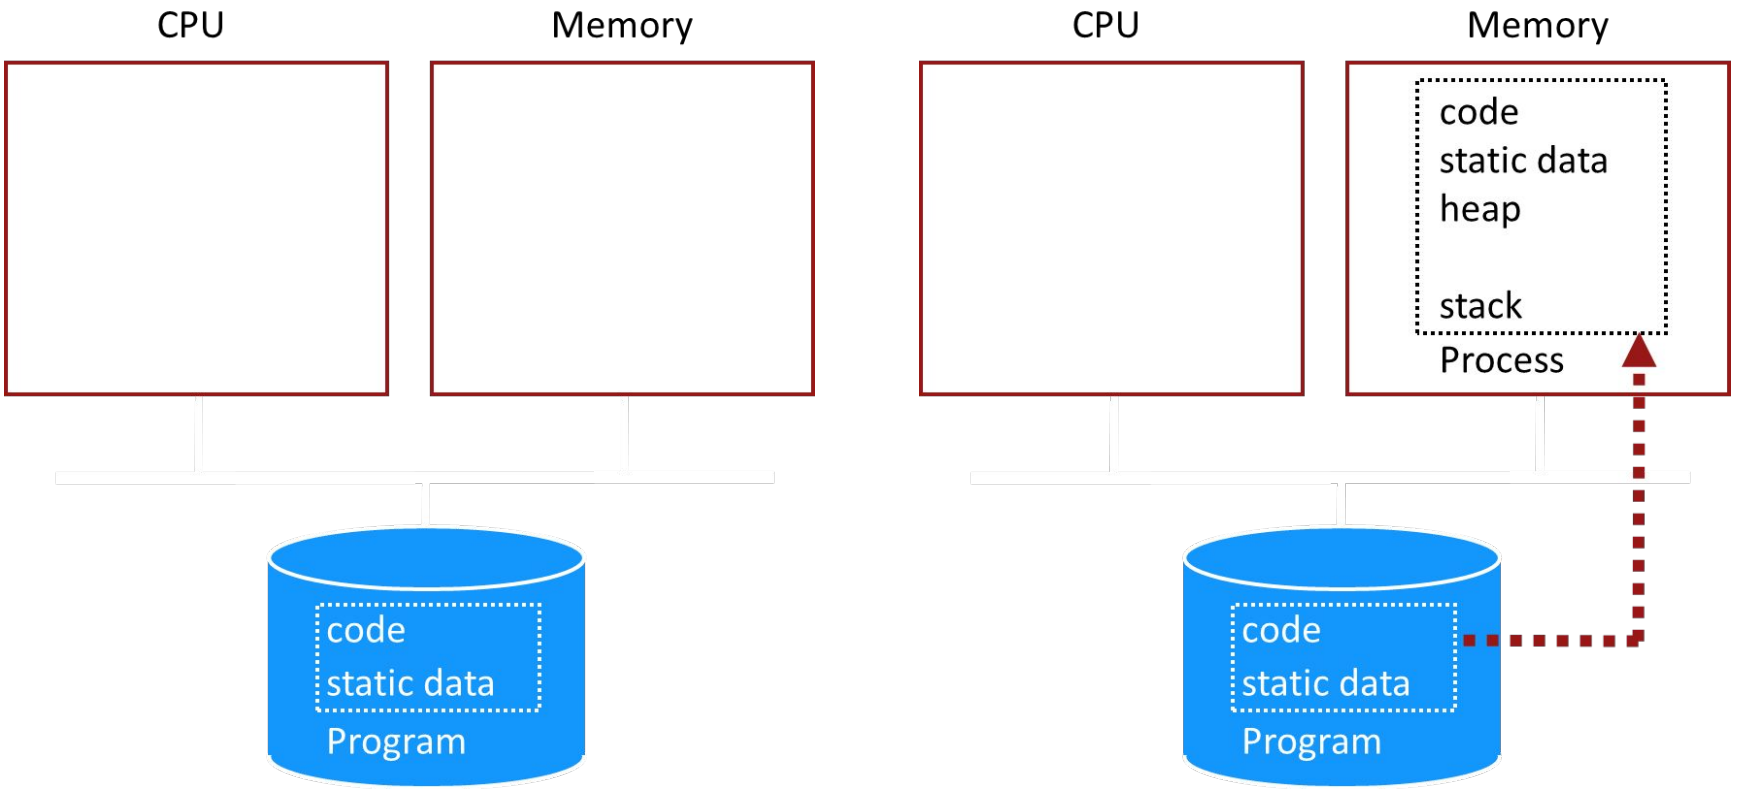

Logical view of process memory

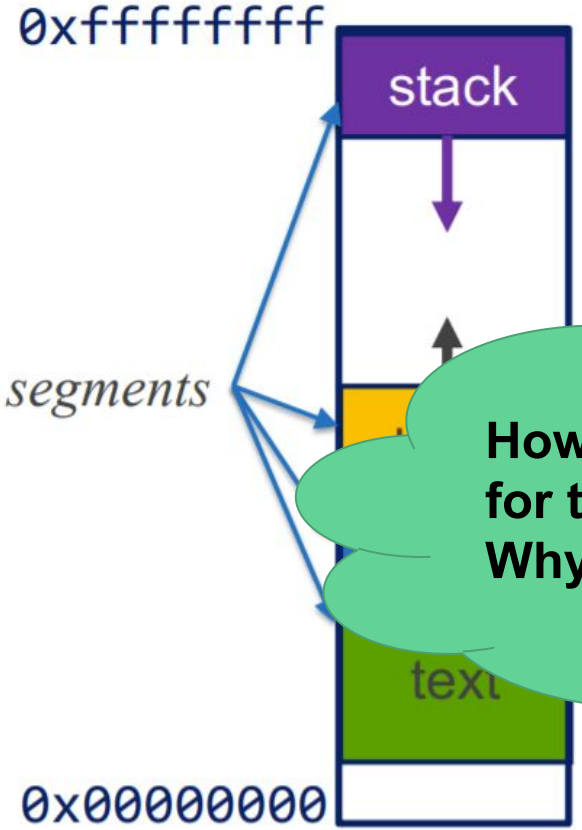

- call stack

How many bits in an address for this CPU?
Why is address 0 not mapped?

n (malloc)

riables

contains code and

onstants

Stack

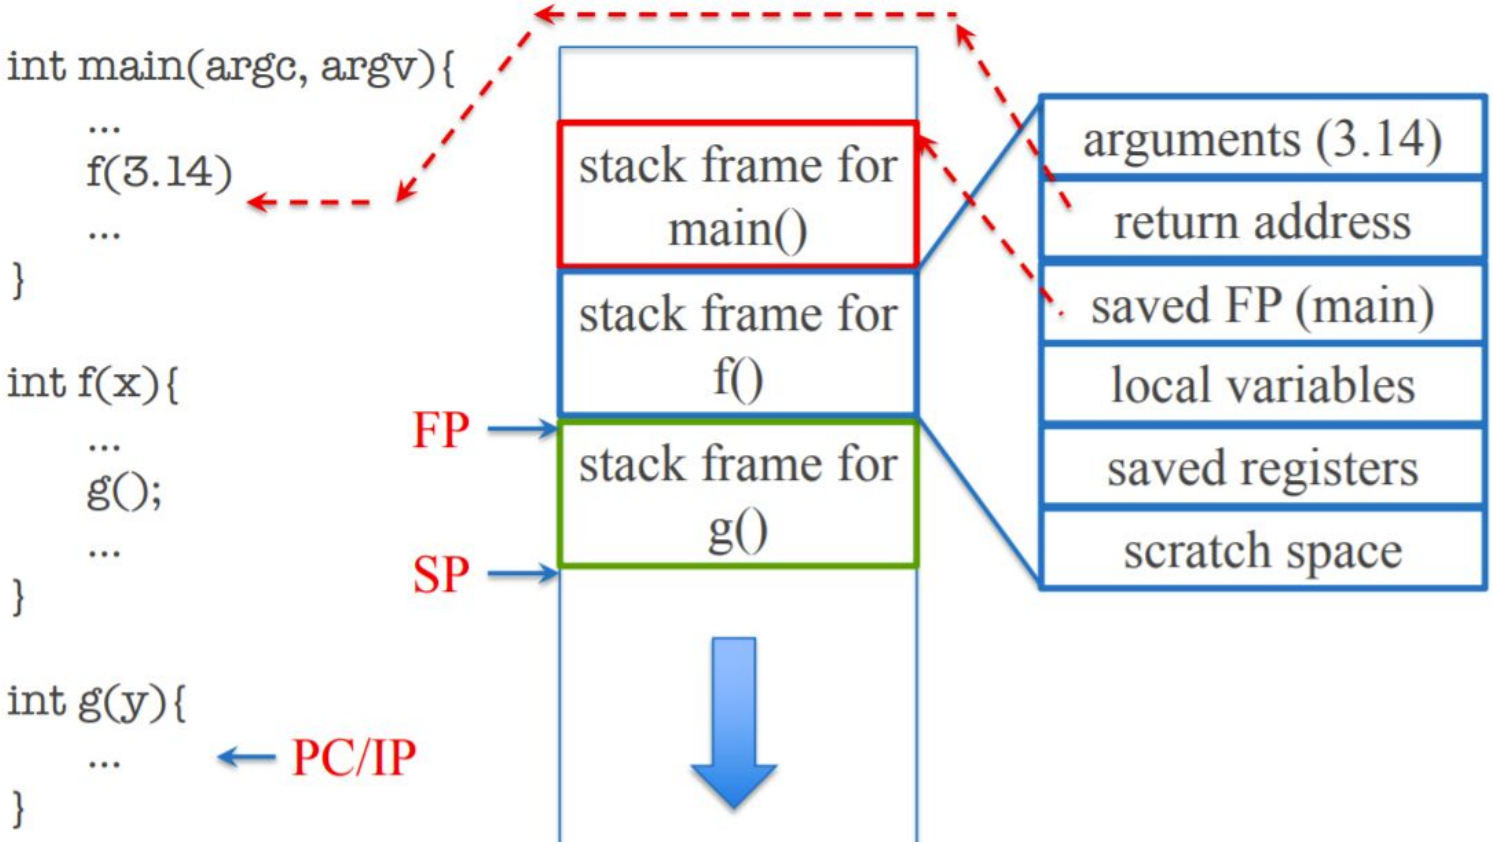
int main(argc, argv){
...
f(3.14)
...
}
int f(x){
...
g();
...
}
int g(y){
...
}
PC/IP
FP
SP
stack frame for main()
stack frame for f()
stack frame for g()
arguments (3.14)
return address
saved FP (main)
local variables
saved registers
scratch space

Virtualizing the CPU

- Goal:
  - Give each process impression it alone is actively using CPU
- Resources can be shared in time and space
- Assume single uniprocessor
  - Time-sharing (multi-processors: advanced issue)
- Memory?
  - Space-sharing (later)
- Disk?
  - Space-sharing (later)

How to Provide Good CPU Performance?

- Direct execution
  - Allow user process to run directly on hardware
  - OS creates process and transfers control to starting point (i.e., main())
- Problems with direct execution?
  - Process could do something restricted
    - Could read/write other process data (disk, memory) or restricted device Process could run forever (slow, buggy, or malicious)
  - OS needs to be able to switch between processes
    - Process could do something slow (like I/O)
- OS wants to use resources efficiently and switch CPU to other process

Problem 1: Restricted Operations

- How can we ensure user process can’t harm others?
- Solution: privilege levels supported by hardware (bit of status)
  - User processes run in user mode (restricted mode)
  - OS runs in kernel mode (not restricted)
    - Instructions for interacting with devices
    - Could have many privilege levels (advanced topic)
- How can process access device?
  - System calls (function call implemented by OS)
  - Change privilege level through system call (trap)

System Call

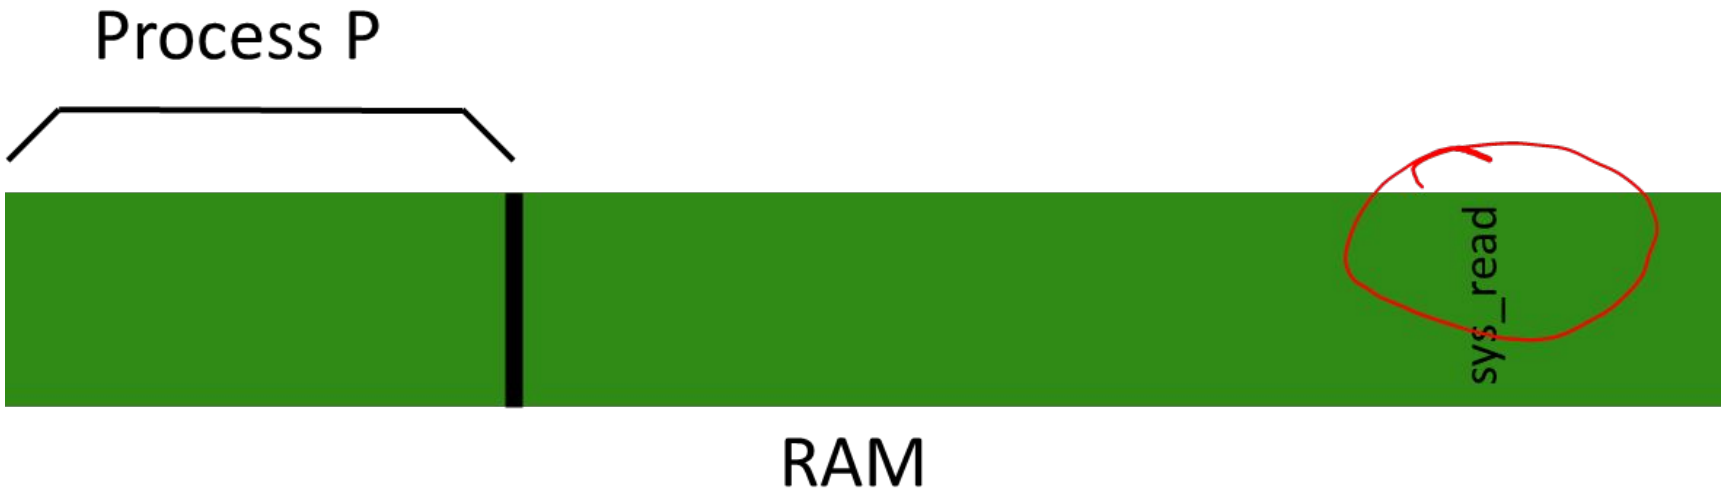

P wants to call read()

System Call

P can only see its own memory because of **user mode**
(other areas, including kernel, are hidden)

System Call

P wants to call read() but no way to call it directly

System Call

- A process runs on CPU
- Can access O.S. kernel through “system calls”
- Skinny interface - Why?

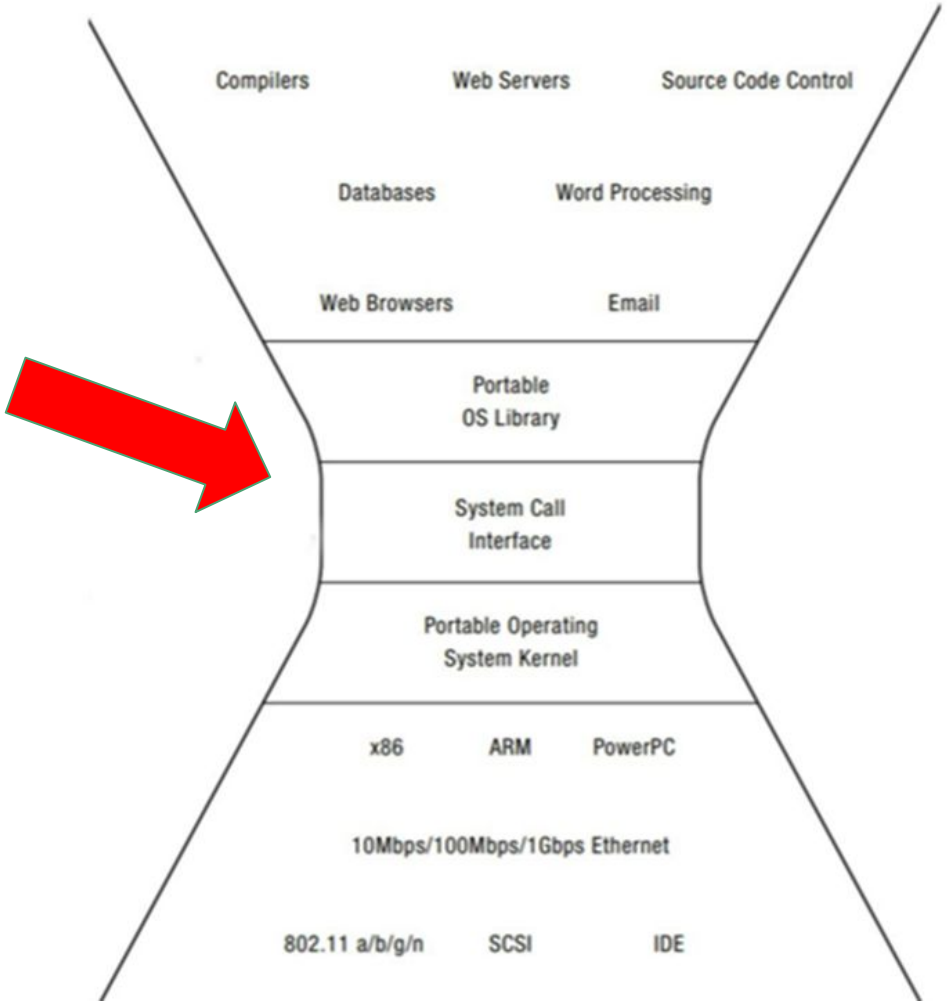

Why a “skinny” interface?

- Portability
  - easier to implement and maintain
  - e.g., many implementations of “Posix” interface
- Security
  - “small attack surface”: easier to protect against vulnerabilities

Executing a system call

- Process:
  - 1. Calls system call function in library
  - 2. Places arguments in registers and/or pushes them onto user stack
  - 3. Places syscall type in a dedicated register
  - 4. Executes **syscall** machine instruction
- Kernel:
  - 5. Executes syscall interrupt handler
  - 6. Places result in dedicated register
  - 7. Executes **return_from_interrupt**
- Process:
  - 8. Executes **return_from_function**

Executing read System Call

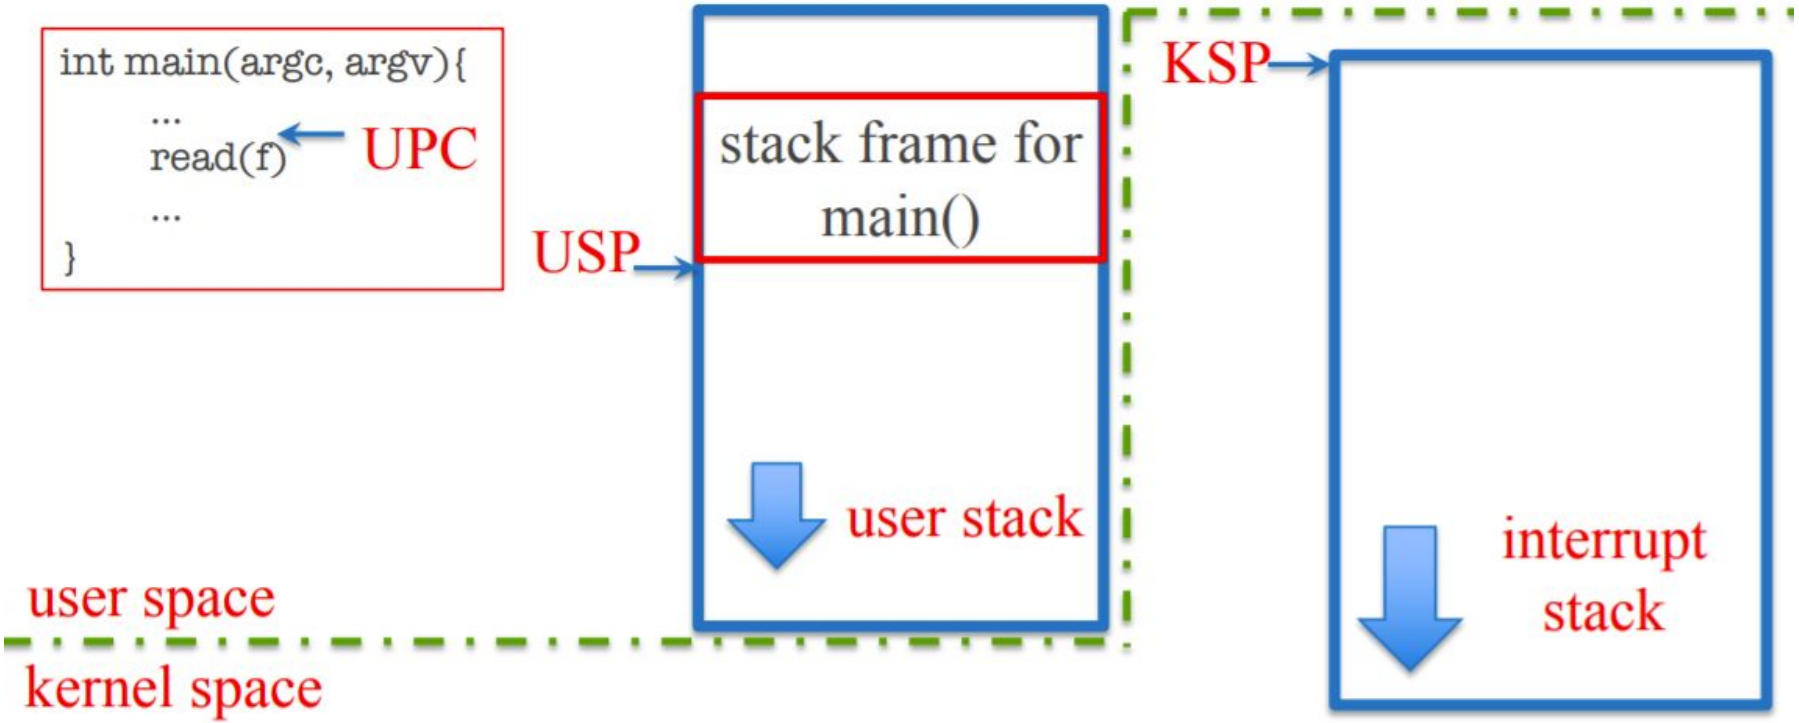

Executing read System Call

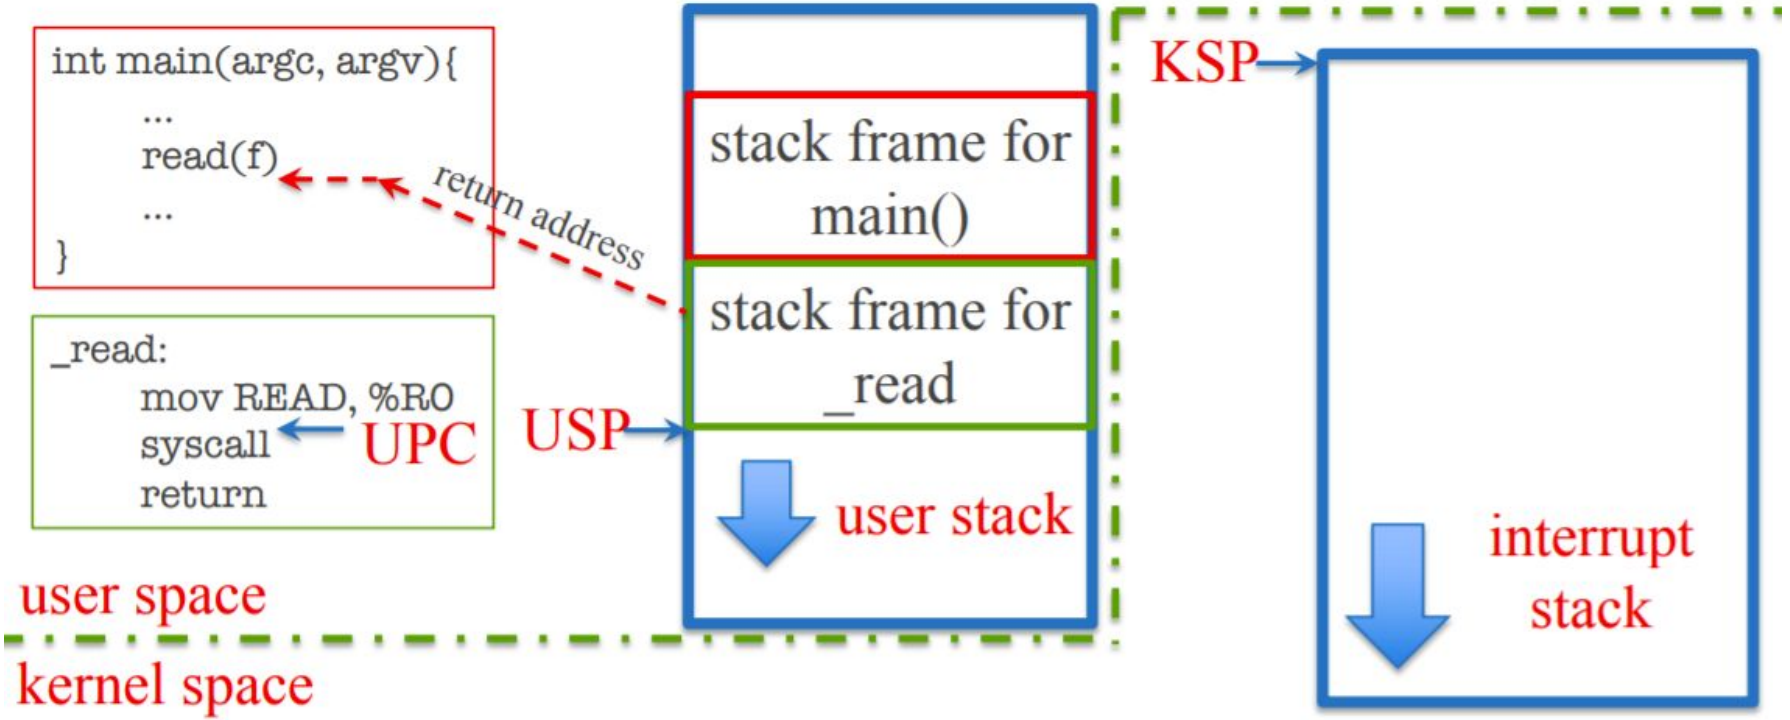

Executing read System Call

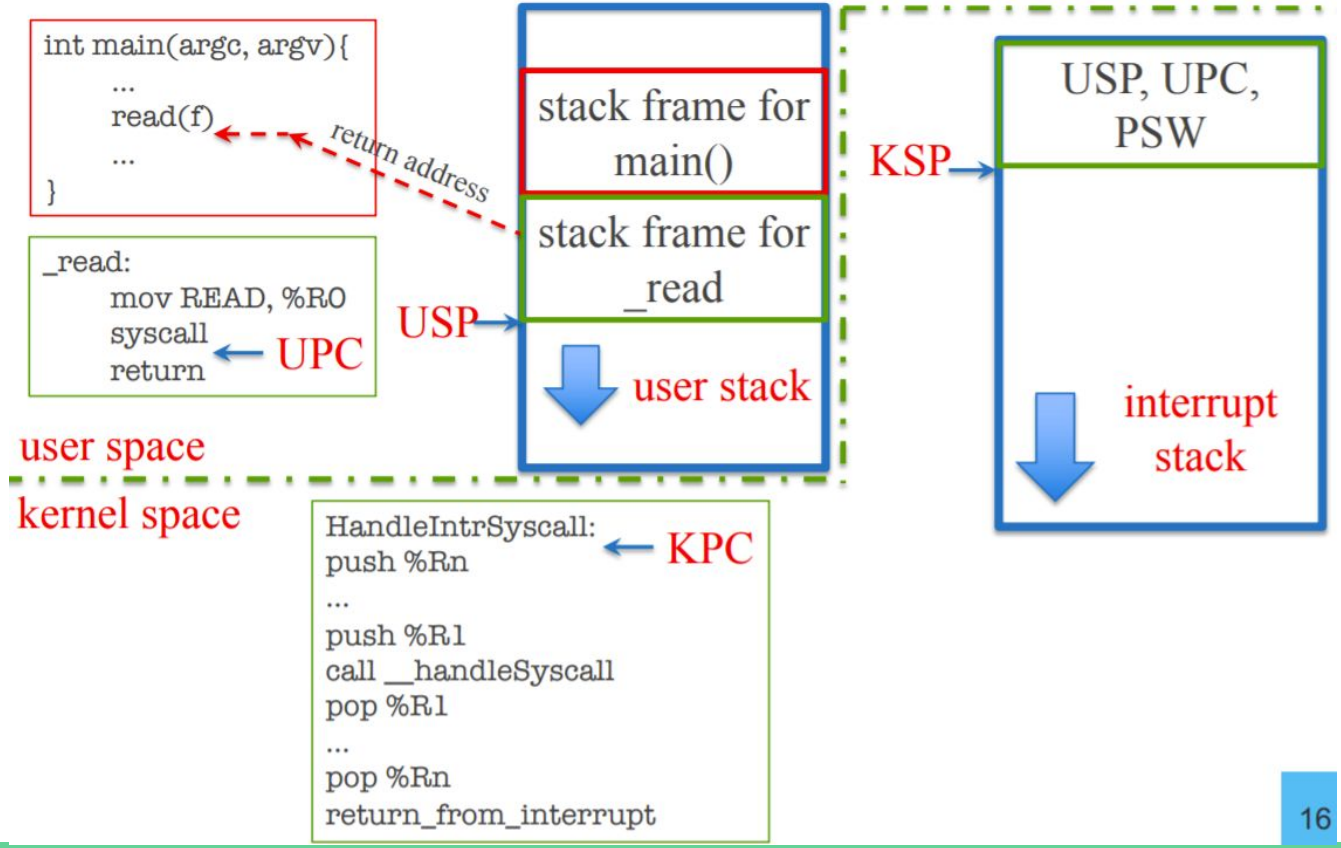

Executing read System Call

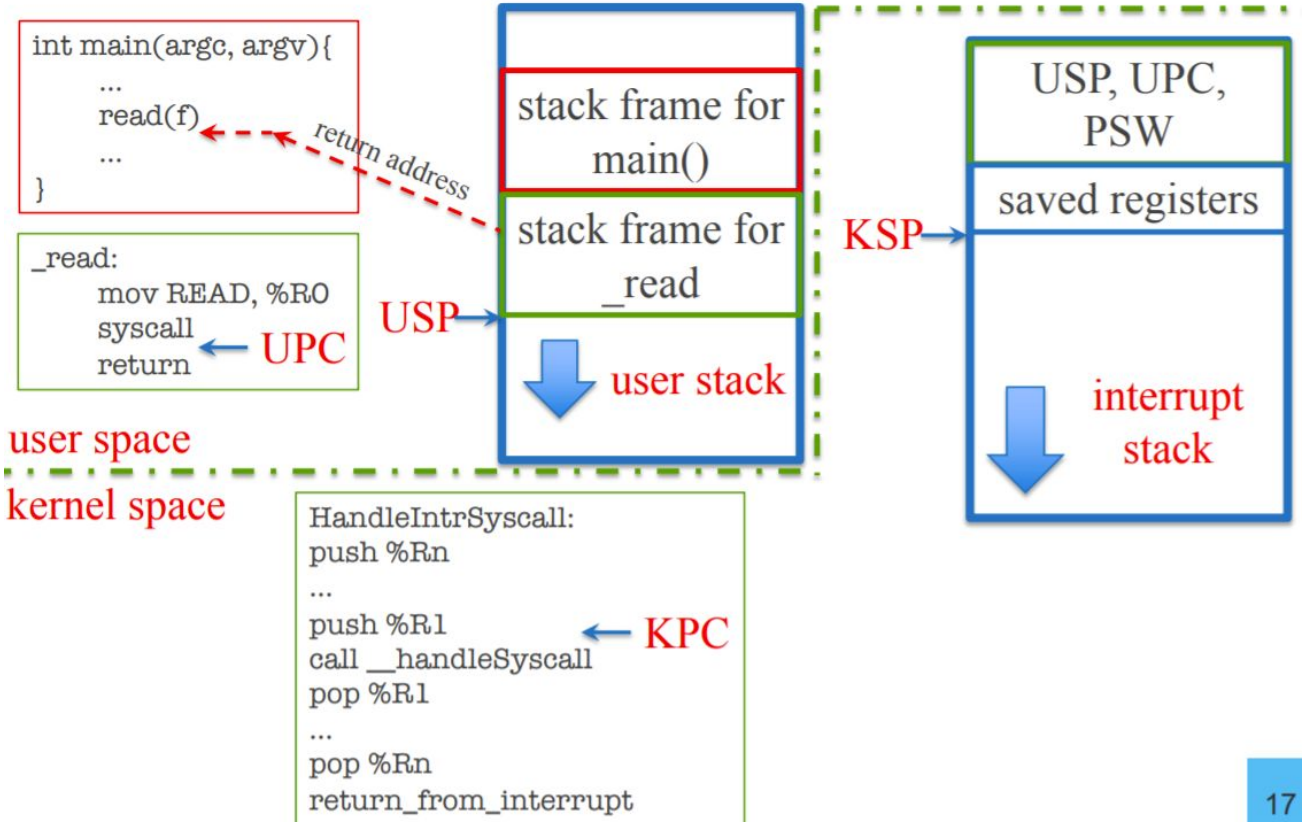

Executing read System Call

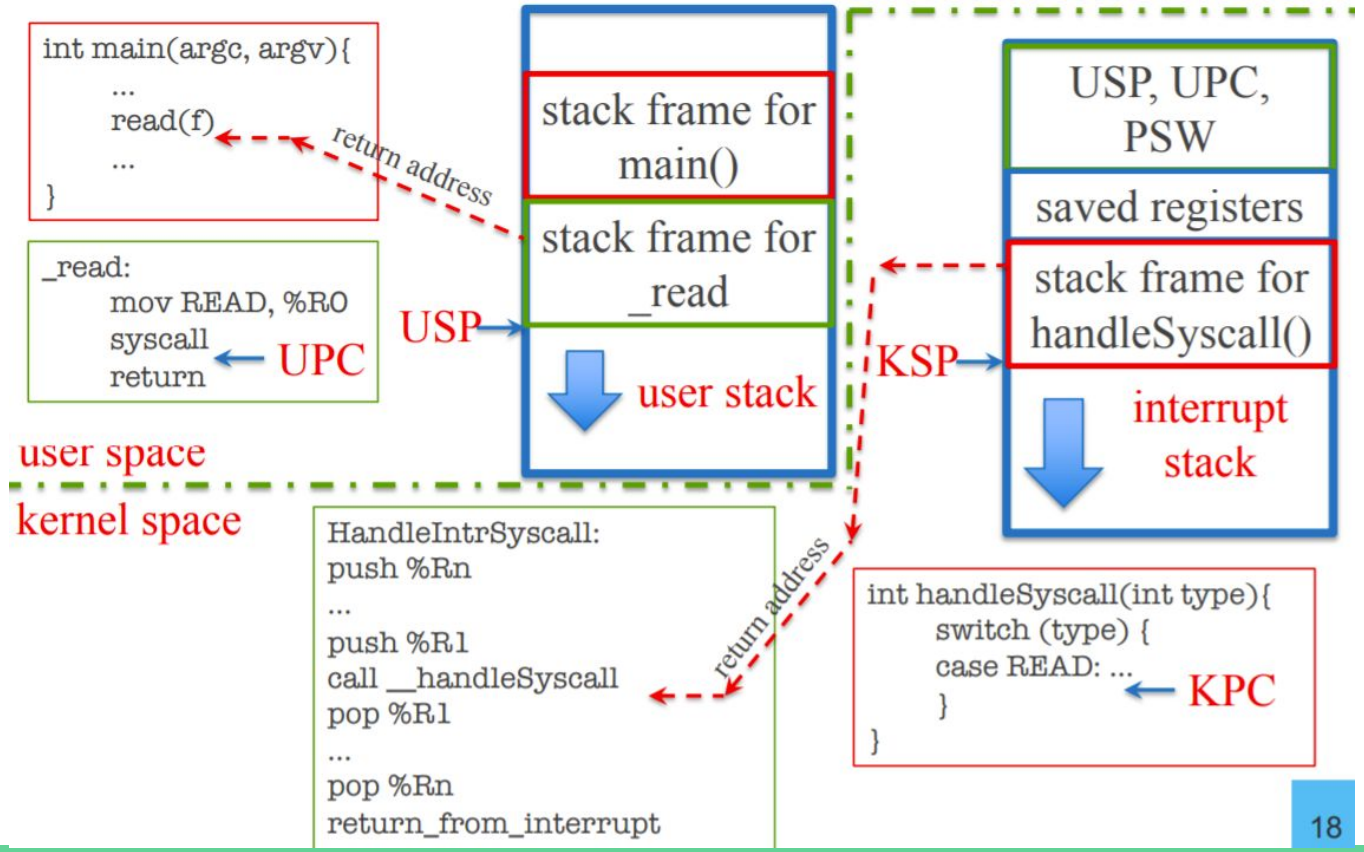

What if read needs to “block”?

- read may need to block if
  - reading from terminal
  - reading from disk and block not in cache
  - reading from remote file server
  - etc

should run another process!

THANK YOU!