



ILLINOIS TECH

College of Computing

CS 450 Operating Systems

Swapping & Intro to IO

Yue Duan

Recap: Swapping



- OS goal: Support multiple processes
 - Single process with very large address space
 - Multiple processes with combined address spaces
- User code should be independent of amount of physical memory
 - Correctness, if not performance
- Virtual memory: OS provides illusion of more physical memory
- Reality: many processes, limited physical memory

Recap: Swapping



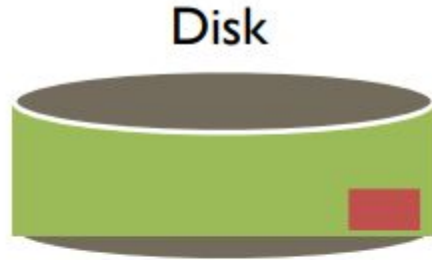
- Idea: OS keeps unreferenced pages on disk
 - Slower, cheaper backing store than memory
- Process can run when not all pages are loaded into main memory
- OS and hardware cooperate to make large disk seem like memory
 - Same behavior as if all of address space in main memory
- Requirements:
 - OS must have **mechanism** to identify location of each page in address space ==> in memory or on disk
 - OS must have **policy** to determine which pages live in memory and which on disk

Virtual Address Space Mechanisms



- Each page in virtual address space maps to one of three locations:
 - Physical main memory: Small, fast, expensive
 - Disk (backing store): Large, slow, cheap
 - Nothing (error): Free
- Extend page tables with an extra bit: present
 - permissions (r/w), valid, present
 - Page in memory: present bit set in PTE
 - Page on disk: present bit cleared
 - PTE points to block on disk
 - Causes trap into OS when page is referenced
 - **Trap: page fault**

Virtual Address Space Mechanisms



Phys Memory



PFN	valid	prot	present
10	1	r-x	1
-	0	-	-
23	1	rw-	0
-	0	-	-
-	0	-	-
-	0	-	-
-	0	-	-
-	0	-	-
-	0	-	-
-	0	-	-
28	1	rw-	0
4	1	rw-	1

What if access vpn 0xb?

Virtual Address Space Mechanisms



- First, hardware checks TLB for virtual address
- if TLB hit, address translation is done; page in physical memory
- Else
 - Hardware or OS walk page tables
 - If PTE designates page is present, then page in physical memory (i.e., present bit is cleared)
 - Else
 - Trap into OS (not handled by hardware)
 - OS selects victim page in memory to replace
 - Write victim page out to disk if modified (use dirty bit in PTE)
 - OS reads referenced page from disk into memory
 - Page table is updated, present bit is set
 - Process continues execution

Swapping Policies



- Goal: Minimize number of page faults
 - Page faults require milliseconds to handle (reading from disk)
 - Implication: Plenty of time for OS to make good decision
- OS has two decisions
 - Page selection
 - When should a page (or pages) on disk be brought into memory?
 - Page replacement
 - Which resident page (or pages) in memory should be thrown out to disk?

Page Selection



- **Demand paging:**
 - Load page only when page fault occurs
 - Intuition: Wait until page must absolutely be in memory
 - When process starts: No pages are loaded in memory
 - Problems: Pay cost of page fault for every newly accessed page
- **Prepaging** (anticipatory, prefetching):
 - Load page before referenced
 - OS predicts future accesses (**oracle**) and brings pages into memory
 - Works well for some access patterns (e.g., sequential)

Page Selection



- Hints: Combine above with user-supplied hints about page references
 - User specifies: may need page in future, don't need this page anymore, or sequential access pattern
 - Example: `madvise()` in Unix
 - `madvise` - give advice about use of memory

Page Replacement



- Which page in main memory should be selected as victim?
 - Write out victim page to disk if modified (dirty bit set)
 - If victim page is not modified (clean), just discard
- **OPT**: Replace page **not used for longest time in future**
 - Advantages: Guaranteed to minimize number of page faults
 - Disadvantages: Requires that OS predict the future; Not practical, but good for comparison

Page Replacement

- **FIFO**: Replace page that has been in memory the longest
 - Intuition: First referenced long time ago, done with it now
 - Advantages: Fair: All pages receive equal residency; Easy to implement
 - Disadvantage: Some pages may always be needed
- **LRU**: Replace page not used for longest time in past
 - Intuition: Use past to predict the future
 - Advantages: With locality, LRU approximates OPT
 - Disadvantages:
 - Harder to implement, must track which pages have been accessed
 - Does not handle all workloads well

Comparison



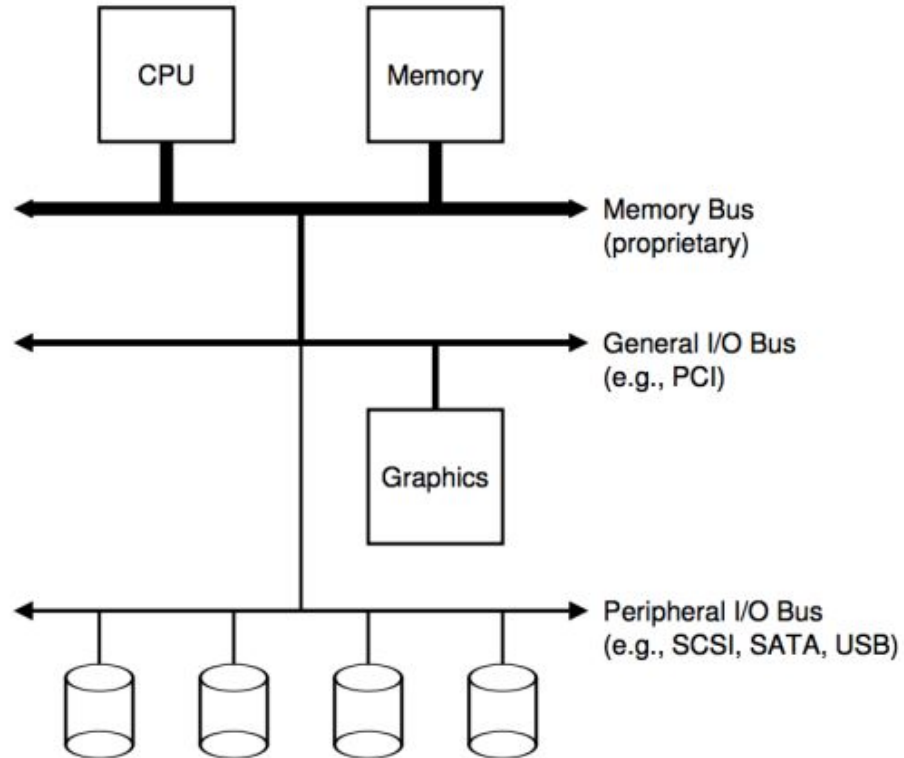
- Add more physical memory, what happens to performance?
- LRU, OPT:
 - Guaranteed to have fewer (or same number of) page faults
 - Smaller memory sizes are guaranteed to contain a subset of larger memory sizes
 - Stack property: smaller cache always subset of bigger
- FIFO:
 - Usually have fewer page faults
 - Belady's anomaly: May actually have more page faults!

I/O Devices

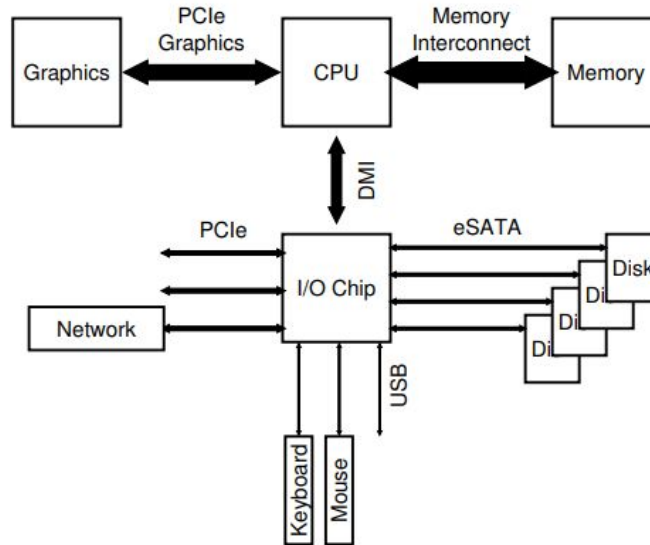


- Three conceptual pieces
 - 1. Virtualization
 - Make each application believe it has each resource to itself CPU and Memory
 - 2. Concurrency
 - Provide mutual exclusion, ordering
 - 3. Persistence

System Architecture

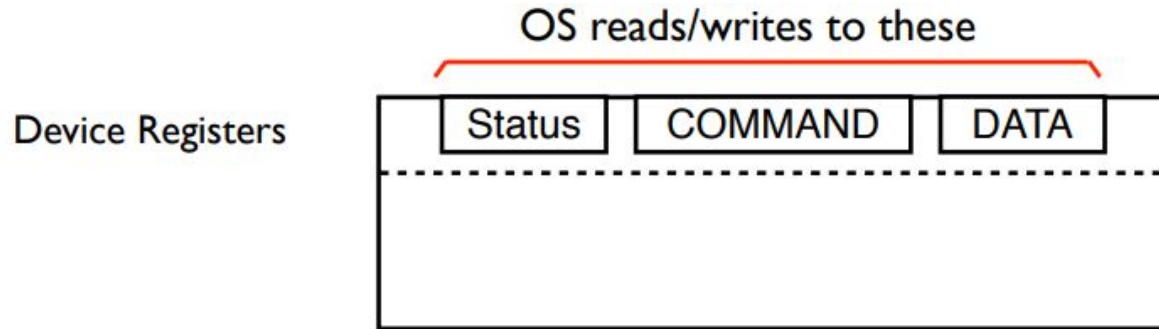


System Architecture




- Modern systems increasingly use specialized chipsets and faster point-to-point interconnects to improve performance

Canonical Device




```
While (STATUS == BUSY)
    ; // wait until device is not busy
Write data to DATA register
Write command to COMMAND register
    (starts the device and executes the command)
While (STATUS == BUSY)
    ; // wait until device is done with your request
```


Lowering CPU Overhead With Interrupts




- P1 issues an I/O request
- CPU spins




- P1 issues an I/O request
- P2 on the CPU while the disk serves P1's request.

Data Movement with DMA



- P1 wishes to write some data to the disk
- CPU copies



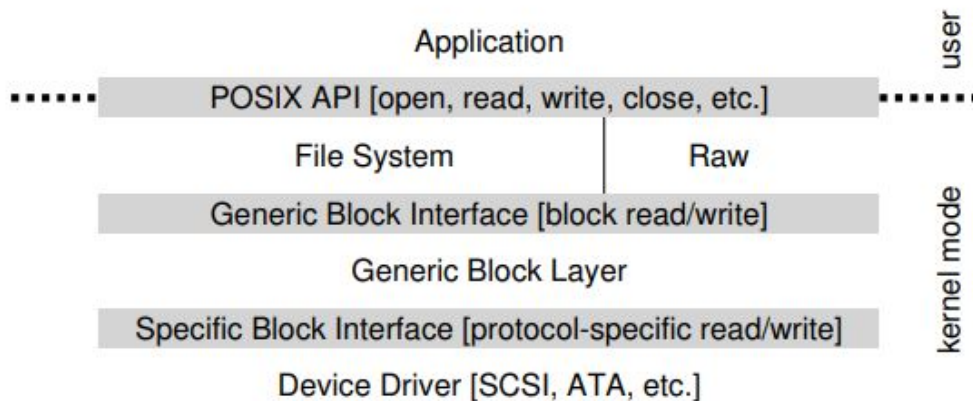
- DMA controller handles the copying
- DMV controller raises an interrupt when finishes

Programmed I/O VS. Directed Memory Access



- PIO (Programmed I/O):
 - CPU directly tells device what the data is
- DMA (Direct Memory Access):
 - Orchestrate transfers between devices and main memory without much CPU intervention
 - CPU leaves data in memory
 - Device reads data directly from memory

Device Drivers



- **a file system** is completely oblivious to the specifics
- **a raw interface** to devices, which enables special applications to directly read and write blocks without using the file abstraction

Hard Disks




- Disk has a sector-addressable address space
 - Appears as an array of sectors
- Sectors are typically 512 bytes
- Main operations: reads + writes to sectors
- Mechanical and slow

Hard Disks




Platter



Hard Disks




Spindle




Hard Disks

- Motor connected to spindle spins platters
- Rate of rotation
 - RPM (rotation per minute)
- 10000 RPM ==>
 - 1 min == 60 sec == 60,000 ms
 - single rotation is 6 ms




Hard Disks

- Surface is divided into rings:
tracks
- Stack of tracks(across platters):
cylinder




Hard Disks

- **Tracks** are further divided into numbered **sectors**




Hard Disks

- **Heads** on a moving **arm** can read from each surface.



Hard Disks

- Operations:
 - seek
 - read
 - write
- Must specify
 - cylinder # (distance from spindle)
 - head #
 - sector #
 - transfer size
 - memory address





THANK YOU!