



ILLINOIS TECH

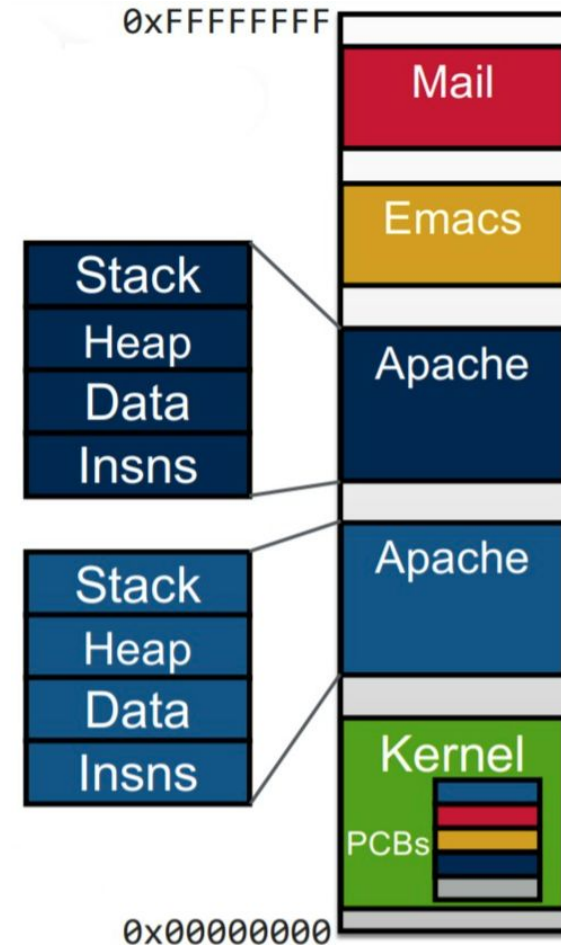
College of Computing

CS 450 Operating Systems Threads

Yue Duan

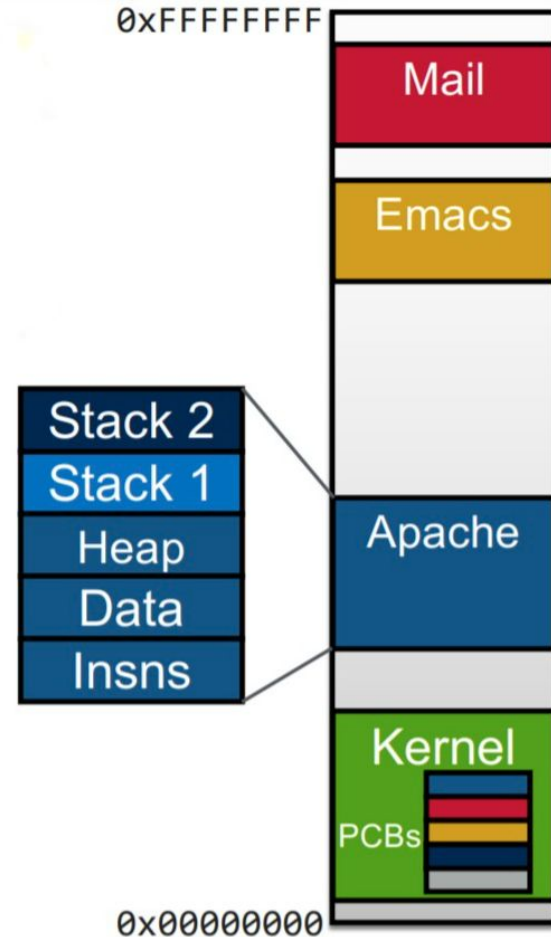
Consider..

- Apache wants to run multiple concurrent computations
 - Two heavyweight address spaces for two concurrent computations
 - Hard to share cache, etc.
-
- *Physical address space*
 - *Each process' address space by color (shown contiguous to look nicer)*



Improvement idea

- Place concurrent computations in the same address space!



Thread



- Also known as
 - Lightweight Process
 - Thread of Control
 - Task
- Major differences between thread and process:
 - lightweight vs. heavyweight
 - share data vs. isolation

Process vs. Thread



- A process is an abstraction of a computer
 - CPU, memory, devices
- A thread is an abstraction of a core
 - registers (incl. PC and SP)

Unbounded #computers, each with unbounded #cores

- Different processes typically have their own (virtual) memory, but different threads share virtual memory.
- Different processes tend to be mutually distrusting, but threads must be mutually trusting. **Why?**

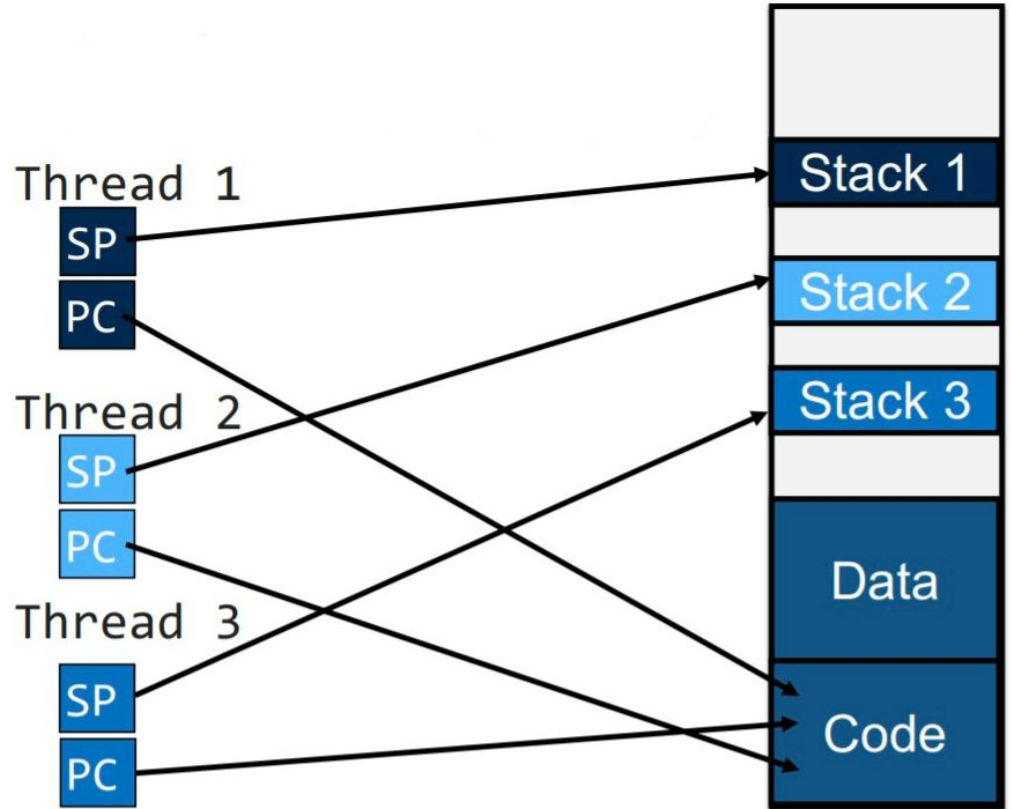
Process vs. Thread



- **Multiple threads** within a single process **share**:
 - Address space
 - Code (instructions)
 - Most data (heap)
 - Open file descriptors
 - Current working directory
 - User and group id
- **Each thread has its own**
 - Thread ID (TID)
 - Set of registers, including Program counter and Stack pointer
 - Stack for local variables and return addresses (in same address space)

Virtual Memory Layout

- Thread stacks are allocated on the heap!



Why Threads?



- Concurrency
 - exploiting multiple CPUs/cores
- Mask long latency of I/O
 - doing useful work while waiting
- Responsiveness
 - high priority GUI threads / low priority work threads
- Encourages natural program structure
 - expressing logically concurrent tasks
 - update screen, fetching data, receive user input

Example



- Web server:
 - 1. get network message (URL) from client
 - 2. get URL data from disk
 - 3. compose response
 - 4. send response

```
for (k = 0; k < n; k++) {  
    a[k] = b[k] × c[k] + d[k] × e[k]  
}
```

Simple Thread API



<code>void thread_create (func, arg)</code>	Creates a new thread that will execute function func with the arguments arg
<code>void thread_yield()</code>	Calling thread gives up processor. Scheduler can resume running this thread at any point.
<code>void thread_exit()</code>	Finish caller

Preemption



- Two kinds of threads:
 - Non-preemptive:
 - explicitly yield to other threads
 - Preemptive:
 - yield automatically upon clock interrupts
- Most modern threading systems are preemptive

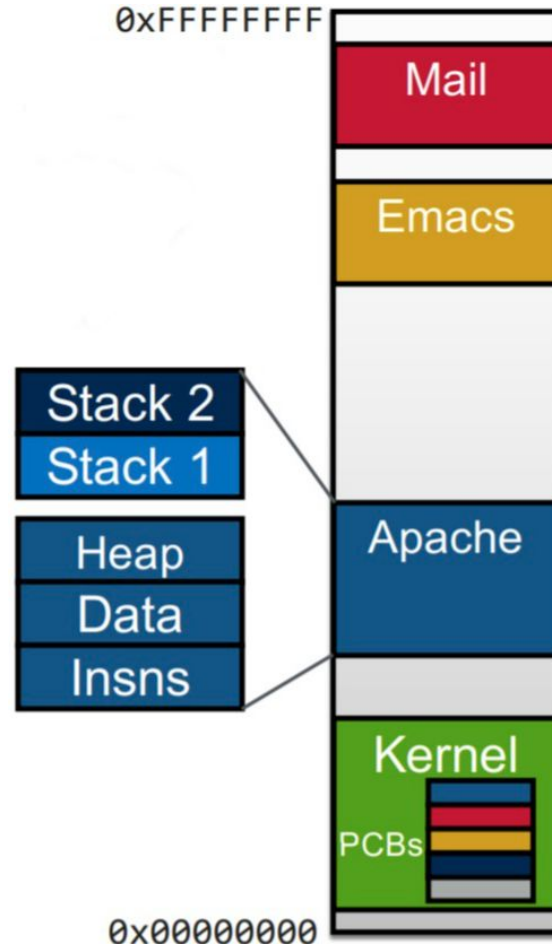
Implementation of Threads



- One abstraction, two implementations:
- 1. “kernel threads”:
 - each thread has its own PCB in the kernel, but the PCBs point to the same physical memory
- 2. “user threads”:
 - one PCB for the process; threads implemented entirely in user space.
 - each thread has its own Thread Control Block (TCB)

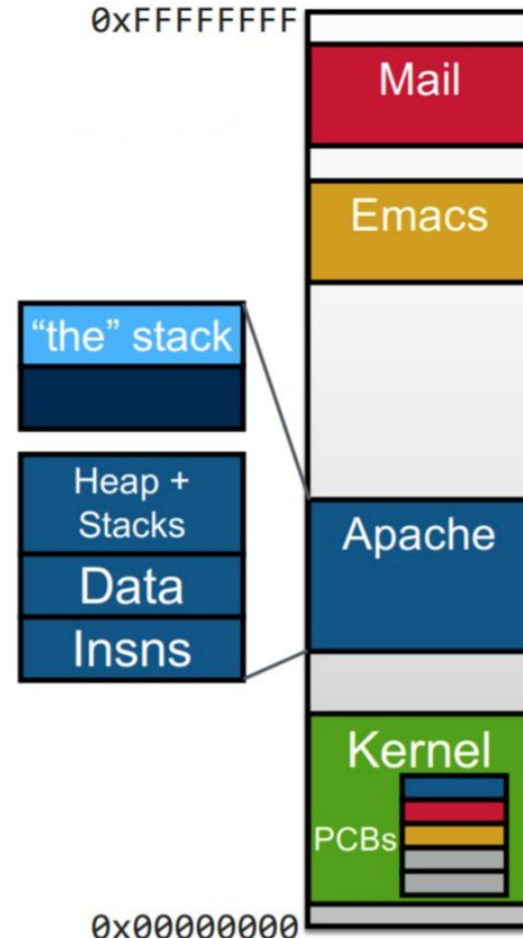
Kernel-Level Threads

- Kernel knows about, schedules threads (just like processes)
- Separate PCB for each thread
- PCBs have:
 - **same:** page table base register
 - **different:** PC, SP, registers, interrupt stack



User-Level Threads

- Run mini-OS in user space
 - Real OS unaware of threads
 - Single PCB
 - Thread Control Block (TCB) for each thread
- Generally more efficient than kernel-level threads (**Why?**)
- But kernel-level threads simplify system call handling and scheduling (**Why?**)

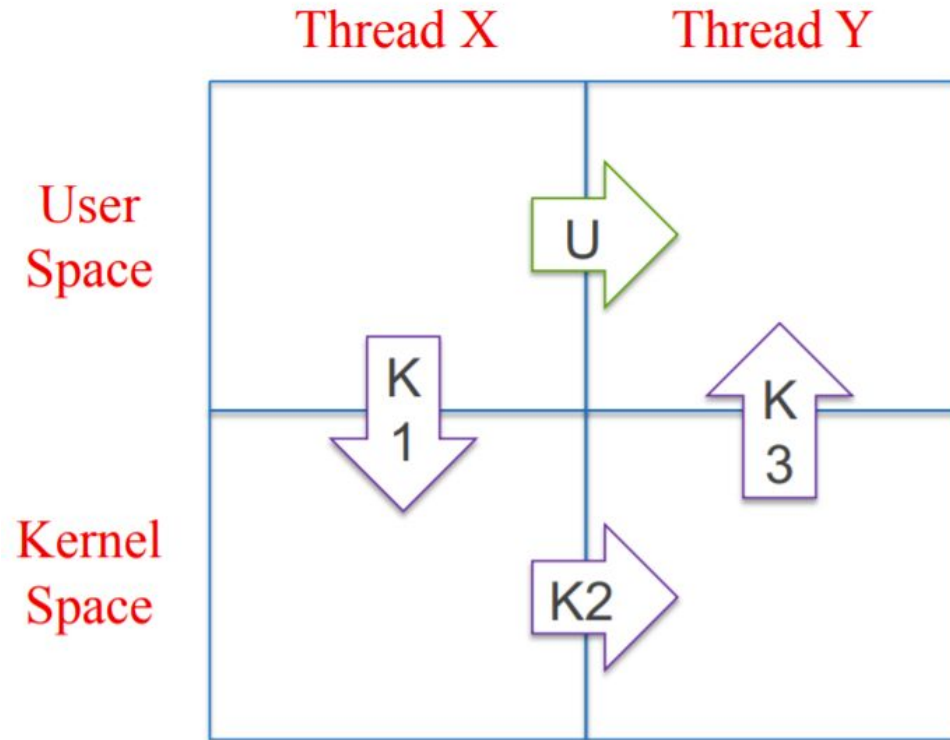


Kernel- vs User-level Threads



Kernel-Level Threads	User-level Threads
<ul style="list-style-type: none">• Easy to implement: just processes with shared page table	<ul style="list-style-type: none">• Requires user-level context switches, scheduler
<ul style="list-style-type: none">• Threads can run blocking system calls concurrently	<ul style="list-style-type: none">• Blocking system call blocks all threads: needs O.S. support for non-blocking system calls
<ul style="list-style-type: none">• Thread switch requires three context switches	<ul style="list-style-type: none">• Thread switch efficiently implemented in user space

Kernel vs User Thread Switch



Thread Schedule 1

- balance = balance + 1; //balance at 0x9cd4

State:

0x9cd4: 100

%eax: ?

%rip = 0x195

process
control
blocks:

Thread 1

%eax: ?
%rip: 0x195

Thread 2

%eax: ?
%rip: 0x195

T1 ➡
0x195 mov 0x9cd4, %eax
0x19a add \$0x1, %eax
0x19d mov %eax, 0x9cd4

Thread Schedule 1

- balance = balance + 1; //balance at 0x9cd4

State:

0x9cd4: 100

%eax: 100

%rip = 0x19a

process
control
blocks:

Thread 1

%eax: ?
%rip: 0x195

Thread 2

%eax: ?
%rip: 0x195

T1 →
0x195 mov 0x9cd4, %eax
0x19a add \$0x1, %eax
0x19d mov %eax, 0x9cd4

Thread Schedule 1

- balance = balance + 1; //balance at 0x9cd4

State:

0x9cd4: 100

%eax: 101

%rip = 0x19d

process
control
blocks:

Thread 1

%eax: ?
%rip: 0x195

Thread 2

%eax: ?
%rip: 0x195

0x195 mov 0x9cd4, %eax

0x19a add \$0x1, %eax

T1 ➡ 0x19d mov %eax, 0x9cd4

Thread Schedule 1

- balance = balance + 1; //balance at 0x9cd4

State:

0x9cd4: 101

%eax: 101

%rip = 0x1a2

process
control
blocks:

Thread 1

%eax: ?
%rip: 0x195

Thread 2

%eax: ?
%rip: 0x195

0x195 mov 0x9cd4, %eax

0x19a add \$0x1, %eax

0x19d mov %eax, 0x9cd4

T1 →

Thread Context Switch

Thread Schedule 1

- balance = balance + 1; //balance at 0x9cd4

State:

0x9cd4: 101

%eax: ?

%rip = 0x195

process
control
blocks:

Thread 1

%eax: 101
%rip: 0x1a2

Thread 2

%eax: ?
%rip: 0x195

T2 ➡

```
0x195  mov 0x9cd4, %eax
0x19a  add $0x1, %eax
0x19d  mov %eax, 0x9cd4
```

Thread Schedule 1

- balance = balance + 1; //balance at 0x9cd4

State:

0x9cd4: 101

%eax: 101

%rip = 0x19a

process
control
blocks:

Thread 1

%eax: 101
%rip: 0x1a2

Thread 2

%eax: ?
%rip: 0x195

T2



0x195 mov 0x9cd4, %eax

0x19a add \$0x1, %eax

0x19d mov %eax, 0x9cd4

Thread Schedule 1

- balance = balance + 1; //balance at 0x9cd4

State:

0x9cd4: 101

%eax: 102

%rip = 0x19d


process
control
blocks:

Thread 1

%eax: 101
%rip: 0x1a2

Thread 2

%eax: ?
%rip: 0x195



T2

0x195 mov 0x9cd4, %eax

0x19a add \$0x1, %eax

0x19d mov %eax, 0x9cd4

Thread Schedule 1

- balance = balance + 1; //balance at 0x9cd4

State:

0x9cd4: 102

%eax: 102

%rip = 0x1a2

process
control
blocks:

Thread 1

%eax: 101
%rip: 0x1a2

Thread 2

%eax: ?
%rip: 0x195

0x195 mov 0x9cd4, %eax

0x19a add \$0x1, %eax

0x19d mov %eax, 0x9cd4



Desired result!

Thread Schedule 2

- `balance = balance + 1; //balance at 0x9cd4`

State:

`0x9cd4`: 100

`%eax`: ?

`%rip` = 0x195

process
control
blocks:

Thread 1

`%eax`: ?
`%rip`: 0x195

Thread 2

`%eax`: ?
`%rip`: 0x195

T1 ➡

```
0x195  mov 0x9cd4, %eax
0x19a  add $0x1, %eax
0x19d  mov %eax, 0x9cd4
```

Thread Schedule 2

- balance = balance + 1; //balance at 0x9cd4

State:

0x9cd4: 100

%eax: 100

%rip = 0x19a

process
control
blocks:

Thread 1

%eax: ?
%rip: 0x195

Thread 2

%eax: ?
%rip: 0x195

T1 →
0x195 mov 0x9cd4, %eax
0x19a add \$0x1, %eax
0x19d mov %eax, 0x9cd4

Thread Schedule 2

- balance = balance + 1; //balance at 0x9cd4

State:

0x9cd4: 100

%eax: 101

%rip = 0x19d

process
control
blocks:

Thread 1

%eax: ?
%rip: 0x195

Thread 2

%eax: ?
%rip: 0x195

0x195 mov 0x9cd4, %eax
0x19a add \$0x1, %eax
T1 → 0x19d mov %eax, 0x9cd4

Thread Context Switch

Thread Schedule 2

- balance = balance + 1; //balance at 0x9cd4

State:

0x9cd4: 100

%eax: ?

%rip = 0x195

process
control
blocks:

Thread 1

%eax: 101
%rip: 0x19d

Thread 2

%eax: ?
%rip: 0x195

T2 ➡

```
0x195  mov 0x9cd4, %eax
0x19a  add $0x1, %eax
0x19d  mov %eax, 0x9cd4
```

Thread Schedule 2

- balance = balance + 1; //balance at 0x9cd4

State:

0x9cd4: 100

%eax: 100

%rip = 0x19a

process
control
blocks:

Thread 1

%eax: 101
%rip: 0x19d

Thread 2

%eax: ?
%rip: 0x195

T2



0x195 mov 0x9cd4, %eax

0x19a add \$0x1, %eax

0x19d mov %eax, 0x9cd4

Thread Schedule 2

- balance = balance + 1; //balance at 0x9cd4

State:

0x9cd4: 100

%eax: 101

%rip = 0x19d

process
control
blocks:

Thread 1

%eax: 101
%rip: 0x19d

Thread 2

%eax: ?
%rip: 0x195

0x195 mov 0x9cd4, %eax

0x19a add \$0x1, %eax



T2

0x19d mov %eax, 0x9cd4

Thread Schedule 2

- balance = balance + 1; //balance at 0x9cd4

State:

0x9cd4: 101

%eax: 101

%rip = 0x1a2

process
control
blocks:

Thread 1

%eax: 101
%rip: 0x19d

Thread 2

%eax: ?
%rip: 0x195

0x195 mov 0x9cd4, %eax

0x19a add \$0x1, %eax

0x19d mov %eax, 0x9cd4



Thread Context Switch

Thread Schedule 2

- balance = balance + 1; //balance at 0x9cd4

State:

0x9cd4: 101

%eax: 101

%rip = 0x19d

process
control
blocks:

Thread 1

%eax: 101
%rip: 0x19d

Thread 2

%eax: 101
%rip: 0x1a2



T1

0x195 mov 0x9cd4, %eax

0x19a add \$0x1, %eax

0x19d mov %eax, 0x9cd4

Thread Schedule 2

- balance = balance + 1; //balance at 0x9cd4

State:

0x9cd4: 101

%eax: 101

%rip = 0x1a2

process
control
blocks:

Thread 1

%eax: 101
%rip: 0x1a2

Thread 2

%eax: 101
%rip: 0x1a2

0x195 mov 0x9cd4, %eax

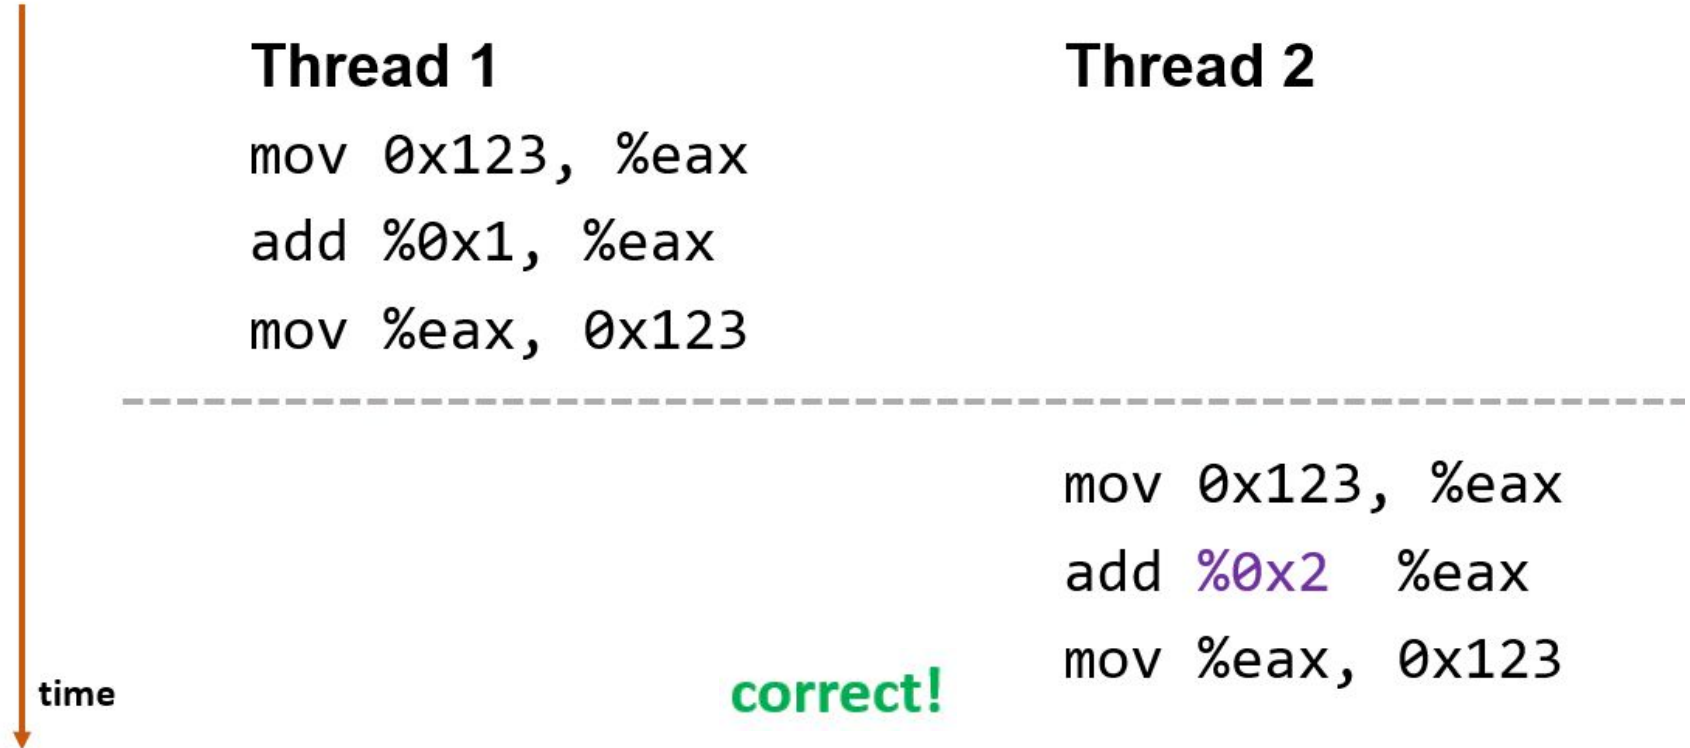
0x19a add \$0x1, %eax

0x19d mov %eax, 0x9cd4


T1 →

WRONG RESULT! Final balance value is 101


Timeline View: Interleaving 1




Timeline View: Interleaving 2




Timeline View: Interleaving 3



Timeline View: Interleaving 4



Timeline View: Interleaving 5



Non-Determinism



- Concurrency leads to non-deterministic results
 - Not deterministic result: different results even with same inputs
 - **race conditions**

What do we want?

- Want 3 instructions to execute as an uninterruptible group
- That is, we want them to be an **atomic** unit

```
mov 0x123, %eax  
add %0x1, %eax  
mov %eax, 0x123
```

critical section

- More general:
 - Need mutual exclusion for critical sections
 - if process A is in critical section C, process B can't be (okay if other processes do unrelated work)



THANK YOU!