

6 points ✓ Saved

- ☒ FIFO and SJF are for non-preemptive scheduling, and RR and SRJF are for preemptive scheduling

- ☒ Compared to FIFO, RR has shorter responsive time but larger turnaround time.
- ☐ RR and SRJF are starvation free, because they are preemptive.
- ☐ SRJF has better job throughput than RR.
- ☐ For RR, a smaller time slice means shorter response time as well as shorter turnaround time
- ☐ Compared to SJF, SRJF has shorter response time and larger turnaround time

7 points  Saved

- ☒ NFS can co-exist with other file systems.
- ☐ NFS's Close-to-Open consistency model guarantees that there is no inconsistency.
- ☒ In NFS, not all requests from clients can be made idempotent.
- ☒ Generation # in NFS will not be decremented.
- ☐ The major reason for FFS to be faster is that it treats disks like RAM.
- ☒ Groups in FFS can make nodes closer to data blocks, and leads to better space efficiency.
- ☒ Smart allocation policy in FFS is to improve runtime performance.

6 points  Saved

- ☐ Integrity property in security prevents sensitive information from unauthorized access.
- ☒ RAID-5 is strictly better than RAID-4.
- ☐ Access control list of a file should be stored only in memory.
- ☐ RAID-1 has the best redundancy among RAID-0, 1, 4, 5.
- ☐ Salting in password storage can make attacks harder.
- ☒ Demand paging in page selection will trigger page fault for every newly accessed page.

8 points ✓ Saved

For the toolbar, press ALT+F10 (PC) or ALT+FN+F10 (Mac).

When we run "sleep 3" the shell will temporarily pause before returning the prompt back. Then running "exec sleep 3" will also temporarily pause the shell however it will not return the prompt back and the terminal session ends.

The difference between the two commands is the fact that when you run a command the shell forks this command into a new subshell process and terminates this subshell when the command completes which allows the parent shell to run without issues. When you use the "exec" command it runs the command directly on the parent shell and therefore when the command terminates it terminates the parent shell which ran the process.

9 points ✓ Saved

```
1). web_server() {
    2). while (1) {
        int sock = accept();
        3). thread_cirete(handle_request, sock);
    }
}

handle_request(int sock) {
    ...process request
    close(sock);
    4).
}
```

```
B I U Paragraph Arial 10pt
web_server(){
    2). sem_init(&threads_left, 99)
    while (1){
        int sock = accept();
        3). sem_wait(&threads_left)
        thread_cirete(handle_request, sock);
    }
}

handle_request(int sock){
    ...process request
    close(sock);
    4). sem_post(&threads_left)
}
```

10 points ✔ Saved

{1, 22, 14, 72, 86, 32, 11, 66, 45, 80}

Given that the arm head is currently positioned at block 75 and is in the midst of moving in the direction towards block 0, what is the sequence in which the enqueued blocks will be read with the i) SCAN, and ii) C-SCAN algorithms (picks up requests on way down only)?

For the toolbar, press ALT+F10 (PC) or ALT+Fn+F10 (Mac).

i) Scan - rotates full way through

Arm Head Starting at Block 75 moving towards Block 0

75 -> 72 -> 66 -> 45 -> 32 -> 22 -> 14 -> 11 -> 1 -> 80 -> 86

i) C-Scan - Gets from starting to end then restarts from beginning

75 -> 72 -> 66 -> 45 -> 32 -> 22 -> 14 -> 11 -> 1 -> 86 -> 80

10 points  Saved

- What is the largest file size that can be indexed in this system?
- How many blocks (including indirect blocks) are needed to address a file of size 100 bytes, 10K bytes and 4G bytes?

7 a. What is the largest file size that can be indexed in this system?

10 Direct Pointers which points to 8K data blocks:  $10 * 8K = 80K$   
1 Single Direct Pointer which points to 8K index block, index block holds:  $8K/4 = 2K$  pointers to data blocks  
Single Indirect Direct Pointer indirectly points to:  $2K * 8K = 16M$   
Double Indirect Pointer indirectly points to  $2K * 2K * 8K = 32G$   
**Sum all pointers:  $(80K + 16M + 32G) = \text{Maximum File Size}$**

b.

100 B = 0.1 KB

**1 Block is enough**

10KB

**2 Blocks are enough**

4G = 4194304 KB

First 80KB stored in 10 blocks, so remaining  $4194304 - 80 = 4194224$  KB

$4194224 \text{ KB} / 8KB = 524278$ , since a single-indirect block can store  $8KB/4B = 20K$ , we only need 1 single-indirect block.

So total # of blocks  $= 10 + 1 + 524278 = \mathbf{524289 \text{ Blocks are enough}}$

## QUESTION 8

10 points ✓ Saved

Explain the first 10 steps in the timeline below for file write.

	data bitmap	inode bitmap	root inode	foo inode	bar inode	root data	foo data	bar data [0]	bar data [1]	bar data [2]
create (/foo/bar)		read write	read	read		read	read			
					read write		write			
write()	read write			read				write		
write()	read write			write read					write	
write()	read write			write read						write

For the toolbar, press ALT+F10 (PC) or ALT+FN+F10 (Mac).

B I U Paragraph Arial 10pt
   
 (1) read root inode to locate root data
   
 (2) read root data to lookup foo and find foo's inode
   
 (3) read foo inode to locate foo data
   
 (4) read foo data to find bar's inode
   
 (5) read inode bitmap data to find a free inode
   
 (6) write inode bitmap to initialize the inode to identify that the inode bitmap is used
   
 (7) write foo data to update entries due to addition of D entry of bar
   
 (8) read bar inode to retrieve bar's metadata as part of file open operation
   
 (9) write bar inode to update its last access timestamp
   
 (10) write foo inode to update metadata of directory's inode, size, and data

P 113 WORDS POWERED BY TINY

## QUESTION 9

10 points ✓ Saved

Assume a single-level page table system with 4KB page size, 64-bit address and 8-byte PTE.

a. How many pages are needed?

b. How much space would the page table take up? Hint: think about how big the address space is; use power-of-two math.

For the toolbar, press ALT+F10 (PC) or ALT+FN+F10 (Mac).

B I U Paragraph Arial 10pt
   
 a) Offset Bits = 12 because  $2^{12} = 4096 = 4\text{KB Page Size}$ 
  
 $64 - 12 = 52$  bits for page numbers
   
**How many pages are needed?:  $2^{52}$  pages**
  
 b)  $2^{52} \times 8 \text{ Bytes} = 36 \text{ (PB) Petabytes}$ 
  
**How much space would the page table take up? 36PB**

P = STRONG 52 WORDS POWERED BY TINY

9 points  Saved

- What is the average rotation time?
- What is the average transfer time?
- What is the disk access time?

a) It takes  $1 / 5400 \text{ RPM}$  for one revolution =  $(1000 \text{ ms/sec} * 60 \text{ sec/min}) / 5400 \text{ RPM} = 11.11 \text{ ms}$   
 The average rotation time is  $1/2$  the time for a complete revolution so  $11.11 \text{ ms} / 2 = 5.56 \text{ ms}$   
**What is the average rotation time? 5.56ms**

b) We have 500 sectors per track. So we use  $1/500$  revolutions of the disk to read the full sector. This is  $11.11 \text{ ms} / 500 = 0.002222 \text{ ms}$   
**What is the average transfer time? 0.002222ms**

c) The disk access time is the sum of average seek time, average rotation time, and average transfer time.  $7 \text{ ms} + 5.56 \text{ ms} + 0.002222 \text{ ms} = 12.56 \text{ ms}$   
**What is the disk access time? 12.56ms**

15 points  Saved

VPN:0	0x0
PageTable	0x0
	0x1
	0x9
	0x7
	0x8
	0
	:
VPN:15	0

```
0x3000: load 0x5320, %eax
0x3004: load 0x4004, %ebx
0x3008: mul %ecx, %eax, %ebx
```

- Offset is 12 bits (because 4KB =  $2^{12}$  bytes)  
VPN is 20 bits  
PFN: 3 = 0x9000
- Physical Memory Accesses with Paging  
Fetch instruction at virtual addr 0x3000  
Mem ref 1: 0x2000 -> VPN 3 -> 1  
Fetch instruction at 0x1000 (Mem ref 2)
- Exec. load from virtual addr 5230 into %eax  
Access page table to get VPN 5->7  
Mem ref 3: 0x2014  
load 0x7320 into %eax
- Fetch instruction = 2  
Fetch val = 2
- The next two instructions will have 4 references and the final instruction will have 6 references
- The total references this is 10 references.**
- b)
- Offset is 12 bits (because 4KB =  $2^{12}$  bytes)  
VPN is 20 bits  
PFN: 3 = 0x9000
- Physical Memory Accesses with Paging  
Fetch instruction hits TLB at virtual addr 0x3000  
Mem ref 1: 0x2000 -> VPN 3 -> 1 [2 Memory References]  
Fetch instruction at 0x1000 (Mem ref 2)
- Exec. load from virtual addr 5230 into %eax hits TLB  
Access page table to get VPN 5->7 [2 Memory References]  
Mem ref 3: 0x2014
- Fetch instruction = 2  
Fetch instruction hits TLB at virtual addr 0x3000 [Note every time we load the references will go from 2->1]
- Fetch val = 2
- The next two instructions will be 3 and the final instruction will be 1 references
- The total references this is 6 references.**